

Title Page

The image shows a book cover for 'Omnillm: The Complete Book'. The cover is white with a blue border. The title 'Omnillm' is in large, bold, black letters. Below it, 'THE COMPLETE BOOK' is in smaller, bold, black letters. The subtitle 'Compare, Route, Orchestrate Every LLM - and Put Them Inside a Robot.' is in a smaller, black font. The author's name 'Akshita-sr' and the year '2026' are at the bottom.

Copyright & Licence

This work documents the **OmniLLM** open-source project (MIT License, copyright © 2026 Akshita-sr) and the *Embodied LLM Arena* research study that the project supports.

This book is intended as **educational and reference material** for the project's owner, future collaborators, thesis examiners, and anyone learning how to build a multi-LLM platform connected to a social robot.

The source code is freely available at <https://github.com/Akshita-sr/OmniLLM>.

Dedication

To everyone who looked at a chatbot and asked — "yes, but what would happen if I plugged it into a robot?"

Preface — How to Read This Book

This book is unusual in one important way: **it is written for three different readers at the same time.**

When you began this project, you were one of those readers. By the time you finish reading the book, you will have walked through the perspective of all three — and you will be able to *talk* to the project, the code, and the research community fluently in any of the three languages.

The Three Readers

Reader	Who they are	What they need
The Owner	You, the person who built this and may now have forgotten parts of it	A quick-take recap, "why did we do X this way", file paths, fast orientation
The Beginner	A friend or junior who has just learned Python and wants to understand the project from zero	Plain English explanations, mini-primers on jargon, worked examples, no skipping steps
The Academic	A thesis examiner, an HRI researcher, an LLM engineer evaluating the work	Motivation, literature context, formal evaluation, why the methodology is sound

The Layered Chapter Structure

Every chapter (from Part II onwards) follows the same three-layer structure so you can read it the way that fits your purpose today:

+-----+	
* AT A GLANCE (Owner's Recap)	
One-sentence summary, file paths, the "why we built it this way"	
+-----+	
(Book) THE WALK-THROUGH (Beginner's Path)	
Plain language, glossary inline, examples you can run,	
Python concepts explained where they appear	
+-----+	
(Academic) ACADEMIC CONTEXT (Researcher's Sidebar)	
Literature, citations, research justification,	
"why this approach is methodologically defensible"	
+-----+	

How to Use This Book

If you are...	Read...
Re-learning your own project	Skim the "At a Glance" boxes; dive into the modules you have forgotten
Setting up the project on a new machine	Part V (Setup & Running) front-to-back
Onboarding a collaborator	Hand them Parts I, IV, and the relevant module chapter from Part III
Writing a thesis or paper	Read the "Academic Context" sidebars; consult Part VII for the experimental design
Connecting to a Pepper robot for the first time	Part IV, then Part V Chapter 12
Trying to understand a single Python file	Find it in the index in Part III; each module has its own chapter
Just curious	Read the preface, then skip to Chapter 1

A Note on Notation

Throughout the book:

- **Bold** marks a term being introduced for the first time. The first introduction is always a definition — not just a use of the word.
- `code-style monospace` is used for filenames, command lines, function names, and code identifiers.

Block-quoted text contains direct quotes from the codebase or short excerpts that are too important to paraphrase.

- (Tip) Tips and rules of thumb are marked with the lightbulb glyph.
- (!) Warnings and gotchas are marked with the warning glyph.
- (Academic) Academic-only material is marked with the mortar-board glyph; you may skim it on a first read.

What This Book Does Not Cover

To keep the book finite, the following are **not** included in detail:

- Generic Python language tutorials — Appendix A has a focused primer for the parts you'll see in this codebase, but for a full tutorial we point you to the Python documentation.
- Generic Linux/Windows administration.
- The internal implementation details of LangGraph, LiteLLM, ChromaDB, or the NAOqi C++ middleware. We use these as black boxes and focus on how OmniLLM uses them.

- The full text of the OmniLLM source code (it lives in the repository). The book quotes excerpts; the code is the source of truth.

Table of Contents

Part I — Foundations (Chapters 1–4)

- **Chapter 1** What Is OmniLLM, Really?
- **Chapter 2** Why a Robot? The Motivation Behind This Project
- **Chapter 3** The Five-Minute Tour of the Repository
- **Chapter 4** The Vocabulary You Need Before Chapter 5

Part II — Architecture & Data Flow (Chapters 5–8)

- **Chapter 5** The Three-Layer System Architecture
- **Chapter 6** The Lifecycle of a Single Spoken Question
- **Chapter 7** The LangGraph Agent Pipeline, Node by Node
- **Chapter 8** Where Every Byte Goes — Endpoint and Provider Map

Part III — Module Deep-Dive (Chapters 9–22)

- **Chapter 9** `gateway.py` — The Door to Every LLM
- **Chapter 10** `router.py` — The Smart Model Selector
- **Chapter 11** `consensus.py` — The LLM Council
- **Chapter 12** `evaluator.py` — The LLM-as-Judge Pipeline
- **Chapter 13** `scorer.py` — The ELO Leaderboard
- **Chapter 14** `cli.py` — Your Terminal Dashboard
- **Chapter 15** `rag/pipeline.py` — Retrieval-Augmented Generation
- **Chapter 16** `hri/classifier.py` and `language_detector.py`
- **Chapter 17** `hri/agent_graph.py` — The Brain of the Robot
- **Chapter 18** `hri/experiment.py` — Conditions, Sessions, Counterbalancing
- **Chapter 19** `robotics/` — Bridges, Gestures, Whisper
- **Chapter 20** `server/app.py` — The Flask AI Server
- **Chapter 21** `server/naoqi_client.py` — The Python 2.7 Side
- **Chapter 22** `utils/` — Logging, Cost Tracking, Questionnaires, Export

Part IV — Pepper, Choregraphe, NAOqi (Chapters 23–27)

- **Chapter 23** Meet Pepper — The Hardware Inside the Plastic Shell

- **Chapter 24** NAOqi 101 — The Operating System That Runs on Pepper
- **Chapter 25** Choregraphe — The Visual Programming Studio
- **Chapter 26** The Python 2.7 / Python 3.x Bridge Problem (and the Solution)
- **Chapter 27** Five Practical Walk-throughs (with and without a real robot)

Part V — Setup & Running (Chapters 28–32)

- **Chapter 28** Installing OmniLLM From Zero
- **Chapter 29** Running Without a Robot (The Default Path)
- **Chapter 30** Running With Choregraphe's Virtual Robot
- **Chapter 31** Running With a Physical Pepper
- **Chapter 32** Running an Experimental Session With a Real Participant

Part VI — Building This From Scratch (Chapters 33–35)

- **Chapter 33** The Build Order — A 13-Week Plan
- **Chapter 34** The Critical First Mile — Foundations Before Anything Else
- **Chapter 35** When (and How) to Add the Robot

Part VII — Features, Scope, Evaluation (Chapters 36–39)

- **Chapter 36** The Complete Feature Reference
- **Chapter 37** The Five Experimental Conditions, in Detail
- **Chapter 38** Evaluation Methodology — Eight Axes, Three Judge Patterns
- **Chapter 39** Scope, Limits, and What Comes Next

Part VIII — Appendices

- **Appendix A** Python Primer — Just the Parts You Need for OmniLLM
- **Appendix B** Complete Terminal Command Reference
- **Appendix C** Complete Glossary
- **Appendix D** Troubleshooting Reference
- **Appendix E** External Resources — Papers, Libraries, Repositories
- **Appendix F** File Index — Every File in the Repository, One Line Each

PART I — FOUNDATIONS

Chapter 1 — What Is OmniLLM, Really?

*** AT A GLANCE. *OmniLLM is a Python platform that lets you talk to many different AI models (OpenAI, Anthropic, Google, DeepSeek, Ollama-local) through one unified interface, automatically picks the best one for each kind of question, can ask several at once and synthesise their answers, scores them on an ELO leaderboard, grounds their replies in your own documents (RAG), and connects all of that to a Pepper humanoid robot for human-robot interaction research. The core innovation: the robot's brain is never locked to a single model**.*

1.1 The One-Sentence Definition

OmniLLM is a multi-LLM orchestration platform with a robot front-end.

That sentence is technically correct, but it hides everything interesting. Let us unpack it for each of our three readers.

1.2 For the Beginner — A Working Analogy

Imagine you are a small startup that needs legal advice. You could hire one lawyer and use them for everything — divorce, tax, intellectual property, employment law. They will be okay at all of those, but excellent at none.

Now imagine, instead, that you have an *office manager* with a Rolodex. The office manager:

1. Listens to your problem.
2. Decides which lawyer in the Rolodex is the right specialist.
3. Forwards the question.
4. Compares quotes from cheap and expensive lawyers when you are budget-bound.
5. When the question is high-stakes, sends it to *three* lawyers and synthesises their answers.
6. Keeps a running scoreboard of each lawyer's track record so the next routing decision is even smarter.

OmniLLM is that office manager. The lawyers are the LLMs.

Office-manager idea	OmniLLM equivalent
The Rolodex	<code>config/models.yaml</code> (the model registry)
Forwarding the question	<code>LLMGateway.query()</code>
Picking the right lawyer	<code>SmartRouter.route()</code>
Asking three lawyers at once	<code>ConsensusEngine.query_council()</code>
Grading the lawyer's answer	<code>Evaluator</code> (LLM-as-Judge)
The track-record scoreboard	<code>EloScorer</code>

The receptionist who repeats back what the lawyer said	The Pepper robot
--	------------------

1.3 For the Owner — What You Actually Built

You built a system with **five logical layers** that sit on top of each other:

```
+-----+
| 5. ROBOT FRONT-END — Pepper / NAO / Buddy / virtual / text-only |
+-----+
| 4. AGENT PIPELINE — LangGraph: STT → classify → RAG → LLM → plan |
+-----+
| 3. INTELLIGENCE LAYER — Router, Council, Evaluator, ELO Scorer |
+-----+
| 2. UNIFIED GATEWAY — LiteLLM, model registry, cost & latency |
+-----+
| 1. PROVIDER LAYER — OpenAI, Anthropic, Google, DeepSeek, Ollama |
+-----+
```

Each layer has a precise responsibility:

1. **Provider layer** — the actual LLM APIs. You did not write these. You wrote the way OmniLLM *talks* to them.
2. **Unified gateway** — your `gateway.py`. It hides every provider's quirks behind one Python coroutine call.
3. **Intelligence layer** — your `router.py`, `consensus.py`, `evaluator.py`, `scorer.py`. This is where "smart" happens.
4. **Agent pipeline** — your `hri/agent_graph.py`. A LangGraph DAG that turns raw audio into a robot action plan.
5. **Robot front-end** — your `robotics/`, `server/app.py`, `server/naoqi_client.py`. The bit that touches hardware (or simulates touching hardware).

(Tip) **Why this layering matters.** Each layer can be swapped without breaking the others. You can replace Ollama with vLLM (Layer 1) and the gateway above doesn't notice. You can replace LiteLLM with raw requests calls (Layer 2) and the router doesn't notice. You can swap Pepper for NAO (Layer 5) and the agent graph doesn't notice. **This separation is the entire engineering achievement of the project.**

1.4 For the Academic — The Research Contribution

(Academic) The current LLM evaluation literature lives almost entirely on text: *MMLU* (Hendrycks 2021), *MT-Bench* (Zheng 2023), *Chatbot Arena* (LMSYS 2023), *LiveBench* (White 2024). The HRI literature, in parallel, has explored single-LLM social robots: Irfan et al. (HRI 2024), Nichols et al. (2024 arXiv), Grassi et al. (HAI 2024), Spitale et al. (Vita, 2024).

No prior work does both at once. No published study has:

1. Compared **multiple** LLMs as interchangeable backends for the *same* social robot in the *same* HRI experimental protocol.
2. Applied **dynamic smart routing** between LLMs during live HRI based on task type, language, and budget.

3. Tested whether the **rankings** produced by embodied HRI agree with the rankings produced by text-only benchmarks.
4. Combined **RAG-augmented** social robotics with multi-LLM backend comparison.
5. Built a reusable **LangGraph-orchestrated** open-source platform for any HRI lab to repeat the study with their own robots, KB, and conditions.

The research gap is genuine; OmniLLM and the *Embodied LLM Arena* fill it.

1.5 Three Honest Questions and Answers

Q. Is OmniLLM "just" a LiteLLM wrapper? A. No. LiteLLM solves "how do I call any provider with one API". OmniLLM solves "given many providers, which should I call, why, with what context, through what robot, and how do I prove the answer is good?". LiteLLM is *Layer 2*. OmniLLM is layers 2 through 5.

Q. Does the robot actually do anything intelligent? A. The robot is a thin client. **All intelligence lives on the AI server.** Pepper records audio, sends it over HTTP, receives a JSON action plan, and executes the plan. This is by design. If the robot were the brain, swapping robots (NAO, Buddy) would mean rewriting the brain.

Q. Can I run this without a robot? A. Yes — and most development happens that way. The Flask server has a `/interact` endpoint that accepts plain text. You can talk to the system with `curl`. The robot is the demonstration target, not the prerequisite. Chapters 29 and 30 walk through robot-free use.

Chapter 2 — Why a Robot? The Motivation Behind This Project

*** AT A GLANCE. A reasonable critic asks "if you're benchmarking LLMs, why involve a 28-kg robot at all? You could test models on a screen." This chapter is the reply. The short version: physical embodiment changes how humans evaluate AI responses (the embodiment effect in HRI literature), exposes failure modes that screens hide (latency feel, gesture-speech sync), and matches the actual deployment context that organisations care about (reception robots, hospital wayfinders, educational companions).

2.1 The Skeptic's Question

You may have heard, or thought, some variant of:

"Adding a robot just adds complexity — Python 2.7 issues, network latency, physical logistics, hardware risk — without changing the underlying claim, which is about LLM quality. Why bother?"

This is a reasonable question. The honest answer takes four points.

2.2 Point 1 — The Embodiment Effect Is Real

Multiple decades of HRI research show that **the same AI text is evaluated differently** when it is delivered through:

- a screen (the typical chatbot)
- a disembodied voice (a smart speaker)
- a physical robot that gestures and makes eye contact

Key references:

Citation	Finding
Wainer et al. (2006)	Robots are perceived as more credible than equivalent on-screen agents on persuasion tasks
Li (2015), <i>IJSR</i> survey	Across 33 HRI experiments, embodied interaction produced reliably different ratings of trust, intelligence, and likeability
Bartneck et al. (2009), <i>Godspeed</i>	Established a validated multidimensional measurement instrument for robot perception
Bainbridge et al. (2011)	Physically present robots elicit greater compliance than the same robot via video

The mechanism is multi-causal: physical presence triggers social-norm evaluations that screens do not; gestures and gaze provide multimodal disambiguation; latency and timing matter more in face-to-face interaction.

The implication for OmniLLM is direct: **a model that scores best on MMLU may not score best when delivered via Pepper**. That divergence — if it exists — is itself a publishable finding.

2.3 Point 2 — Embodied Tasks Surface Failures Text Hides

Consider Task T2 ("Navigation / Guidance"). The user asks:

"Where is Room 305?"

A text-only system can produce a perfectly graded answer:

"Room 305 is on the third floor of Building C. Take the elevator on your left, then turn left at the corridor."

A robot must do more. It must produce that answer **and** point with the correct arm **and** turn its head **and** light its eyes a colour appropriate to navigation guidance. A model that produces the right text but the wrong gesture (says "left", points right) is a robot failure. **No text benchmark will ever detect this failure**. Multimodal coherence is invisible until the robot is actually moving.

Similarly: a 3-second response is fine in text but ruinous in face-to-face conversation. Latency is a quality variable only in embodied settings.

2.4 Point 3 — Real Deployment Targets Are Embodied

Universities, hospitals, hotels, and retail spaces increasingly deploy Pepper-class robots as reception, wayfinding, and triage agents. The question those organisations ask is not "which model has the highest MMLU score" — it is "which model produces the best **face-to-face experience** on **our** robot for **our** visitors talking about **our** information?"

OmniLLM's experiment design answers exactly that question. The five experimental conditions (A–E, see Chapter 37) directly map to the choices a deployment team must make: cloud vs local, single model vs ensemble, RAG vs no RAG.

2.5 Point 4 — The Methodological Bridge

(Academic) OmniLLM bridges two communities that rarely talk. The **NLP / LLM** community publishes leaderboards (LMSYS, LiveBench, HELM). The **HRI** community publishes user studies with validated psychological instruments (Godspeed, NARS, RoSAS). OmniLLM applies the *NLP* benchmarking infrastructure inside the *HRI* experimental framework. The Embodied LLM Arena is, in effect, **Chatbot Arena rendered through a robot**.

A useful side-effect: any HRI lab can adopt the OmniLLM platform, swap in their own robot bridge (the abstract RobotBridge makes this a 200-line job) and their own knowledge base, and replicate the methodology. The platform is itself a research contribution.

2.6 When You *Can* Skip the Robot

Be honest with yourself about scope. The robot is essential when:

- You are evaluating **embodied user experience** (the central thesis claim).
- You need **multimodal coherence** (gesture + speech + LED).

- You want **latency-aware** ratings (real-time feel).
- You are running **the experiment** with real participants.

The robot is **not** essential when:

- You are debugging the AI server or RAG pipeline.
- You are evaluating model quality on text benchmarks.
- You are demoing the routing logic.
- You are training a new component locally.

In other words: build the AI server first, debug it without Pepper, and only involve the robot when the question being asked actually requires embodiment. This sequence is reflected in the build order of Chapter 33.

Chapter 3 — The Five-Minute Tour of the Repository

***** AT A GLANCE.**** The repository has one Python package (omnilm/) split into seven sub-packages, one config directory, one knowledge-base directory, one tests directory, and a small constellation of Markdown docs. This chapter is a guided walk through the file tree so you know where to look later.

3.1 The Top-Level Layout

When you `cd` into the project root you see this:

```
OmnILLM/
+-- README.md                # Project front page
+-- GETTING_STARTED.md       # Beginner setup guide
+-- EXPLANATION.md           # Detailed file-by-file explanation
+-- ARCHITECTURE.md          # Architecture reference (large)
+-- PEPPER_CHOREGRAPHE_GUIDE.md # Pepper-specific walkthrough
+-- OmnilLM_Complete_Beginners_Guide.md # Choregraphe + virtual robot path
+-- LICENSE                  # MIT licence
+-- pyproject.toml           # Python package metadata + extras
+-- requirements.txt          # Flat dependency list
+-- .env.example              # Template for API keys
+-- .gitignore                # Files git should ignore
|
+-- config/                  # YAML configuration
|   +-- models.yaml          # ★ The model registry (every LLM lives here)
|   +-- tasks/                # Benchmark task YAML files
|       +-- reasoning.yaml
|       +-- knowledge.yaml
|       +-- code.yaml
|       +-- instruction.yaml
|       +-- safety.yaml
|       +-- robot.yaml
|
+-- omnilm/                  # ★ Main Python package
|   +-- __init__.py           # Public API re-exports
|   +-- gateway.py            # LLMGateway – unified LiteLLM wrapper
|   +-- router.py              # SmartRouter – strategy-based selection
|   +-- consensus.py          # ConsensusEngine – multi-model council
|   +-- evaluator.py           # Evaluator – LLM-as-Judge patterns
|   +-- scorer.py              # EloScorer – chess-style leaderboard
|   +-- cli.py                 # Click+Rich terminal dashboard
|   |
|   +-- hri/                  # Human-Robot Interaction module
|       +-- __init__.py
|       +-- classifier.py     # T1-T4 task classifier
|       +-- language_detector.py # Multilingual routing
|       +-- experiment.py     # Conditions A-E, sessions
|       +-- agent_graph.py    # ★ LangGraph multi-node pipeline
|
|   +-- rag/                  # Retrieval-Augmented Generation
|       +-- __init__.py
|       +-- pipeline.py       # ChromaDB + keyword fallback
|
|   +-- robotics/             # Robot bridges + helpers
|       +-- __init__.py
|       +-- bridge.py         # Abstract RobotBridge + RobotAction
|       +-- pepper.py         # Pepper HTTP bridge
|       +-- nao.py            # NAO HTTP bridge
|       +-- buddy.py          # Buddy WebSocket bridge
|       +-- gesture_planner.py # Task → gesture mapping
|       +-- whisper_stt.py    # Whisper STT wrapper
|
|   +-- server/               # AI server + NAOqi client
|       +-- __init__.py
```

```

| | +-- app.py # ★ Flask AI server (Python 3.x)
| | +-- naoqi_client.py # ★ NAOqi client (Python 2.7)
| |
| +-- tasks/ # Built-in benchmark tasks
| | +-- __init__.py
| | +-- loader.py # YAML/JSON task loader
| | +-- sample_tasks.py # 8 axes of built-in eval tasks
| |
| +-- utils/ # Shared utilities
| | +-- __init__.py
| | +-- cost_tracker.py # USD spend per model/session
| | +-- experiment_logger.py # InteractionRecord logging
| | +-- questionnaire.py # Likert + Godspeed + pairwise
| | +-- export.py # CSV / JSON / Markdown exporters
| |
+-- knowledge_base/ # ★ RAG document store
| +-- lab_info.txt
| +-- faq.txt
| +-- visitor_profiles.csv
| +-- event_schedule.csv
| +-- research_projects.txt
| +-- university_map.txt
|
+-- tests/ # 278 pytest test cases
| +-- test_gateway.py
| +-- test_router.py
| +-- test_evaluator.py
| +-- test_scorer.py
| +-- test_consensus.py
| +-- test_rag.py
| +-- test_hri.py
| +-- test_gesture_planner.py
| +-- test_experiment_logger.py
| +-- test_new_components.py
|
+-- results/ # Auto-created – eval JSON, CSV exports

```

The stars (★) mark the **most important files**. If you are reading unfamiliar source code, start with those seven and the rest follows.

3.2 The Five Files To Know By Heart

If you only have time to remember five files, make them these:

#	File	What it does	Reason
1	config/models.yaml	Registry of every model	Adding a new LLM is 7 YAML lines, zero code
2	omnillm/gateway.py	One function call to any provider	Everything else depends on it
3	omnillm/hri/agent_graph.py	The robot's brain, as a graph	Where audio becomes a robot action
4	omnillm/server/app.py	Flask HTTP API for Pepper	The contact surface with the robot

5	omnillm/server/naoqi_client.py	Python 2.7 robot side	The only Python 2 file in the project
---	--------------------------------	-----------------------	---------------------------------------

Files 4 and 5 talk to each other over HTTP. Files 1, 2, and 3 talk to each other through Python imports. The whole project is, in essence, those five files plus their helpers.

3.3 How The Documentation Files Relate

You may have noticed there are six(!) Markdown files at the root. They serve different audiences and were written at different times:

File	Audience	What it covers	Status
README.md	First-time visitor	Project overview, install, key features	Current
GETTING_STARTED.md	Beginner setting up	Step-by-step install + first commands	Current
EXPLANATION.md	Deep-dive reader	File-by-file explanation, every command	Current
ARCHITECTURE.md	Architect	Diagrams, data flow, endpoint maps	Current
PEPPER_CHOREGRAPHE_GUIDE.md	Robot-stage user	Pepper + Choregraphe walk-through	Current
OmniLLM_Complete_Beginners_Guide.md	Local-only user	Choregraphe virtual robot, no API keys	Current

This book is a *consolidation* of everything in those documents plus extensive new material derived directly from the source code. When the book and a doc disagree, the **source code** is the truth (the docs may have aged).

3.4 What Lives Outside the Repository

Some things you'll need are **not** in the repo and never will be:

- **API keys** for OpenAI, Anthropic, Google, DeepSeek, Qwen — they live in your `.env` file (which is in `.gitignore`).
- **Ollama models** (Llama 3, Qwen, Mistral, etc.) — stored by Ollama in its own directory, often `~/.ollama/`.
- **The Pepper robot** itself, obviously. But also the **NAOqi SDK** which is downloaded from SoftBank's developer portal (you have it at `C:\pynaoqi\`).
- **Choregraphe** — installed as a desktop application at `C:\Program Files (x86)\Aldebaran\Choregraphe Suite 2.5\`.
- **Python 2.7** — needed only for the NAOqi client; lives at `C:\Python27\python.exe`.

Chapter 4 — The Vocabulary You Need Before Chapter 5

**** AT A GLANCE.** This is the project's Rosetta Stone. Every term you'll meet later in the book is defined here once. Skim it; come back when a word trips you up. Each entry has a one-line definition and, where useful, a "first appears in" pointer.*

(Book) Read this chapter as a glossary, not a narrative. It will make later chapters dramatically easier.

4.1 AI / LLM Vocabulary

LLM — Large Language Model. An AI system that consumes and produces text. Examples: GPT-4o, Claude Sonnet, Gemini Flash, Llama 3.

Provider — A company or service that hosts an LLM. OpenAI, Anthropic, Google, DeepSeek, Alibaba (Qwen), and Ollama (free local) are the providers OmniLLM uses.

Model ID (in OmniLLM) — A short string used inside the project to refer to a model. Defined in `config/models.yaml`. Example: `openai-gpt4o-mini`.

Token — Roughly one short word, or a sub-word fragment. LLMs charge in tokens. "Hello world" is two tokens; "antidisestablishmentarianism" might be five.

Prompt — The input you send to an LLM. In OmniLLM, prompts are *messages* in OpenAI chat format: `[{"role": "user", "content": "Hello"}]`.

System prompt — A leading message with `role: "system"` that sets the model's behaviour. OmniLLM uses one to make the LLM "be Pepper".

Temperature — A 0.0-to-1.0 number controlling how random the model is. 0 = deterministic, 1 = creative. OmniLLM defaults to 0.7 for chat, 0.0 for judge calls.

API key — A secret string proving you have an account with a provider. Lives in `.env`. Never commit to git.

LiteLLM — A Python library that gives all providers a unified API. OmniLLM uses it as the foundation of `gateway.py`.

Ollama — A free, local LLM runtime. Runs models like Llama 3 on your own GPU or CPU. Has no API key. Listens on `http://localhost:11434`.

Whisper — OpenAI's open-source speech-to-text model. Comes in sizes `tiny`, `base`, `small`, `medium`, `large`. OmniLLM can use Whisper *locally* (free, no API call) or via the OpenAI Whisper API.

4.2 RAG and Evaluation Vocabulary

RAG — Retrieval-Augmented Generation. Search a document store for relevant text chunks, then prepend them to the LLM prompt. Reduces hallucination dramatically.

Vector store / Vector database — A specialised database that stores *embeddings* (lists of numbers representing text meaning) and supports "find the K most similar items" queries. OmniLLM uses **ChromaDB** as its vector store.

Embedding — A list of (typically 384–1536) floating-point numbers that encodes the semantic meaning of a piece of text. Two pieces of text with similar meaning have similar embeddings.

Chunk — A short slice of a longer document, used as the unit of retrieval. OmniLLM chunks documents into 512-character pieces with 64-character overlap.

Faithfulness — A 0–1 score measuring whether an LLM's answer actually uses the retrieved context (vs making things up). OmniLLM computes this via LLM-as-Judge.

Hallucination — When the LLM confidently states something not in the retrieved context. OmniLLM uses a word-overlap heuristic to flag candidates.

LLM-as-Judge — Using one LLM to evaluate another LLM's response. OmniLLM implements three variants: Referenceless (G-Eval), Reference-Based, and Pairwise.

ELO — A rating system from chess. After many head-to-head matches, models get a numeric rating where $+100 \approx 64\%$ expected win-rate. Used by LMSYS Chatbot Arena. OmniLLM has its own `EloScorer`.

Pairwise comparison — A judge sees two responses and picks the better one. OmniLLM runs each comparison **twice** with positions swapped, to cancel position bias.

Position bias — The well-documented tendency of LLM judges to prefer the *first* response they see. Mitigated by swap-and-aggregate.

4.3 Robotics / HRI Vocabulary

HRI — Human-Robot Interaction. The academic field studying how people interact with robots.

Pepper — A 120 cm humanoid social robot from SoftBank Robotics (originally Aldebaran). 20 degrees of freedom, eye LEDs, chest tablet, four microphones, two speakers. Runs Python 2.7 NAOqi.

NAO — Pepper's smaller sibling (58 cm, 25 DoF). Same NAOqi OS. Common in education and RoboCup.

Buddy — A different companion robot from Blue Frog Robotics. Android-based, no NAOqi. OmniLLM has a separate WebSocket bridge for it.

NAOqi — Pepper / NAO's middleware operating system. Provides services like `ALAnimatedSpeech`, `ALMotion`, `ALLeds`. **Locked to Python 2.7**, which is the source of the Python-version bridge problem.

ALAnimatedSpeech — The NAOqi service that makes Pepper speak with synchronised body animations. Cleaner than `ALTextToSpeech` (which is voice-only).

ALMotion — The NAOqi service controlling Pepper's joints (motors). You must call `motion.wakeUp()` before the robot can move.

ALLeds — The NAOqi service controlling the eye and ear LEDs. `fadeRGB("FaceLeds", r, g, b, duration)` changes eye colour.

ALBehaviorManager — The NAOqi service that runs pre-built animations ("behaviors") like waving, bowing, pointing.

Behavior — A named animation file installed on the robot. Lives at paths like `animations/Stand/Gestures/Hey_1` (which is the wave gesture).

Choregraphe — A drag-and-drop desktop IDE from SoftBank for programming Pepper / NAO. Includes a virtual robot simulator. **Only works with NAOqi 2.5** (the version Pepper uses).

Pepper SDK / pynaoqi — The Python 2.7 binding to NAOqi services. Imported as `import qi` (modern) or `import naoqi` (legacy).

RobotAction — In OmniLLM, a dataclass with fields `speech`, `gesture`, `movement`, `emotion_led`, `nav2_goal`, `metadata`. The contract between the brain (Python 3) and the body (Python 2.7).

Embodiment effect — The HRI finding that the *same* AI text is rated differently when delivered through a physical robot versus a screen.

Godspeed Questionnaire Series — Bartneck et al.'s validated five-subscale instrument for measuring perceived anthropomorphism, animacy, likeability, intelligence, and safety of robots. OmniLLM has a dataclass for it.

Likert scale — A 1–7 (or 1–5) rating of agreement. OmniLLM uses 1–7 for custom questionnaire items, 1–5 for Godspeed.

Within-subjects design — An experimental design where every participant experiences every condition. OmniLLM uses this for the 5 conditions A–E.

Counterbalancing / Latin square — A method of varying the *order* in which participants meet the conditions, to cancel out fatigue and order-of-presentation effects.

4.4 Software Engineering Vocabulary

Async / await / coroutine — Python's syntax for non-blocking I/O. `async def` declares a coroutine; `await` pauses it until the awaited operation completes; `asyncio.gather()` runs many coroutines concurrently. OmniLLM uses `async` heavily because LLM calls are I/O-bound.

Dataclass — A Python class created with `@dataclass` decorator. Auto-generates `__init__`, `__repr__`, etc. OmniLLM uses dataclasses everywhere for structured data.

LangGraph — A library on top of LangChain for building multi-node agent graphs (DAGs of functions sharing state). Used in `hri/agent_graph.py`.

LangChain — A broader framework for LLM apps. OmniLLM uses LangChain's document loaders for PDF and CSV ingestion in the RAG pipeline.

ChromaDB — A pure-Python embedded vector database. Used in `rag/pipeline.py` when available; the pipeline falls back to keyword search if not installed.

Flask — A minimalist Python web framework. Used in `server/app.py` to expose the AI server's HTTP API.

Click — A Python library for building command-line interfaces. Used in `cli.py` to build the `omni_llm` command.

Rich — A Python library for beautiful coloured terminal output. Used in `cli.py` for tables, panels, and progress bars.

Pyproject.toml — The modern Python project metadata file. Defines package name, version, dependencies, optional extras (`[dev]`, `[hri]`, `[robotics]`, `[all]`).

Virtual environment / venv — An isolated Python sandbox per project. OmniLLM lives in `.venv/` after `python -m venv .venv`.

Editable install — `pip install -e .` makes a package importable without copying files; changes to source code take effect immediately. OmniLLM is always installed editable for development.

HTTP / REST — The web protocol used between OmniLLM's Flask server and Pepper's NAOqi client. All payloads are JSON.

WebSocket — A persistent bi-directional protocol used between OmniLLM and the Buddy robot for token-by-token streaming.

JSON — JavaScript Object Notation. OmniLLM uses it as the universal serialisation format for inter-process communication.

WAV / PCM — The audio format used for voice transport. 16 kHz mono, 16-bit signed PCM, optionally base64-encoded for HTTP transport.

Base64 — A way to encode binary bytes as ASCII text so they fit in JSON.

PART II — ARCHITECTURE & DATA FLOW

Chapter 5 — The Three-Layer System Architecture

*** AT A GLANCE. *OmniLLM has three physical tiers. (1) The human + robot tier — a person speaking to Pepper. (2) The AI server tier — a Python 3.x Flask process running the LangGraph pipeline, RAG, gateway, router, council, evaluator. (3) The provider tier — external LLM APIs (paid clouds) and local Ollama (free). The robot is a thin client; all intelligence lives in tier 2; tier 3 is interchangeable.*

5.1 The High-Level Picture



5.2 Why Three Tiers (and not two, or four)?

The three-tier choice is forced by a single hard constraint: **NAOqi is locked to Python 2.7, and every modern AI library requires Python 3.11+.** You cannot run them in the same process. That single constraint creates the need for the AI server tier as a *separate process* communicating over HTTP.

Could we have a *two-tier* architecture by putting the AI directly on Pepper? No: Pepper's onboard computer (Intel Atom E3845, 4 GB RAM, 8 GB flash) cannot host LLMs, ChromaDB, LangGraph, and Whisper.

Could we have a *four*-tier architecture by adding a "perception preprocessor" between robot and server? You could, but the marginal benefit is small and the engineering cost is large. The current three-tier design has been used in 15+ published Pepper-LLM integration projects (see Chapter 26) and is the established pattern.

5.3 Cross-Tier Contracts

Each tier boundary has a **contract** — a precise data format both sides agree on. These contracts are the project's most important invariants. If you want to swap a tier (e.g., replace Pepper with NAO), you only need to honour the contracts.

Boundary	Direction	Format	What it carries
Human → Robot	analogue speech	sound waves	The user's question
Robot → AI server	HTTP POST	JSON with base64 WAV	{audio, participant_id, session_id, condition, rag_enabled}
AI server → Robot	HTTP response	JSON	{speech, gesture, emotion_led, metadata}
AI server → Provider	HTTPS POST	OpenAI chat format	{model, messages, temperature, max_tokens}
Provider → AI server	HTTPS response	Provider-specific	{choices, usage, ...} (LiteLLM normalises this)
AI server → ChromaDB	in-process	Python API	Embedding vector + chunk text

The contract that matters most for *your* code is the **robot ↔ AI server** contract because everything else either runs inside one process or uses a well-known external API. The robot ↔ AI server contract is the one OmniLLM defines and you may need to extend.

5.4 The Robot-↔-Server JSON Contract

Going **into** the server (POST `/interact`):

```
{
  "audio":          "<base64 string of 16 kHz mono WAV>",
  "text":           "(optional – used in text-only mode instead of audio)",
  "participant_id": "P001",
  "session_id":     "s-abc-123",
  "condition":      "A" | "B" | "C" | "D" | "E",
  "rag_enabled":    true | false
}
```

Coming **out** of the server (200 OK):

```
{
  "speech": "Room 305 is on the third floor, on your left.",
  "gesture": "point_left",
  "emotion_led": "#00AAFF",
  "metadata": {
    "task_type": "navigation",
    "model_id": "openai-gpt4o-mini",
    "rag_enabled": true
  }
}
```

The robot's NAOqi client (`server/naoqi_client.py`) reads this JSON, sets the LED colour via `ALLeds.faderGB`, runs the gesture as a NAOqi behaviour via `ALBehaviorManager.runBehavior(GESTURE_TO_BEHAVIOR[gesture])`, and speaks the text via `ALAnimatedSpeech.say(speech)`. Everything Pepper does is determined by these three fields plus optional metadata.

5.5 What If a Tier Fails?

A robust system degrades, it does not crash. Here is what happens at each tier failure:

Failure	Effect	Mitigation built into OmniLLM
Pepper offline	No HRI possible	Use text-only mode (<code>{"text": "..."} — chapter 29</code>
AI server down	Robot speaks nothing	NAOqi client retries; user hears "I could not connect to my AI brain"
Cloud LLM API error	Specific provider unreachable	Gateway returns <code>ModelResponse</code> with <code>error</code> field; router can fall back to a different model
Ollama down	Local models unreachable	Cloud fallback (if API keys configured)
ChromaDB import fails	No vector search	RAG pipeline auto-falls-back to keyword search
LangGraph import fails	No agent graph	Server falls back to direct gateway call (<code>_fallback_interact</code> in <code>server/app.py</code>)
Whisper not installed	No audio transcription	Text-only mode still works

This **graceful degradation** is intentional. The system has many optional components. Each failure isolates rather than cascading.

Chapter 6 — The Lifecycle of a Single Spoken Question


**** AT A GLANCE.** This chapter follows one sentence — "Where is Room 305?" — through the entire system, byte by byte, from the moment the user opens their mouth to the moment Pepper points left. Twenty-one steps. The whole journey takes about 1–3 seconds.*

6.1 The 21 Steps

#	Where	What happens	Data shape	File / line
1	Air	Human speaks	sound waves	—
2	Pepper mic	ALAudioDevice records	16 kHz mono PCM	naoqi_client.py:196
3	Pepper CPU	Bytes encoded base64	ASCII string	naoqi_client.py:228
4	Wi-Fi	HTTP POST to /interact	JSON body	naoqi_client.py:233
5	AI server	Flask receives, decodes b64	bytes	app.py:237
6	AI server	LangGraph invoked	HRIGraphState dict	app.py:260
7	LangGraph node	Whisper STT	"Where is Room 305?"	agent_graph.py:132
8	LangGraph node	Language detected	"en"	agent_graph.py:157
9	LangGraph node	Task classified	"navigation" (T2), conf 0.92	agent_graph.py:175
10	LangGraph router	Conditional branch	→ nav_rag	agent_graph.py:497
11	RAG	ChromaDB similarity search	top-4 chunks	pipeline.py:488
12	LLM	Augmented prompt sent	messages = [system, user_with_context]	pipeline.py:292
13	Gateway	Provider call (OpenAI / etc.)	OpenAI chat format	gateway.py:219
14	LLM	Generates answer	"Room 305 is on..."	— (provider)
15	Gesture planner	"on your left" → point_left	(gesture, led)	gesture_planner.py:112
16	Smart router (Cond C/D)	Per-task model swap or council	new model_id	agent_graph.py:355

17	LangGraph node	Build action plan	RobotAction dict	agent_graph.py:430
18	LangGraph node	Log interaction	InteractionRecord saved	agent_graph.py:460
19	Flask	Return 200 OK with JSON	{speech, gesture, emotion_led}	app.py:275
20	Pepper CPU	Receive JSON, dispatch	—	naoqi_client.py:269
21	Pepper hardware	LED, gesture, speech execute	physical actions	naoqi_client.py:278–286

6.2 The Same Story as a Sequence Diagram



6.3 Timing Budget

The user's perception of "responsive" caps at about 3 seconds in face-to-face conversation. Here is how OmniLLM stays inside that budget:

Step	Typical	Worst case
------	---------	------------

Whisper STT (local, base model)	200–500 ms	800 ms
Whisper STT (API, OpenAI)	300–800 ms	1500 ms
Language detection (rule-based)	5–10 ms	20 ms
Task classification (rule-based)	10–20 ms	50 ms
ChromaDB retrieval	50–200 ms	400 ms
LLM generation (cloud, mini-tier)	500–1500 ms	3000 ms
LLM generation (Ollama local)	1000–3000 ms	5000 ms
Gesture planning	5–10 ms	20 ms
Network round-trips (LAN)	10–50 ms	100 ms
Sum (typical, cloud)	~1.0 s	~3 s
Sum (typical, local)	~2.0 s	~6 s

The biggest variable is the LLM call itself. This is why **Condition D (Consensus)** is slowest — it makes three concurrent LLM calls plus a synthesis call. It pays a latency tax for higher quality.

(Tip) The latency budget is also why OmniLLM uses `asyncio.gather()` everywhere — querying three council models concurrently is the same wall-clock time as querying one. If we did them sequentially, Condition D would take 6+ seconds.

6.4 Where Time Is Spent (a real measurement)

For a typical T2 (Navigation) interaction with `condition=C`, GPT-4o-mini, RAG enabled, English input, the breakdown looks like:

```

                                Latency contribution (ms)
Network (Pepper → server)      | 35
Base64 decode + dispatch       | 5
Whisper transcription          ##### 380
Language detect                 | 5
Task classify                   | 12
ChromaDB retrieve top-4        || 90
LLM call (GPT-4o-mini)         ##### 850
Gesture planning               | 4
Build RobotAction              | 1
Logging                       | 2
JSON serialise + return        | 8
Network (server → Pepper)     | 32
=====
Total: ~1.4 s

```

The LLM call dominates. Almost everything else is essentially free.

Chapter 7 — The LangGraph Agent Pipeline, Node by Node

***** AT A GLANCE.** `omnilm/hri/agent_graph.py` **builds a** directed graph** of 9 nodes that share a state dictionary. Audio comes in at the top; a `RobotAction` dict comes out at the bottom. Each node is an `async` function that reads state, does one thing, and returns state updates. `LangGraph` handles the wiring.

7.1 Why a Graph?

You could write all this logic as a single 200-line `async` function with `if/elif` branches. The result would work but be:

- Hard to test (no node-level boundaries)
- Hard to read (every concern interleaved)
- Hard to extend (where do you put a new step?)
- Hard to visualise (no graph to render)


`LangGraph` gives you **named nodes**, **typed state**, and **conditional edges**. The graph is, literally, a programmable state machine. It produces a graph diagram you can show your supervisor; it lets you mock individual nodes in tests; it lets you add a new branch without touching the others.

7.2 The Nodes

```
graph TD
  START([START]) --> A[A[transcribe_audio<br/>Whisper STT]]
  A --> B[B[detect_language<br/>script + n-gram + langdetect]]
  B --> C[C[classify_task<br/>T1/T2/T3/T4 keyword scoring]]
  C -->|info_retrieval +<br/>rag_enabled| D[D[rag<br/>ChromaDB + LLM]]
  C -->|navigation +<br/>rag_enabled| E[E[nav_rag<br/>RAG + gesture planner]]
  C -->|social_conversation| F[F[direct_llm<br/>system prompt = Pepper]]
  C -->|multilingual| G[G[multilingual_llm<br/>language-optimal model]]
  C -->|condition=E<br/>rag-off| F
  D --> H[H[smart_router<br/>Cond C/D = override]]
  E --> H
  F --> H
  G --> H
  H --> I[I[generate_action_plan<br/>build RobotAction dict]]
  I --> J[J[log_interaction<br/>ExperimentLogger]]
  J --> END([END])
```

(If your PDF renderer does not support Mermaid, here is the same as ASCII:)





7.3 The State Object

Every node sees and updates one shared object: `HRIGraphState`. Here is the full dataclass, with a comment for each field:

```

@dataclass
class HRIGraphState:
    # -- Input fields (set by the caller) -----
    audio_bytes: bytes = b"" # raw WAV from Pepper (or empty)
    participant_id: str = "" # e.g. "P001"
    session_id: str = "" # UUID for this session
    condition: str = "A" # experimental condition A-E
    rag_enabled: bool = True # may be overridden by condition E

    # -- Intermediate fields (filled progressively by nodes) -----
    utterance: str = "" # transcribed text
    detected_language: str = "en" # ISO 639-1
    task_type: str = "info_retrieval" # T1/T2/T3/T4
    task_confidence: float = 0.0 # 0-1

    rag_context: str = "" # formatted context block
    rag_chunks: list[Any] = [] # DocumentChunk list
    rag_faithfulness: float = -1.0 # -1 = not scored

    model_id: str = "" # model that ultimately answered
    response_text: str = "" # the answer Pepper will speak
    gesture: str = None # e.g. "point_left"
    led_color: str = None # e.g. "#00AAFF"

    robot_action: dict = {} # the final dict returned to Pepper

    judge_score: float = -1.0
    latency_ms: float = 0.0
    input_tokens: int = 0
    output_tokens: int = 0
    cost_usd: float = 0.0

    error: str | None = None
    _start_time: float = field(default_factory=time.monotonic)

```

The state grows as the graph runs. By the time it reaches `END`, every relevant field is filled.

7.4 Each Node, Briefly

We will spend a whole chapter (Chapter 17) on this file, but here is the node-by-node summary so you can read the rest of the architecture:

Node 1 — transcribe_audio. If `audio_bytes` is non-empty, instantiate `WhisperSTT`, call `await stt.transcribe(audio_bytes)`, store the text in `utterance`. If `audio_bytes` is empty (text-only mode), skip — `utterance` was already set by the caller.

Node 2 — detect_language. Run `LanguageDetector().detect(utterance)`. Store ISO 639-1 code in `detected_language`.

Node 3 — classify_task. Run `HRI_TaskClassifier().classify(utterance, detected_language)`. Returns a `ClassificationResult` whose `.task_type` and `.confidence` are stored.

Conditional edge after Node 3. A function `_route_by_task_type` examines `task_type`, `condition`, and `rag_enabled` and returns the next node name:

- Condition E (RAG-off control) always → `direct_llm`.
- T1 (info_retrieval) + RAG → `rag`.
- T2 (navigation) + RAG → `nav_rag`.
- T3 (social_conversation) → `direct_llm`.
- T4 (multilingual) → `multilingual_llm`.
- Anything else → `direct_llm`.

Node 4a — rag. Run the RAG pipeline. Produces `rag_context`, `rag_chunks`, `rag_faithfulness`, `response_text`, `model_id`, `latency_ms`.

Node 4b — nav_rag. Same as `rag`, plus calls the `GesturePlanner` with `task_type="navigation"` to set `gesture` and `led_color` based on direction words in the response.

Node 4c — direct_llm. Build a Pepper system prompt ("You are Pepper, a friendly social robot..."). Call `gateway.query(model_id, messages, temperature=0.8)`. Run `GesturePlanner` for `social_conversation`.

Node 4d — multilingual_llm. Re-detect language (in case the LLM-based detector found something the rule-based missed). Build a prompt instructing the LLM to respond in the detected language. Call the gateway with the language-optimal model.

All four 4-nodes then converge on:

Node 5 — smart_router. Inspects `condition`. For Condition A/B/E this is a no-op (the model has already answered). For Condition C, override the choice using `SmartRouter.route_for_hri_task(task_type)` and re-call. For Condition D, query a 3-model council via `ConsensusEngine`, synthesise.

Node 6 — generate_action_plan. Combine `response_text`, `gesture`, `led_color`, and metadata into the `robot_action` dict that Pepper will receive.

Node 7 — log_interaction. If an `ExperimentLogger` was provided when the graph was built, write an `InteractionRecord` with everything we know about this turn: model, latency, tokens, cost, RAG metrics, condition, task type, language, gesture used.

The graph then terminates. The Flask route returns `result["robot_action"]` as JSON to Pepper.

7.5 Why "Conditional Edges" Not "If Statements"

In LangGraph, you do not write `if task_type == "navigation": ...` inside a node. You write a pure function that returns a string node name, and you register conditional edges:

```
builder.add_conditional_edges(  
    "classify_task",          # source node  
    _route_by_task_type,     # the function returning a node name  
    {  
        "rag": "rag",  
        "nav_rag": "nav_rag",  
        "direct_llm": "direct_llm",  
        "multilingual_llm": "multilingual_llm",  
    },  
)
```

This separation matters because:

1. The router can be tested in isolation (it's a pure function).
2. The graph is statically inspectable — LangGraph can produce a DOT-format diagram of all possible paths.
3. Adding a new branch is a one-line change to the routing function plus a new node registration, without touching existing nodes.

Chapter 8 — Where Every Byte Goes — Endpoint and Provider Map

*** AT A GLANCE.** This is the reference page for "what URL does what". Every HTTP endpoint in OmniLLM, every external API the project calls, every protocol used. Bookmark this for debugging.

8.1 The AI Server's HTTP Endpoints

These are the URLs `server/app.py` exposes. By default the server listens on `http://0.0.0.0:5000`.

Method	Path	What it does	Request body	Response body
GET	/health	Liveness probe	(none)	<code>{"status": "ok", "version": "0.1.0"}</code>
GET	/status	Server config dump	(none)	model list, RAG flag, KB path, langgraph flag
POST	/interact	Main HRI endpoint	<code>{audio?, text?, participant_id, session_id, condition, rag_enabled}</code>	<code>{speech, gesture, emotion_led, metadata}</code>
POST	/transcribe	STT only	<code>{audio: <base64>}</code>	<code>{text, language}</code>
POST	/evaluate	Submit questionnaire	<code>{session_id, participant_id, condition, scores}</code>	<code>{}</code>
GET	/export	Dump session logs	(none)	<code>{count, records: [...]}</code>

8.2 External Endpoints OmniLLM Calls

These are URLs OmniLLM calls **outwards** to LLM providers. Note that LiteLLM hides the differences; the gateway just builds a model string.

Provider	Endpoint	Auth	Cost basis
----------	----------	------	------------

OpenAI	https://api.openai.com/Authorization: tions Bearer \$OPENAI_API_KEY	per-1M tokens
Anthropic	https://api.anthropic.cx-api-key: ages \$ANTHROPIC_API_KEY	per-1M tokens
Google Gemini	https://generativelanguage.googleapis.com/v1/models/gemini-1.5-pro-latest:generateContent API key in URL or header	per-1M tokens
DeepSeek	https://api.deepseek.coAuthorization: letions Bearer \$DEEPSEEK_API_KEY	per-1M tokens
Qwen / DashScope	https://dashscope.aliyuncs.com/compatible-models/qwen1.5-72b-chat API key	per-1M chat completions
Ollama (local)	http://localhost:11434/(no auth)/completions	\$0
OpenAI Whisper API	https://api.openai.com/Authorization: criptions Bearer \$OPENAI_API_KEY	\$0.006/min

8.3 The Pricing Cheat-Sheet

For the models OmniLLM has registered, here is the cost per 1 million tokens (values from config/models.yaml, current as of April 2026):

OmniLLM ID	Provider	Input \$/1M	Output \$/1M	Type
openai-gpt54	OpenAI GPT-5.4	2.50	15.00	cloud
openai-gpt4o	OpenAI GPT-4o	2.50	10.00	cloud
openai-o1	OpenAI o1	15.00	60.00	cloud
openai-o3-mini	OpenAI o3-mini	1.10	4.40	cloud
openai-gpt4o-mini	OpenAI GPT-4o-mini	0.15	0.60	cloud
claude-sonnet	Claude Sonnet 4.6	3.00	15.00	cloud
claude-opus	Claude Opus 4.6	15.00	75.00	cloud
claude-haiku	Claude Haiku 4.5	0.80	4.00	cloud
gemini-2.5-pro	Gemini 2.5 Pro	1.25	10.00	cloud

gemini-2.5-flash	Gemini 2.5 Flash	0.15	0.60	cloud
gemini-flash	Gemini 2.0 Flash	0.10	0.40	cloud
deepseek-v3-cloud	DeepSeek V3	0.27	1.10	cloud
deepseek-r1-cloud	DeepSeek R1	0.55	2.19	cloud
qwen-2.5-72b	Qwen 2.5 72B	1.10	3.40	cloud
llama3-8b-local	Llama 3 8B (Ollama)	0.00	0.00	local
qwen-2.5-7b-local	Qwen 2.5 7B (Ollama)	0.00	0.00	local
qwen-2.5-local	Qwen 2.5 3B (Ollama)	0.00	0.00	local
llama3.2-local	Llama 3.2 3B (Ollama)	0.00	0.00	local

A typical 500-token interaction (250 in / 250 out) costs:

- GPT-4o-mini: about **\$0.0002** (one-fifth of one US cent).
- Claude Haiku: about **\$0.0012**.
- Gemini Flash 2.0: about **\$0.000125** (the cheapest paid option).
- Llama 3 8B (Ollama): **\$0**.
- GPT-5.4: about **\$0.0044** — twenty times more than mini.

For a 240-interaction experimental study (48 participants × 5 conditions × 4 tasks), the total spend on `openai-gpt4o-mini` is around 5 US cents.

8.4 Other Inter-Process Channels

OmniLLM uses a few channels that are not HTTP:

Channel	Used for	Where defined
Python in-process call	Gateway → LiteLLM	<code>gateway.py:219</code>
ChromaDB Python API	RAG → vector store	<code>rag/pipeline.py:478</code>
Whisper Python API	STT	<code>robotics/whisper_stt.py:152</code>
WebSocket	AI server → Buddy robot	<code>robotics/buddy.py:78</code>
Filesystem JSON	EloScorer save/load	<code>scorer.py:237</code>

Filesystem JSON/CSV	ExperimentLogger save	utils/experiment_logger.py:359
stdout (CLI)	User-facing tables & panels	cli.py everywhere

8.5 A Test You Can Run Right Now

If you want to confirm the network architecture is wired correctly, here is a five-line test:

```
# 1. Start the AI server in a terminal
python -m omnillm.server.app

# 2. In a second terminal, ping it
curl http://localhost:5000/health

# 3. Send a text-only "interaction"
curl -X POST http://localhost:5000/interact \
  -H "Content-Type: application/json" \
  -d '{"text": "Where is Room 305?", "participant_id": "TEST", "condition": "B"}'
```

You should see a JSON response with speech, gesture, and emotion_led. That confirms tiers 2 and 3 are healthy. Tier 1 (the robot) is optional.

PART III — MODULE DEEP-DIVE

This is the longest part of the book — one chapter per source-code module. The chapters follow a uniform template:

1. **At-a-glance** — what the file is and why it exists.
2. **File facts** — line count, dependencies, public API surface.
3. **Walk-through** — beginner-friendly explanation of the code.
4. **Real code** — annotated excerpts from the actual source file.
5. **Worked example** — usage you can run.
6. **Academic context** — research and design rationale (when relevant).

Chapter 9 — gateway.py — The Door to Every LLM

***** AT A GLANCE. `omnilm/gateway.py` is the most important file in the project**. It exposes one async method, `query(model_id, messages)`, which can talk to any registered LLM (cloud or local) by reading `config/models.yaml` and using LiteLLM as the universal adapter. Returns a structured `ModelResponse` with content, tokens, latency, cost, and any error. `query_multiple()` runs many models concurrently via `asyncio.gather`.**

9.1 File Facts

Attribute	Value
Path	<code>omnilm/gateway.py</code>
Lines	271
External deps	<code>litellm</code> , <code>pyyaml</code>
Public classes	<code>LLMGateway</code> , <code>ModelResponse</code>
Imports <code>models.yaml</code> from	<code><repo_root>/config/models.yaml</code>

9.2 The 30-Second Walk-Through

`LLMGateway` does five things:

1. **Loads** `config/models.yaml` on `__init__`.
2. **Lists** registered models (`list_models`, `list_cloud_models`, `list_local_models`).
3. **Builds** the LiteLLM model string for each provider (`ollama/llama3:8b` for local, just `gpt-4o` for OpenAI, `gemini/<model>` for Google, etc.).
4. **Calls** `litellm.acompletion()` with API key from environment.
5. **Returns** a `ModelResponse` dataclass — never raises; errors are captured in the `error` field.

That's it. Everything else in OmniLLM depends on this contract.

9.3 The `ModelResponse` Contract

This dataclass is the universal return type. Every higher layer (router, council, evaluator, RAG) reads exactly these fields:

```
@dataclass
class ModelResponse:
    model_id: str
    content: str
    input_tokens: int = 0
    output_tokens: int = 0
    latency_ms: float = 0.0
    cost_usd: float = 0.0
    time_to_first_token_ms: float = 0.0
    error: str | None = None
```

```

metadata: dict[str, Any] = field(default_factory=dict)

@property
def total_tokens(self) -> int:
    return self.input_tokens + self.output_tokens

@property
def is_error(self) -> bool:
    return self.error is not None

```

(Book) **Why a dataclass and not a dict?** With a dataclass, your IDE auto-completes field names, your type-checker catches typos, and the structure is documented in one place. A dict would work but every consumer would have to remember the keys.

9.4 How the Model String Is Built

This is the only provider-specific code in the whole gateway:

```

def _build_model_string(self, model_id: str) -> str:
    cfg = self._models[model_id]
    provider = cfg.get("provider", "openai")
    model = cfg["model"]

    if provider == "ollama": return f"ollama/{model}"
    if provider == "openai_compatible": return f"openai/{model}"
    if provider == "deepseek": return f"openai/{model}"
    if provider == "google": return f"gemini/{model}"
    if provider == "anthropic": return f"anthropic/{model}"
    return model # default: OpenAI uses bare model names

```

LiteLLM uses the prefix to dispatch to the correct provider library. DeepSeek and Qwen are OpenAI-compatible, so they use the openai/ prefix plus a custom api_base URL. Ollama gets its own prefix. Google's API uses gemini/. Anthropic uses its own SDK under anthropic/.

9.5 The Heart of the File — query()

Annotated, the core method looks like this:

```

async def query(
    self,
    model_id: str,
    messages: list[dict[str, str]],
    temperature: float = 0.7,
    max_tokens: int = 2048,
) -> ModelResponse:
    # 1. Validate the model is registered.
    if model_id not in self._models:
        return ModelResponse(model_id=model_id, content="",
                               error=f"Model '{model_id}' not found in registry.")

    cfg = self._models[model_id]

    # 2. Build the LiteLLM model string, e.g. "ollama/llama3:8b".
    kwargs = {
        "model": self._build_model_string(model_id),
        "messages": messages,
        "temperature": temperature,
        "max_tokens": max_tokens,
    }

    # 3. Per-model API base / key from YAML and environment.
    if cfg.get("api_base"):

```

```

        kwargs["api_base"] = cfg["api_base"]
    if cfg.get("api_key_env"):
        api_key = os.environ.get(cfg["api_key_env"], "")
        if api_key:
            kwargs["api_key"] = api_key

    # 4. Time the call so we can record latency_ms.
    start = time.perf_counter()
    try:
        resp = await litellm.acompletion(**kwargs)
        latency_ms = (time.perf_counter() - start) * 1000

        content = resp.choices[0].message.content or ""
        usage = getattr(resp, "usage", None)
        input_tokens = getattr(usage, "prompt_tokens", 0) or 0
        output_tokens = getattr(usage, "completion_tokens", 0) or 0
        cost = self._calculate_cost(model_id, input_tokens, output_tokens)

        return ModelResponse(
            model_id=model_id, content=content,
            input_tokens=input_tokens, output_tokens=output_tokens,
            latency_ms=latency_ms, cost_usd=cost,
        )
    except Exception as exc:
        # 5. Never raise - return a structured error response.
        latency_ms = (time.perf_counter() - start) * 1000
        return ModelResponse(model_id=model_id, content="",
                               latency_ms=latency_ms, error=str(exc))

```

9.6 Concurrent Querying

This is one of the small details that makes OmniLLM fast:

```

async def query_multiple(self, model_ids, messages, temperature=0.7,
                        max_tokens=2048):
    tasks = [self.query(mid, messages, temperature, max_tokens)
              for mid in model_ids]
    return list(await asyncio.gather(*tasks))

```

`asyncio.gather` dispatches all the queries **simultaneously**. Querying 5 models takes the wall-clock time of the slowest one, not the sum. This is critical for the LLM Council (Condition D) and for the multi-model evaluation runs.

9.7 The `litellm.drop_params = True` Footnote

At the top of the file you will find this line:

```
litellm.drop_params = True
```

This is **important** and easy to miss. Some new models (OpenAI's o1, o3, gpt-5 reasoning models) **do not accept** the `temperature` parameter. By default LiteLLM raises `UnsupportedParamsError`. With `drop_params=True`, LiteLLM silently strips unsupported parameters and continues. This single line makes OmniLLM future-proof against model API changes.

9.8 Worked Example

A complete, runnable script:

```

import asyncio
from omnillm.gateway import LLMGateway

async def main():
    gw = LLMGateway()

    # Single model
    resp = await gw.query("openai-gpt4o-mini",
                          [{"role": "user", "content": "Hi!"}])
    print(resp.content)
    print(f"Cost: ${resp.cost_usd:.6f}")

    # Many models in parallel
    responses = await gw.query_multiple(
        ["openai-gpt4o-mini", "claude-haiku", "gemini-flash"],
        [{"role": "user", "content": "What is the speed of light?"}]
    )
    for r in responses:
        print(f"{r.model_id} ({r.latency_ms:.0f} ms): {r.content[:80]}")

asyncio.run(main())

```

9.9 Adding a New Model — Zero Code

This is the payoff of YAML-driven design. To add a new model you append to `config/models.yaml`:

```

my-new-model:
  id: my-new-model
  provider: openai           # or anthropic | google | ollama | deepseek
  model: my-actual-model-name
  api_key_env: MY_API_KEY
  cost_per_1m_input: 1.00
  cost_per_1m_output: 3.00
  type: cloud
  description: "What this model is for"
  hri_strengths: ["info_retrieval"] # optional, for HRI routing

```

...and the model is now usable across the entire system: gateway, router, council, evaluator, CLI, RAG, agent graph. **No Python changes.** This is why YAML lives in `config/`, not as Python constants.

Chapter 10 — router.py — The Smart Model Selector

***** AT A GLANCE.** *** omnillm/router.py answers the question "given a task and constraints, which model should I use?" Six strategies: BEST_QUALITY, LOWEST_COST, LOWEST_LATENCY, BEST_VALUE (composite), LOCAL_PREFERRED, and TASK_TYPE (HRI-specific). Hard constraints (budget_usd, max_latency_ms) are filters; the strategy then picks the best of what survives.*

10.1 File Facts

Attribute	Value
Path	omnillm/router.py
Lines	372
Public classes	SmartRouter, RouteDecision, RoutingStrategy
Reads	config/models.yaml, optional results/*.json

10.2 Strategy Cheat-Sheet

Strategy	Picks	Used for
BEST_QUALITY	model with highest quality score	research, complex problems
LOWEST_COST	cheapest model meeting quality threshold	bulk, demos
LOWEST_LATENCY	historically fastest model	real-time HRI
BEST_VALUE	composite (quality 50% + cost 30% + latency 20%)	general
LOCAL_PREFERRED	Ollama first, cloud fallback	privacy / offline / free
TASK_TYPE	HRI: maps T1/T2/T3/T4 to specific models	Embodied LLM Arena

10.3 The Composite Value Formula

BEST_VALUE is the default strategy. Its composite score is:

```
def _calculate_value_score(self, quality, cost, latency_ms):
    # quality already in [0, 1]
    cost_score = max(0.0, 1.0 - cost / 0.05)           # $0 → 1, $0.05 → 0
    latency_score = max(0.0, 1.0 - (latency_ms - 500) / 9500) # 500ms→1, 10s→0
```

```
return 0.50 * quality + 0.30 * cost_score + 0.20 * latency_score
```

The weights (50/30/20) are deliberate. Quality matters more than cost and latency for *most* user-facing tasks. You can change these in code if your deployment cares more about money or speed.

10.4 How the Router Learns

The router has two sources of model scores:

1. **Default scores** — hardcoded in `_DEFAULT_SCORES` at the top of the module. These are the cold-start values before any evaluation has run.
2. **Evaluation results** — when `results_path` is passed at construction, the router reads past evaluation JSON files and updates a running average per (category, model) of quality and latency.

Each call to `update_scores()` appends new evaluation data. The scores **improve over time** as more evaluations are completed. This is how "learns from past evaluations" is realised — there is no neural network; it is a maintained running average.

10.5 TASK_TYPE Routing for HRI

Condition C of the Embodied LLM Arena uses this. It reads the `hri_task_routing` block of `models.yaml`:

```
routing:
  hri_task_routing:
    info_retrieval:    openai-gpt4o-mini    # T1
    navigation:        gemini-flash         # T2
    social_conversation: claude-haiku       # T3
    multilingual:      gemini-flash         # T4
```

The mapping was chosen based on each model's strengths: GPT-4o-mini is factually accurate and good with RAG; Gemini Flash is fast and good with spatial language; Claude Haiku is empathetic and natural.

The implementation in `route()` is:

```
elif strategy == RoutingStrategy.TASK_TYPE:
    hri_map = self._routing_cfg.get("hri_task_routing", {})
    preferred = hri_map.get(task_category)
    if preferred and preferred in self._models and preferred in candidates:
        candidates = [preferred]
    else:
        # Fallback: models that list this task in their hri_strengths
        strength_matches = [m for m in candidates
                           if task_category in self._models[m].get("hri_strengths", [])]
        if strength_matches:
            candidates = strength_matches
```

10.6 Worked Example

```
from omnillm.router import SmartRouter, RoutingStrategy

router = SmartRouter()

# Strategy 1: Best quality, no constraints
d1 = router.route(strategy=RoutingStrategy.BEST_QUALITY)
print(f"Best quality model: {d1.model_id}")
# → openai-o1 (highest default quality score)

# Strategy 2: Cheap mode with budget
d2 = router.route(budget_usd=0.001, strategy=RoutingStrategy.LOWEST_COST)
print(f"Cheap model under $0.001: {d2.model_id}")
# → gemini-flash or an Ollama model

# Strategy 3: HRI navigation task
d3 = router.route_for_hri_task("navigation")
print(f"Best for navigation: {d3.model_id}")
# → gemini-flash (per hri_task_routing config)

# Strategy 4: Composite value with latency cap
d4 = router.route(max_latency_ms=2000, strategy=RoutingStrategy.BEST_VALUE)
print(f"Best value under 2s: {d4.model_id}")
# → likely gemini-flash or claude-haiku
```

Chapter 11 — `consensus.py` — The LLM Council

**** AT A GLANCE.** ConsensusEngine queries N models in parallel, then applies one of three strategies — majority_vote (Jaccard-clustered), weighted (static weights), or synthesis (a judge LLM combines the answers). Returns ConsensusResult with the final answer, individual responses, agreement score, dissenting models, and synthesis reasoning.*

11.1 Why a Council?

A single LLM can be **confidently wrong**. If three independent models all give the same answer, you have much higher confidence the answer is correct. For robotics this is critical — a single hallucinated "turn left" command can be a safety problem. Consensus filters errors that appear in only one model.

References for the design: - arXiv:2601.07245 — "Learning to Trust the Crowd" - Karpathy's "LLM Council" concept - MDPI 2024 — "Multiple Large AI Models' Consensus for Object Detection"

11.2 The Three Strategies

Strategy 1: `majority_vote`

Cluster the responses by **Jaccard similarity** of word tokens. Pick the representative of the largest cluster.

```
def _jaccard_similarity(text_a, text_b):
    tok_a = set(re.findall(r"\b\w+\b", text_a.lower()))
    tok_b = set(re.findall(r"\b\w+\b", text_b.lower()))
    return len(tok_a & tok_b) / len(tok_a | tok_b)
```

Threshold = 0.15. Cheap and dependency-free (no embedding model needed). Works well for short factual answers, less well for long creative answers.

Strategy 2: `weighted`

Each model has a static quality weight (`_WEIGHTS` dict in source). Pick the response from the highest-weighted model. Effectively a "trust the strongest model" fallback. Production should replace static weights with EloScorer ratings.

Strategy 3: `synthesis` (default)

The most powerful. A **judge LLM** sees all council responses and is asked to synthesise the best combined answer. The judge prompt:

```
You are a synthesis judge reviewing multiple AI responses to a question.
Your job is to produce the BEST possible answer by combining insights from
all responses.

**Original Question:**
<the question>

**Council Responses:**
### model-id-1
<response 1>
```

```

### model-id-2
<response 2>
...

Instructions:
1. Identify areas of agreement across responses (likely correct).
2. Note any disagreements or unique insights.
3. Synthesise a final answer that is more accurate and complete than any
   individual response.

Respond with valid JSON: {final_answer, agreement_score, reasoning,
dissenting_models}.

```

This is what Condition D of the Embodied LLM Arena uses.

11.3 Worked Example

```

import asyncio
from omnillm.gateway import LLMGateway
from omnillm.consensus import ConsensusConfig, ConsensusEngine

async def main():
    gw = LLMGateway()
    cfg = ConsensusConfig(
        council_models=["openai-gpt4o-mini", "claude-haiku", "gemini-flash"],
        judge_model="openai-gpt4o-mini",
        strategy="synthesis",
    )
    engine = ConsensusEngine(gw, cfg)
    result = await engine.query_council(
        [{"role": "user", "content": "Is consciousness emergent or fundamental?"}]
    )
    print("Final answer:", result.final_answer)
    print(f"Agreement: {result.agreement_score:.0%}")
    print("Dissenting:", result.dissenting_models)

asyncio.run(main())

```

Chapter 12 — `evaluator.py` — The LLM-as-Judge Pipeline

*** AT A GLANCE. *Three judge patterns, all implemented as one class. Referenceless (G-Eval) scores quality without a gold answer. Reference-Based compares to a known-correct answer. Pairwise picks a winner between two responses, with swap-and-aggregate** to cancel position bias. The judge is itself an LLM — by default GPT-4o.*

12.1 Why LLM-as-Judge?

Human evaluation does not scale. Programmatic evaluation (regex, exact match) only works for tightly-constrained answers. The middle ground is to have an LLM evaluate the response. A 2024 survey (arXiv:2412.05579) found LLM judges agree with human raters about as often as humans agree with each other on most tasks — making LLM-as-Judge methodologically defensible for benchmark studies.

12.2 The Three Judge Patterns

Pattern	Inputs	Output	Use case
Referenceless (G-Eval)	prompt, response	score 0–1 + reasoning	open-ended, no gold answer
Reference-Based	prompt, response, reference	score 0–1 + reasoning	factual, gold known
Pairwise	prompt, two responses	winner: A / B / tie	leaderboard / ELO

12.3 Referenceless Prompt (the actual one)

You are an expert evaluator. Evaluate the following AI response on a scale from 0.0 to 1.0 where 1.0 is perfect.

****Original question:****
<prompt>

****AI Response:****
<response>

Evaluate for: accuracy, completeness, clarity, and helpfulness.

Respond ONLY with a valid JSON object in this exact format:
{ "score": <float 0.0-1.0>, "reasoning": "<one-paragraph explanation>" }

12.4 Reference-Based Prompt

Same idea but the judge sees the gold answer and is asked to rate similarity:

****Question:**** <prompt>
****Reference Answer:**** <gold>
****AI Response:**** <response>

Score from 0.0 (wrong) to 1.0 (perfectly correct).

```
{"score": ..., "reasoning": ...}
```

12.5 Pairwise — The Position-Bias Trick

LLM judges have a documented preference for the **first** response shown. OmniLLM cancels this by running the comparison twice with positions swapped:

```
async def evaluate_pairwise(self, task, model_a, model_b):
    resp_a, resp_b = await asyncio.gather(
        self.gateway.query(model_a, messages),
        self.gateway.query(model_b, messages),
    )
    w1, _ = await self._judge_pairwise(task, resp_a, resp_b, swap=False)
    w2, _ = await self._judge_pairwise(task, resp_a, resp_b, swap=True)
    final = w1 if w1 == w2 else "tie"
    return PairwiseResult(...)
```

If both runs agree, that's the winner. If they disagree, the comparison is declared a **tie** — the position bias was the deciding factor, so neither response is reliably better.

12.6 The Programmatic Escape Hatch

For tasks with a deterministic correctness check (e.g. "does this response parse as JSON with the right schema?"), you can bypass the judge entirely by setting `task.grading_fn`:

```
EvalTask(
    id="cost-yes-no",
    prompt="Is Python interpreted? Answer only Yes or No.",
    reference_answer="Yes",
    grading_fn=lambda r: 1.0 if "yes" in r.strip().lower()[:10] else 0.0,
    judge_pattern="referenceless", # ignored when grading_fn is set
)
```

`grading_fn` takes priority over the judge.

12.7 Running a Full Benchmark

The evaluator concurrent-runs the full task × model matrix:

```
async def run_benchmark(self, tasks, model_ids, max_concurrent=5):
    semaphore = asyncio.Semaphore(max_concurrent)
    async def _limited(task, model_id):
        async with semaphore:
            return await self.evaluate_task(task, model_id)
    coros = [_limited(t, m) for t in tasks for m in model_ids]
    return list(await asyncio.gather(*coros))
```

The semaphore caps simultaneous calls so you don't trigger rate limits.

Chapter 13 — `scorer.py` — The ELO Leaderboard

**** AT A GLANCE.** Implements chess-style ELO ratings. Every pairwise match updates two ratings. K-factor 32, default 1500. Supports per-category leaderboards, rating history, and JSON save/load. Same methodology as LMSYS Chatbot Arena.*

13.1 The ELO Formula

For two players with ratings R_a and R_b , the **expected score** for A is:

$$E_a = \frac{1}{1 + 10^{(R_b - R_a) / 400}}$$

After a match where the actual score for A is $S_a \in \{0, 0.5, 1\}$ (loss / tie / win), the new rating is:

$$R'_a = R_a + K \cdot (S_a - E_a)$$

K is the volatility constant (32 in OmniLLM). Higher K = ratings move faster but are noisier. 100 ELO points $\approx$ 64% expected win rate.

13.2 Why ELO and Not Plain Win-Rate?

Plain win-rate has a flaw: beating a strong opponent is worth the same as beating a weak one. ELO rewards strength of opposition. After enough matches, ELO converges to a stable relative ranking even if matches are not uniformly distributed.

13.3 Implementation Walk-Through

```
class EloScorer:
    def __init__(self, k_factor=32, default_rating=1500):
        self.k_factor = k_factor
        self.default_rating = default_rating
        self.ratings: dict[str, float] = {}
        self.matches_played: dict[str, int] = {}
        self._history: list[_MatchRecord] = []
        self._rating_history: dict[str, list[(str, float)]] = {}
        self._category_ratings: dict[str, dict[str, float]] = {}

    def expected_score(self, ra, rb):
        return 1.0 / (1.0 + math.pow(10, (rb - ra) / 400.0))

    def record_match(self, model_a, model_b, winner, category="general"):
        # Initialise ratings if first match
        for m in (model_a, model_b):
            self.ratings.setdefault(m, self.default_rating)
            self.matches_played.setdefault(m, 0)

        ra_before = self.ratings[model_a]
        rb_before = self.ratings[model_b]
        ra_after, rb_after = self.update_ratings(ra_before, rb_before, winner)

        # Update overall + category + history
        self.ratings[model_a] = ra_after
        self.ratings[model_b] = rb_after
        self.matches_played[model_a] += 1
        self.matches_played[model_b] += 1
        # ... category and history bookkeeping ...
```

13.4 Per-Category Leaderboards

Each match is also tracked under a category (e.g. "reasoning" or "embodied_hri"). Per-category ratings start fresh at 1500 and converge independently.

```
scorer.record_match("gpt-4o", "claude-sonnet", "model_a", category="reasoning")
scorer.record_match("gpt-4o", "claude-sonnet", "model_b", category="code")

scorer.get_leaderboard()           # overall
scorer.get_category_leaderboard("reasoning") # only reasoning matches
scorer.get_category_leaderboard("embodied_hri") # only HRI matches
```

For the Embodied LLM Arena, every pairwise participant preference produces a match in `category="embodied_hri"`, building the *Embodied LLM Leaderboard*.

Chapter 14 — cli.py — Your Terminal Dashboard

***** AT A GLANCE.** *** A Click CLI with 10 commands and Rich-formatted output. Lets you talk to every layer of OmniLLM from the terminal: list models, ask one or many models, run benchmarks, compare two models, run the LLM Council, demo the smart router, see the ELO leaderboard, see costs, and export results. Loaded as the `omnillm` shell command after install.*

14.1 The Command Map

Command	What it does	Most-used flags
<code>omnillm models</code>	List registered models	<code>--type cloud local</code>
<code>omnillm ask "<prompt>"</code>	Send to one or many models	<code>--all, -m id</code>
<code>omnillm evaluate</code>	Run benchmark on tasks × models	<code>-m, -c, -o</code>
<code>omnillm compare "<prompt>"</code>	Pairwise judge	<code>--model-a, --model-b</code>
<code>omnillm council "<prompt>"</code>	LLM Council	<code>--strategy, -m</code>
<code>omnillm route "<prompt>"</code>	Show routing decision	<code>--budget, --strategy</code>
<code>omnillm leaderboard</code>	ELO rankings	<code>--category</code>
<code>omnillm costs</code>	USD spend per model	<code>--data-file</code>
<code>omnillm export</code>	CSV / JSON / Markdown	<code>--format, --input, -o</code>

14.2 How Click + Rich Work Together

Click handles **argument parsing and validation**. Rich handles **output formatting**. Each command is a function decorated with `@cli.command(...)` that builds a `rich.Table` or `rich.Panel` and prints it.

```
@cli.command("models")
@click.option("--config", "-c", default=None)
@click.option("--type", "model_type", default=None,
              type=click.Choice(["cloud", "local"]))
def models_cmd(config, model_type):
    gateway = _get_gateway(config)
    table = Table(title="Registered LLM Models", header_style="bold cyan")
    table.add_column("ID", style="cyan", no_wrap=True)
    table.add_column("Type", style="magenta")
    # ...
    for model_id in gateway.list_models():
        info = gateway.get_model_info(model_id)
        if model_type and info.get("type") != model_type:
            continue
        table.add_row(model_id, info["type"], info["provider"], ...)
    console.print(table)
```

14.3 Async-from-Sync Pattern

Click runs synchronously but our gateway is async. Each command wraps the async work in a local helper:

```
async def _run():
    return await gateway.query_multiple(model_ids, messages)
responses = asyncio.run(_run())
```

`asyncio.run()` creates a fresh event loop, runs the coroutine, and closes the loop. This is the safe way to call async code from a sync entry point.

14.4 Beautiful Output Tricks

A few Rich tricks worth knowing:

- **Coloured score column**: `green ≥ 0.8, yellow ≥ 0.5, red otherwise`.
- **Progress spinner**: `Progress(SpinnerColumn(), TextColumn("..."))` for multi-second waits.
- **Panels**: `Panel(content, title=..., border_style="cyan")` for boxed responses.
- **Medals**: `1st 2nd 3rd` for the top 3 leaderboard rows.

Chapter 15 — rag/pipeline.py — Retrieval-Augmented Generation

***** AT A GLANCE.** Indexes documents (TXT, CSV, PDF) into ChromaDB. On query, retrieves top-k chunks via cosine similarity, prepends them to the LLM prompt as context, and asks the LLM to answer using only the context. Optional faithfulness scoring (LLM-as-Judge) and hallucination detection (word-overlap heuristic). Falls back gracefully to keyword search when ChromaDB is not installed.


15.1 Why RAG?

A general-purpose LLM does not know your specific lab — your room numbers, your staff, your seminar schedule, your wifi password. Two options:

- **Fine-tune a model on your data.** Expensive, slow, must redo for every update.
- **Show the model your documents at query time.** Cheap, fast, instantly reflects updates. This is RAG.

For OmniLLM, RAG enables Tasks T1 (Information Retrieval) and T2 (Navigation) — both depend on facts in `knowledge_base/`.

15.2 The Three Phases



15.3 Chunking — Why 512 / 64?

The pipeline uses a sliding-window chunker:

```
def _split_text(self, text, source, metadata):
    chunks = []
    start = 0; idx = 0
    while start < len(text):
        end = min(start + self.chunk_size, len(text)) # 512
        chunk_text = text[start:end].strip()
        if chunk_text:
            chunks.append(DocumentChunk(text=chunk_text, source=source, ...))
        start += self.chunk_size - self.chunk_overlap # advance 512 - 64 = 448
        idx += 1
    return chunks
```

- **512 characters** is a useful unit: long enough to contain a complete thought, short enough that several can fit in one prompt.
- **64-character overlap** prevents losing context at chunk boundaries (a sentence split across chunks still appears intact in at least one).

15.4 Retrieval — the Dual Path

```
def retrieve(self, query: str) -> list[DocumentChunk]:
    if self._collection is not None: # ChromaDB available
        return self._retrieve_chromadb(query)
    return self._retrieve_keyword(query) # fallback path
```

ChromaDB path: convert query to embedding (handled by ChromaDB's default), do cosine similarity search, return top-4 chunks with similarity scores in [0, 1].

Keyword fallback: simple Jaccard similarity of stemmed words. Less accurate but **always works**, even on a stripped-down install where `pip install chromadb` failed.

15.5 The Augmented Prompt

```
system = "You are a helpful assistant. Answer the user's question using ONLY "  
         "the provided context. If the answer is not in the context, say so clearly."  
  
user_content = (  
    "Context:\n"  
    "[1] (Source: lab_info.txt)\nThe IRAI Lab is located in Building C, Room 305...\n\n"  
    "[2] (Source: faq.txt)\nQ: What are the lab opening hours?\nA: Mon-Fri 08:00-20:00...\n\n"  
    "Question: What are the lab hours?"  
)  
  
messages = [  
    {"role": "system", "content": system},  
    {"role": "user", "content": user_content},  
]
```

The "ONLY" instruction is critical. It tells the LLM to refuse to invent information — the foundation of grounded answering.

15.6 Faithfulness Scoring

Optional. Costs an extra LLM call but gives a quantitative measure of how well the answer used the retrieved context:

```
prompt = (
    "You are a factuality judge. Score how faithfully the answer uses ONLY "
    "information from the given context (0.0 = completely hallucinated, "
    "1.0 = every claim is supported by the context).\n\n"
    f"Context:\n{context}\n\nQuestion: {question}\n\nAnswer: {answer}\n\n"
    'Respond ONLY with valid JSON: {"score": <float>, "reasoning": "..."}'
)
```

The score is recorded in the experiment log alongside latency and cost.

15.7 Hallucination Detection (Heuristic)

A cheap word-overlap check:

```
def _detect_hallucination(self, answer, chunks):
    if not chunks:
        return False
    all_context = " ".join(c.text.lower() for c in chunks)
    answer_words = re.findall(r"\b\w{5,}\b", answer.lower()) # words ≥5 chars
    if not answer_words:
        return False
    found = sum(1 for w in answer_words if w in all_context)
    return (found / len(answer_words)) < 0.20 # < 20% overlap = suspect
```

This is a *flag*, not a verdict. False positives are common (the answer might paraphrase the context). For final scoring, prefer the faithfulness LLM-as- Judge.

Chapter 16 — `hri/classifier.py` and `hri/language_detector.py`

***** AT A GLANCE.**** Two short files together produce the input metadata the agent graph needs. The classifier maps an utterance to one of T1–T4 using keyword + regex scoring (rule-based, instant, no LLM call). The language detector uses Unicode script analysis, n-gram word matching, and optionally `langdetect`, to identify the language and recommend a model.

16.1 Why Rule-Based and Not LLM?

These are both pre-LLM steps. They run on every interaction. Latency is critical. A rule-based classifier returns in 10–20 ms; an LLM call takes 500+ ms. For tasks this simple, the rules are accurate enough (~92% in testing) that the latency saving is worth it.

That said, both modules **also** offer LLM-based variants for high-stakes or ambiguous inputs (`classify_with_llm()`).

16.2 The Classifier — Keyword Scoring

Each task category has a keyword list and a regex-pattern list:

```
_NAVIGATION_KEYWORDS = frozenset([
    "where", "room", "floor", "cafeteria", "toilet", "elevator", "lift",
    "point", "direction", "guide", "navigate", "take me", "show me",
    "find", "locate", "map", "route",
    # ...
])

_NAVIGATION_PATTERNS = [
    re.compile(r"\broom\s+\d+\b", re.IGNORECASE),          # "Room 305"
    re.compile(r"\bwhere\s+(is|are|can i find)\b", re.IGNORECASE),
    re.compile(r"\b(go to|take me to|guide me to)\b", re.IGNORECASE),
    # ...
]
```

The score for a category is computed by combining:

1. **Keyword overlap**: number of words in the utterance matching the keyword set, capped to 0.3.
2. **Phrase keyword bonus**: 0.15 per multi-word keyword like "how do i get to".
3. **Pattern match bonus**: 0.25 per regex hit.

Score is capped at 1.0 per category. The category with the highest score wins. If no category exceeds 0.1, default to `INFO_RETRIEVAL` with low confidence.

Confidence is normalised so a score of 0.1 maps to 0.5, and scores ≥ 0.5 map to 0.95.

16.3 Why Multilingual Beats Everything


If `detected_language != "en"`, the classifier short-circuits to `MULTILINGUAL` immediately:

```
if detected_language != "en":
    return ClassificationResult(
        task_type=HRITaskType.MULTILINGUAL, confidence=0.99,
        reasoning=f"Non-English input detected (language={detected_language})",
        ...
    )
```

)

This is a deliberate design choice. A French-speaker asking "where is room 305?" will route to T4, not T2, because the **language barrier** is the dominant problem to solve.

16.4 The Language Detector — Three-Tier Cascade



Tier 1 catches Arabic, Chinese, Japanese, Korean, Russian, Greek, Hebrew, Thai instantly via Unicode block analysis. Tier 2 catches French, German, Spanish, Italian, Portuguese, Dutch, Turkish via high-frequency word detection ("le", "der", "el", etc.). Tier 3 falls back to the `langdetect` library when neither is decisive.

16.5 Language → Model Mapping

After detecting the language, the detector recommends a model:

```
_LANGUAGE_MODEL_MAP = {
    "en": "openai-gpt4o-mini", # English → cheap and accurate
    "fr": "gemini-flash",      # French → strong multilingual
    "de": "gemini-flash",
    "ja": "gemini-flash",
    "zh": "gemini-flash",
    "ar": "gemini-flash",
    # ...
}
```

The choice of Gemini Flash for non-English is empirical: it consistently performs strongly across the multilingual axis at very low cost.

Chapter 17 — `hri/agent_graph.py` — The Brain of the Robot

**** AT A GLANCE.** This is the most complex file in the project. It builds the LangGraph state machine described in Chapter 7, with one factory function per node so the graph builder can wire them up. Every node is a closure over the gateway, RAG pipeline, and logger so it has the dependencies it needs without global state.*

17.1 The `build_hri_graph()` Top Level

```
def build_hri_graph(gateway, rag=None, logger=None,
                    default_model="openai-gpt4o-mini"):
    from langgraph.graph import StateGraph, END

    builder = StateGraph(dict)  # state is a plain dict

    builder.add_node("transcribe_audio", _make_transcribe_node(gateway))
    builder.add_node("detect_language", _make_detect_language_node())
    builder.add_node("classify_task", _make_classify_task_node())
    if rag is not None:
        builder.add_node("rag", _make_rag_node(rag))
        builder.add_node("nav_rag", _make_nav_rag_node(rag))
    else:
        builder.add_node("rag", _make_direct_llm_node(gateway, default_model))
        builder.add_node("nav_rag", _make_direct_llm_node(gateway, default_model))
    builder.add_node("direct_llm", _make_direct_llm_node(gateway, default_model))
    builder.add_node("multilingual_llm", _make_multilingual_llm_node(gateway))
    builder.add_node("smart_router", _make_smart_router_node(gateway))
    builder.add_node("generate_action_plan", _make_action_plan_node())
    builder.add_node("log_interaction", _make_log_node(logger))

    builder.set_entry_point("transcribe_audio")
    builder.add_edge("transcribe_audio", "detect_language")
    builder.add_edge("detect_language", "classify_task")

    builder.add_conditional_edges("classify_task", _route_by_task_type, {
        "rag": "rag",
        "nav_rag": "nav_rag",
        "direct_llm": "direct_llm",
        "multilingual_llm": "multilingual_llm",
    })

    for n in ("rag", "nav_rag", "direct_llm", "multilingual_llm"):
        builder.add_edge(n, "smart_router")
    builder.add_edge("smart_router", "generate_action_plan")
    builder.add_edge("generate_action_plan", "log_interaction")
    builder.add_edge("log_interaction", END)

    return builder.compile()
```

17.2 The Closure Pattern

Each `_make*_node` returns an async function (the actual node) that has captured its dependencies. Example:

```
def _make_transcribe_node(gateway):
    async def transcribe_audio(state: dict) -> dict:
        audio = state.get("audio_bytes", b"")
        if not audio:
            # Text-only mode: utterance already set
            return {"_start_time": time.monotonic()}
        try:
            from omnillm.robotics.whisper_stt import WhisperSTT
```

```

        stt = WhisperSTT()
        utterance = await stt.transcribe(audio)
        return {"utterance": utterance, "_start_time": time.monotonic()}
    except Exception as exc:
        return {"error": f"Transcription failed: {exc}",
                "_start_time": time.monotonic()}
    return transcribe_audio

```

Closures are why this file is short despite each node being conceptually distinct — there is no parameter passing, no global state, no class hierarchy. Each node is a pure function of state.

17.3 The Conditional Edge Function

This is the dispatching logic that gives Conditions A–E their behaviour:

```

def _route_by_task_type(state: dict) -> str:
    task_type = state.get("task_type", "info_retrieval")
    condition = state.get("condition", "A")
    rag_enabled = state.get("rag_enabled", True)

    # Condition E: RAG-Off Control always goes to direct LLM
    if condition == "E":
        return "direct_llm"

    routing = {
        "info_retrieval": "rag" if rag_enabled else "direct_llm",
        "navigation": "nav_rag" if rag_enabled else "direct_llm",
        "social_conversation": "direct_llm",
        "multilingual": "multilingual_llm",
    }
    return routing.get(task_type, "direct_llm")

```

Three things to notice:

1. **Condition E** wins over task type. The whole point of E is to disable RAG, so we always route to `direct_llm`.
2. **Other conditions** still respect task type. Condition A (cloud) and Condition B (local) both follow the task → branch routing. They differ only in *which model* the branch ends up calling — that decision is made in `smart_router` or by the default model.
3. `rag_enabled=False` also triggers `direct_llm` even within Conditions A/B/C/D (e.g., for ablation studies).

17.4 The Smart-Router Node

This is where Conditions C and D actually take effect:

```

def _make_smart_router_node(gateway):
    async def smart_router(state: dict) -> dict:
        condition = state.get("condition", "A")
        if state.get("response_text") and condition not in ("C", "D"):
            return {} # already answered, no further routing

        if condition == "D":
            # Consensus: query 3 models, synthesise
            from omnillm.consensus import ConsensusEngine
            engine = ConsensusEngine(gateway=gateway)
            council = ["openai-gpt4o-mini", "gemini-2.5-flash", "claude-haiku"]
            ...
            council_resp = await engine.query(council, messages)
            return {"response_text": council_resp.synthesis,

```

```

        "model_id": "council:" + "+".join(council)}

    elif condition == "C":
        # Smart-route per task type
        from omnillm.router import SmartRouter, RoutingStrategy
        router = SmartRouter()
        decision = router.route_for_hri_task(
            task_type=state.get("task_type"),
            strategy=RoutingStrategy.TASK_TYPE,
        )
        target = decision.model_id
        ...
        resp = await gateway.query(target, messages)
        return {"response_text": resp.content, "model_id": resp.model_id, ...}
    return smart_router

```

17.5 Logging — the Final Node

This is where `ExperimentLogger` records the interaction. Wrapped in a try/except because **logging failure should never crash the interaction**:

```

def _make_log_node(logger):
    async def log_interaction(state: dict) -> dict:
        if logger is None:
            return {}
        start = state.get("_start_time", time.monotonic())
        total_latency = (time.monotonic() - start) * 1000
        try:
            logger.log_interaction(
                session_id=state.get("session_id", ""),
                participant_id=state.get("participant_id", ""),
                condition=state.get("condition", "A"),
                task_type=state.get("task_type"),
                utterance=state.get("utterance", ""),
                response=state.get("response_text", ""),
                model_id=state.get("model_id", ""),
                latency_ms=total_latency,
                # ... full set of fields ...
            )
        except Exception:
            pass
        return {"latency_ms": total_latency}
    return log_interaction

```

17.6 Testing the Graph

The graph is testable end-to-end without any robot:

```

import asyncio
from omnillm.gateway import LLMGateway
from omnillm.rag import RAGPipeline
from omnillm.hri import build_hri_graph

async def main():
    gw = LLMGateway()
    rag = RAGPipeline(gateway=gw)
    rag.index_directory("knowledge_base/")
    graph = build_hri_graph(gateway=gw, rag=rag)

    result = await graph.ainvoke({
        "utterance": "Where is Room 305?",
        "participant_id": "TEST",
        "session_id": "test-1",
        "condition": "C",
        "rag_enabled": True,
    })

```

```
print(result["response_text"])
print(result["robot_action"])

asyncio.run(main())
```

Chapter 18 — `hri/experiment.py` — Conditions, Sessions, Counterbalancing

**** AT A GLANCE.** Defines the five experimental conditions A–E as dataclasses, manages participant sessions (UUIDs, start/end times, per-task logs), and stores per-condition `ConditionConfig` objects describing which model, RAG state, council, and routing strategy each condition uses.*

18.1 The `ExperimentCondition` Enum

A simple string enum:

```
class ExperimentCondition(str, Enum):
    A = "A" # Fixed Cloud LLM (GPT-4o-mini)
    B = "B" # Fixed Local LLM (Llama3:8b via Ollama)
    C = "C" # Smart-Routed (per task type)
    D = "D" # Consensus (3-model council)
    E = "E" # RAG-Off Control (same as A but RAG=off)
```

18.2 The `conditionConfig` Dataclass

Each condition has a configuration object:

```
@dataclass
class ConditionConfig:
    condition: ExperimentCondition
    model_id: str | None # None = dynamic (C/D)
    rag_enabled: bool
    use_consensus: bool = False
    use_smart_routing: bool = False
    council_models: list[str] = field(default_factory=list)
    description: str = ""
```

The five default configs are wired up at module load:

Cond	model_id	RAG	smart_routing	consensus
A	openai-gpt4o-mini	✓	—	—
B	llama3-8b-local	✓	—	—
C	None (dynamic)	✓	✓	—
D	None (dynamic)	✓	—	✓ (gpt4o-mini + claude-haiku + gemini-flash)
E	openai-gpt4o-mini	✗	—	—

18.3 The `ParticipantSession` Dataclass

A live record of one participant's session under one condition:

```

@dataclass
class ParticipantSession:
    session_id: str
    participant_id: str
    condition: ExperimentCondition
    condition_config: ConditionConfig
    start_time: str = ... # ISO8601
    end_time: str | None = None
    task_log: list[dict] = []
    questionnaire_scores: dict[str, float] = {}
    notes: str = ""

    def complete(self):
        self.end_time = datetime.now(timezone.utc).isoformat()

    def log_task(self, task_record):
        task_record.setdefault("timestamp", ...)
        self.task_log.append(task_record)

    def to_dict(self): ...

```

18.4 The ExperimentManager

The manager creates sessions, returns config for any condition, and provides filters for analysis:

```

manager = ExperimentManager()
session = manager.create_session("P001", "C")
print(session.condition_config.use_smart_routing) # True
session.log_task({"task_type": "info_retrieval",
                 "utterance": "What time does the lab open?", ...})
session.complete()

# Analysis filters
manager.get_sessions_by_participant("P001")
manager.get_sessions_by_condition("C")
manager.summarise() # totals, per-condition counts
manager.export_data() # list of dicts for stats

```

18.5 Counterbalancing

The current implementation does **not** automatically generate Latin-square condition orders — that's typically driven by a spreadsheet during experiment planning. The manager simply tags each session with the condition the experimenter chose, and the analysis step (R, SPSS, JASP) handles the within-subjects ANOVA / Friedman test on the exported data.

If you want automated counterbalancing, you can add an `assigned_conditions` method that returns a Latin-square slice for a given participant index — but in practice, with $N = 15\text{--}25$ participants, manually managing the order in a CSV is faster than coding it.

Chapter 19 — `robotics/` — Bridges, Gestures, Whisper

**** AT A GLANCE.** This sub-package contains everything that talks to a robot. The abstract `RobotBridge` defines the interface; concrete implementations exist for Pepper (`pepper.py`), NAO (`nao.py`), and Buddy (`buddy.py`). `gesture_planner.py` maps task type + response text → gesture name + LED hex. `whisper_stt.py` wraps OpenAI Whisper for both local and API backends.*

19.1 The Abstract Bridge

```
class RobotBridge(ABC):
    @abstractmethod
    async def connect(self) -> bool: ...
    @abstractmethod
    async def disconnect(self) -> bool: ...
    @abstractmethod
    async def execute_action(self, action: RobotAction) -> bool: ...
    @abstractmethod
    async def get_sensor_data(self) -> RobotSensorData: ...
    @abstractmethod
    async def say(self, text: str) -> bool: ...
    @abstractmethod
    async def gesture(self, name: str) -> bool: ...
```

The `RobotAction` dataclass is the universal command:

```
@dataclass
class RobotAction:
    speech: str = ""
    gesture: str | None = None
    movement: dict | None = None # {direction, speed, distance_m}
    emotion_led: str | None = None # "#00FF00" hex
    nav2_goal: dict | None = None # {position, orientation}
    metadata: dict = field(default_factory=dict)
```

This is the **boundary between brain and body**. The brain produces a `RobotAction`. The body (one of the bridge implementations) executes it.

19.2 The Pepper Bridge

Bridge	Robot	Transport	NAOqi Python 2.7?
<code>PepperBridge</code>	Pepper	HTTP to NAOqi bridge server	Yes (separate process)

Earlier drafts of this codebase shipped two additional bridges (`NAOBridge` for the smaller NAO robot, `BuddyBridge` for the Android-based Buddy companion). They were removed to keep the surface area focused on the platform actually used in the experiments. The bridge interface (`omnillem/robotics/bridge.py`) is unchanged, so adding a new robot back later is a one-file exercise.

19.3 Why HTTP to NAOqi (Not Direct)?

You might think "the AI server could `import naoqi` and call NAOqi directly". This does not work because:

- **NAOqi is Python 2.7 only.**

- **Modern AI libraries (LangGraph, LiteLLM) require Python 3.11+.**
- **You cannot import both into the same Python process.**

So the AI server (Python 3) HTTPs into a small Flask process (Python 2.7) running on or near the robot. The Python 2.7 side imports `naoqi` and performs the actual hardware calls. See Chapter 26 for the full bridge discussion.

19.4 `gesture_planner.py`

A small but high-leverage file. Three priority levels, in order:

1. **Direction patterns** (only for navigation tasks). If the response says "on your left", the gesture is `point_left`. Direct mapping wins over everything else.
2. **Content keywords**. "hello/hi/welcome" → `wave`. "goodbye" → `wave_goodbye`. "thinking/let me check" → `think`. "sorry/don't know" → `confused`. "absolutely/certainly/yes" → `nod`.
3. **Task-type defaults**. `info_retrieval` → `nod`. `navigation` → `point_forward`. `social_conversation` → `wave`. `multilingual` → `nod`.

Each gesture also carries a default LED colour:

Gesture	LED
wave / nod / happy	#00FF88 (friendly green)
wave_goodbye	#FF8800 (warm orange)
point_*	#00AAFF (calm blue, navigation mode)
show_tablet	#FFFFFF (white)
think	#FFFF00 (yellow)
confused	#FF4400 (red-orange)
neutral	#44AAFF (default blue)

19.5 `whisper_stt.py` — Local vs API

Two backends, same interface:

```
stt = WhisperSTT(backend="local", model_size="base")
stt = WhisperSTT(backend="api")

text = await stt.transcribe(wav_bytes)
```

Backend	Cost	Privacy	Performance	Install
local	\$0	private	needs a GPU for "base" or larger	<code>pip install openai-whisper</code>

api	~\$0.006/min	sent to OpenAI	fast on any hardware	<code>pip install openai</code> + key
-----	--------------	----------------	----------------------	---------------------------------------

Local Whisper sizes (in trade-off order):

Size	RAM	Speed	Quality
tiny	~1 GB	fastest	lowest
base	~1.5 GB	fast	good (default)
small	~2 GB	moderate	better
medium	~5 GB	slow	very good
large-v3	~10 GB	slowest	best


For Pepper interactions, **base** is usually the right default — accurate enough on simple lab utterances, fast enough to fit in the latency budget.

Chapter 20 — `server/app.py` — The Flask AI Server

***** AT A GLANCE.** *Flask app with six endpoints. On startup it loads `config/models.yaml`, indexes `knowledge_base/`, and lazily compiles the LangGraph pipeline. `POST /interact` is the main endpoint — accepts audio or text, runs the graph, returns the RobotAction JSON. Falls back to `_fallback_interact` if LangGraph is not installed.*

19.1 — Where the Server Sits

This is the **central nervous system** of the running project:



20.2 The `create_app()` Factory

The whole Flask app is built by a factory function — useful for testing and for gunicorn:

```
def create_app(knowledge_base_dir=None, default_model=None,
               enable_rag=True, log_to_file=None):
    from flask import Flask, jsonify, request

    app = Flask(__name__)

    # Configuration via env vars, with sensible defaults
    _default_model = (default_model
                      or os.getenv("OMNILLM_DEFAULT_MODEL", "openai-gpt4o-mini"))
    _kb_dir = Path(knowledge_base_dir
                  or os.getenv("OMNILLM_KNOWLEDGE_BASE",
                              str(Path(__file__).parent.parent.parent / "knowledge_base")))

    # Initialise components
    gateway = LLMGateway()
    exp_logger = ExperimentLogger()
    rag = None
    if enable_rag:
        rag = RAGPipeline(gateway=gateway, model_id=_default_model)
        _load_knowledge_base(rag, _kb_dir)

    # Lazy LangGraph build (so tests don't need it)
    _graph = None
    def _get_graph():
        nonlocal _graph
```

```

    if _graph is None:
        try:
            from omnillm.hri.agent_graph import build_hri_graph
            _graph = build_hri_graph(gateway=gateway, rag=rag,
                                     logger=exp_logger,
                                     default_model=_default_model)
        except ImportError:
            logger.warning("LangGraph not installed – falling back.")
    return _graph

# ... endpoints registered below ...

return app

```

20.3 The `/interact` Endpoint, Annotated

This is the most important endpoint:

```

@app.route("/interact", methods=["POST"])
def interact():
    data = request.get_json(force=True) or {}

    utterance = data.get("text", "")
    audio_b64 = data.get("audio", "")
    participant_id = data.get("participant_id", "unknown")
    session_id = data.get("session_id", "")
    condition = data.get("condition", "A")
    rag_enabled = bool(data.get("rag_enabled", True) and rag is not None)

    if not utterance and not audio_b64:
        return jsonify({"error": "Provide 'text' or 'audio' field"}), 400

    audio_bytes = base64.b64decode(audio_b64) if audio_b64 else b""

    graph = _get_graph()
    if graph is not None:
        state = {
            "utterance": utterance,
            "audio_bytes": audio_bytes,
            "participant_id": participant_id,
            "session_id": session_id,
            "condition": condition,
            "rag_enabled": rag_enabled,
            "model_id": _default_model,
        }
        result = _run_async(graph.ainvoke(state))
        return jsonify(result.get("robot_action",
                                   {"speech": result.get("response_text", "")}))

    # Fallback if LangGraph not installed
    return _fallback_interact(...)

```

20.4 The `_run_async` Helper

Flask is sync; LangGraph is async. The bridge:

```

def _run_async(coro):
    loop = asyncio.new_event_loop()
    try:
        return loop.run_until_complete(coro)
    finally:
        loop.close()

```

A new event loop per request keeps the request-handling thread isolated. For a higher-throughput deployment you would use an ASGI server (FastAPI + uvicorn) instead — but for a single-robot HRI

study, this is plenty.

20.5 The Other Endpoints

Endpoint	Purpose	Notes
GET /health	Liveness probe	returns {status: "ok", version: ...}
GET /status	Configuration dump	model list, RAG flag, KB path, etc.
POST /transcribe	STT only	for clients that do their own classification
POST /evaluate	Submit questionnaire	logged as a special interaction
GET /export	Dump all logs as JSON	used by experimenter at study end

20.6 Running the Server

From the command line:

```
python -m omnillm.server.app # localhost:5000
python -m omnillm.server.app --host 0.0.0.0 --port 5000 # accept LAN
python -m omnillm.server.app --no-rag # disable RAG (for debugging)
python -m omnillm.server.app --model gemini-flash # change default
python -m omnillm.server.app --kb /path/to/my-kb # custom KB
python -m omnillm.server.app --debug # Flask debug mode
```

For production-ish deployment:

```
gunicorn 'omnillm.server.app:create_app()' --bind 0.0.0.0:5000 --workers 1
```

(!) **Use `--workers 1`** when serving a single robot. Multiple workers do not share the in-memory ChromaDB collection, the experiment logger, or the compiled LangGraph instance. Multiple workers would mean inconsistent state across requests.


Chapter 21 — `server/naoqi_client.py` — The Python 2.7 Side

***** AT A GLANCE. *This is the* only Python 2.7 file** in the project. It runs on or near Pepper, imports `qi` or `naoqi`, captures audio, POSTs it to the AI server over HTTP, receives the `RobotAction` JSON, and executes speech / gesture / LED on the robot. Falls back to "simulation mode" if NAOqi is not installed — useful for testing on a non-NAOqi machine.**

21.1 Why Python 2.7

Pepper's NAOqi SDK is locked to Python 2.7. There is no Python 3 port that runs all NAOqi services. SoftBank announced and abandoned a Python 3 binding called `qi 3.1.5` but several core services are broken in it (touch detection, audio callbacks). For production, **assume Python 2.7**.

21.2 Two-Process Architecture



The Python 2.7 process is **deliberately small**. It does I/O with the robot and HTTP with the server. It does not import any AI library. Everything that *could* break stays on the Python 3 side.

21.3 The Class Skeleton

```
class PepperNAOqiClient(object):
    def __init__(self, robot_ip, robot_port, server_ip, server_port,
                 participant_id, condition):
        self.robot_ip = robot_ip
        self.robot_port = robot_port
        self.server_url = "http://{}:{ {}".format(server_ip, server_port)
        self.participant_id = participant_id
        self.condition = condition
        self.session_id = str(uuid.uuid4())

        self._app = None
        self._audio_device = None
        self._animated_speech = None
        self._motion = None
        self._leds = None
        self._tablet = None
        self._behavior = None
        self._face_detection = None
        self._connected = False

    def connect(self): ...
    def run(self): ...
    def _interaction_loop(self): ...
    def _record_audio(self, duration_seconds=5): ...
    def _send_audio(self, audio_bytes): ...
    def _send_text(self, text): ...
    def _execute_action(self, action): ...
    def _speak(self, text): ...
    def _set_leds(self, hex_color): ...
    def _run_behavior_async(self, behavior_name): ...
    def _cleanup(self): ...
```

21.4 Connecting via NAOqi

```
def connect(self):
    if not NAOqi_AVAILABLE:
        print("[SIM] NAOqi not available – simulation mode")
        self._connected = True
        return True
    try:
        self._app = qi.Application(["PepperClient",
                                     "--qi-url={}:{}".format(self.robot_ip, self.robot_port)])
        self._app.start()
        session = self._app.session

        self._audio_device = session.service("ALAudioDevice")
        self._animated_speech = session.service("ALAnimatedSpeech")
        self._motion = session.service("ALMotion")
        self._leds = session.service("ALLeds")
        self._tablet = session.service("ALTabletService")
        self._behavior = session.service("ALBehaviorManager")
        self._face_detection = session.service("ALFaceDetection")

        self._motion.wakeUp() # turns on stiffness – must happen before any move
        self._connected = True
        return True
    except Exception as exc:
        print("[ERR] Failed to connect: {}".format(exc), file=sys.stderr)
        return False
```

21.5 The Main Loop

```
def run(self):
    if not self._connected and not self.connect():
        return
    self._speak("Hello! I am Pepper, powered by OmniLLM. How can I help?")
    try:
        while True:
            self._interaction_loop()
    except KeyboardInterrupt:
        pass
    finally:
        self._cleanup()

def _interaction_loop(self):
    audio_bytes = self._record_audio(duration_seconds=5)
    if not audio_bytes:
        return
    response = self._send_audio(audio_bytes)
    if response is None:
        self._speak("I'm sorry, I could not connect to my AI brain.")
        return
    self._execute_action(response)
```

21.6 Executing the RobotAction

```
def _execute_action(self, action):
    speech = action.get("speech", "")
    gesture = action.get("gesture")
    led_color = action.get("emotion_led")

    if led_color:
        self._set_leds(led_color)
    if gesture and gesture in GESTURE_TO_BEHAVIOR:
        self._run_behavior_async(GESTURE_TO_BEHAVIOR[gesture])
    if speech:
        self._speak(speech)

def _speak(self, text):
    config = {"bodyLanguageMode": "contextual"}
    self._animated_speech.say(str(text), config)

def _set_leds(self, hex_color):
    hex_color = hex_color.lstrip("#")
    r, g, b = (int(hex_color[i:i+2], 16) / 255.0 for i in (0, 2, 4))
    self._leds.fadeRGB("FaceLeds", r, g, b, 0.3)
```

The `bodyLanguageMode: "contextual"` setting is what makes Pepper move its arms and head while speaking. Without it, you get a stiff, talking-head robot. Chapter 23 explains why this matters.

21.7 The `GESTURE_TO_BEHAVIOR` Dictionary

This is the bridge between OmniLLM's gesture *names* (e.g. `"point_left"`) and the actual NAOqi *behaviour paths* (e.g. `"animations/Stand/Gestures/Explain_8"`):

```
GESTURE_TO_BEHAVIOR = {
    "wave": "animations/Stand/Gestures/Hey_1",
    "bow": "animations/Stand/Gestures/BowShort_1",
    "wave_goodbye": "animations/Stand/Gestures/Farewells_1",
    "point_left": "animations/Stand/Gestures/Explain_8",
    "point_right": "animations/Stand/Gestures/Explain_7",
    "point_forward": "animations/Stand/Gestures/Explain_1",
    "point_up": "animations/Stand/Gestures/Explain_6",
    "show_tablet": "animations/Stand/Gestures/ShowTablet_1",
    "nod": "animations/Stand/Emotions/Positive/Enthusiastic_1",
    "think": "animations/Stand/Emotions/Neutral/Thinking_1",
    "happy": "animations/Stand/Emotions/Positive/Happy_4",
    "confused": "animations/Stand/Emotions/Negative/Confused_1",
}
```

If you install custom Choregraphe behaviours, add them here.

21.8 Audio Capture — Currently a Placeholder

Look closely at `_record_audio`:

```
def _record_audio(self, duration_seconds=5):
    if not NAOQI_AVAILABLE or self._audio_device is None:
        return b""
    try:
        self._audio_device.setClientPreferences(
            "OmniLLMCapture", 16000, 3, 0 # 16 kHz, front mic, deinterleaved=0
        )
        time.sleep(duration_seconds)
        # In real implementation: use ALAudioRecorder or callback-based capture
        return b"" # ← placeholder
    except Exception as exc:
```

```
return b""
```

This is **deliberately incomplete** in the current source. To make it production-ready you have two options:

1. **ALAudioRecorder** — record to a file on Pepper, SCP it back to your PC, then base64-encode and send. Simpler but adds a per-utterance file transfer.
2. **processRemote callback** — subclass `ALModule`, implement `processRemote(nbChannels, nbSamples, timeStamp, inputBuffer)`, accumulate raw int16 samples in a numpy array, write a WAV header, send. Real-time but requires more NAOqi knowledge.

For the present project, the workaround during development is to use the **text-mode** path (`_send_text`) instead of audio capture.

21.9 Running the Client

```
# Set the pynaoqi path so Python 2.7 finds NAOqi
set PYTHONPATH=C:\pynaoqi\pynaoqi-python2.7-2.5.5.5-win32-vs2013\lib;%PYTHONPATH%

# Run
C:\Python27\python.exe omnillm/server/naoqi_client.py \
  --robot-ip 192.168.1.100 \
  --server-ip localhost \
  --participant P001 \
  --condition C
```

Chapter 22 — `utils/` — Logging, Cost Tracking, Questionnaires, Export

**** AT A GLANCE.** Four small utility modules. `cost_tracker.py` accumulates USD per model per session. `experiment_logger.py` records every HRI interaction with full metadata. `questionnaire.py` defines the 1–7 Likert, Godspeed, pairwise, and observer dataclasses. `export.py` writes CSV, JSON, or Markdown reports.*

22.1 `cost_tracker.py`

A flat list of `_Record` objects with summarisation methods:

```
tracker = CostTracker()
tracker.record("openai-gpt4o", input_tokens=500, output_tokens=300,
              cost_usd=0.0045, session_id="eval-1")
print(tracker.get_total_cost())           # 0.0045
print(tracker.get_cost_by_model())        # {'openai-gpt4o': 0.0045}
print(tracker.get_cost_by_session())      # {'eval-1': 0.0045}
print(tracker.get_token_usage())          # {'openai-gpt4o': {'input': 500, 'output': 300, 'total': 800}}
tracker.save("results/costs.json")
```

The CLI command `omnillm costs` builds and prints a Rich table from this tracker.

22.2 `experiment_logger.py` — The Star of the Sub-package

This file produces the dataset that the entire experiment is about. Every interaction creates an `InteractionRecord`:

```
@dataclass
class InteractionRecord:
    session_id: str
    participant_id: str
    condition: str # "A"-"E"
    task_type: str # T1-T4
    utterance: str
    response: str
    model_id: str
    latency_ms: float = 0.0
    input_tokens: int = 0
    output_tokens: int = 0
    cost_usd: float = 0.0
    rag_enabled: bool = False
    rag_faithfulness: float = -1.0 # -1 = not scored
    rag_chunk_count: int = 0
    judge_score: float = -1.0
    task_success: bool | None = None
    language: str = "en"
    gesture_used: str | None = None
    timestamp: str = ... # ISO8601
    notes: str = ""
```

The logger has a `log_interaction(...)` method called from the `LangGraph`'s final node, plus convenience methods for analysis:

```
logger.get_records(condition="C")
logger.get_records(model_id="claude-haiku")
logger.get_summary() # by_model, by_condition, by_task_type
```

```

logger.save("results.json")
logger.save_csv("results.csv")
logger.load("results.json")

```

The CSV export is **flat** — every interaction is one row, every field a column. This is the format that R, SPSS, and JASP love.

22.3 questionnaire.py — Four Dataclasses

Dataclass	Scale	Items	When
InteractionQuestionnaire	1–7 Likert	5 (accuracy, naturalness, trust, gesture, speed)	After each condition
GodspeedResponse	1–5	5 subscales (anthropomorphism, animacy, likeability, perceived intelligence, safety)	After each condition
PairwisePreference	A / B / tie	"Which Pepper did you prefer?"	End of session
ObserverRating	1–5 + counts	gesture-sync quality, breakdown count, task completion	Live during interaction

A typical post-session collection flow:

```

collector = QuestionnaireCollector()
collector.add_interaction_response(InteractionQuestionnaire(
    session_id="s1", participant_id="P001", condition="C",
    accuracy=6, naturalness=5, trust=6,
    gesture_appropriateness=5, response_speed=7,
))
collector.add_godspeed(GodspeedResponse(
    session_id="s1", participant_id="P001", condition="C",
    anthropomorphism=3.2, animacy=3.4, likeability=4.1,
    perceived_intelligence=4.0, perceived_safety=4.5,
))
collector.add_pairwise_preference(PairwisePreference(
    session_id="s1", participant_id="P001",
    condition_a="A", condition_b="C", preferred="C",
))

print(collector.summary_by_condition()) # mean Likert per condition
print(collector.pairwise_win_rates())   # for ELO updates

collector.save("results/questionnaires.json")
collector.to_csv("results/questionnaires.csv")

```

22.4 `export.py` — Three Output Formats

```
exporter = ResultExporter()
exporter.to_json(results, "results/eval.json")
exporter.to_csv(results, "results/eval.csv")
exporter.to_markdown(results, "results/eval.md")
table = exporter.to_rich_table(results)  # for terminal display
stats = exporter.get_summary_stats(results)
```

The CLI command `omnillm export --format csv --input ... -o ...` is just a thin wrapper around these methods.

PART IV — PEPPER, CHOREGRAPHE, NAOqi

This part is the practical guide to the robot side of the project. We start with the hardware, work up through NAOqi the operating system, the Choregraphe IDE, the Python 2.7 problem and its solution, and end with five walk-throughs you can actually run.


Chapter 23 — Meet Pepper — The Hardware Inside the Plastic Shell

**** AT A GLANCE.** Pepper is a 120 cm, 28 kg humanoid robot from SoftBank Robotics. 20 degrees of freedom. Intel Atom CPU, 4 GB RAM. Four microphones in the head. Two ear speakers. A 10.1" chest tablet. Eye LEDs. Three omnidirectional wheels at the base. ~10 hours of battery. Knowing this hardware shapes how you think about the AI server's job.*

23.1 Physical Specs

Attribute	Value
Height	120 cm
Weight	28 kg
Degrees of freedom	20 (head 2, each arm/hand 6, hip 2, knee 1, base wheels 3)
Battery	30 Ah / 795 Wh lithium-ion (~8–10 hours active)
Onboard CPU	Intel Atom E3845 quad-core @ 1.91 GHz
Onboard RAM	4 GB DDR3
Onboard storage	8 GB flash + microSD slot
Operating system	NAOqi OS (modified Gentoo Linux)
Onboard Python	Python 2.7

23.2 Sensors



The cameras and depth sensor are not used by OmniLLM in the current design, but they are available — see "Future Work" in Chapter 39.

23.3 The Tablet

The chest tablet is a separate Android computer (1.3 GHz quad-core ARM Cortex-A7, 1 GB RAM, 32 GB storage) communicating with the head computer over an internal network at IP 198.18.0.1. From OmniLLM's perspective it is just a service: `ALTabletService`. You can `loadUrl()` to display web content, `showImage()` to display an image, or `executeJS()` to run JavaScript in the browser.

For the Embodied LLM Arena, the tablet is mostly unused (most lab tasks do not need a screen), but it is available for showing maps during navigation tasks.

23.4 The Speaker / Microphone Pair

Pepper has **four microphones** in its head, and the NAOqi audio device exposes all four channels at 48 kHz, or **a single mixed-down channel at 16 kHz**. OmniLLM uses the 16 kHz mono channel — the same format Whisper expects.

The "front" channel (channel 3 in the four-channel layout) is the most useful for one-on-one conversation. `ALAudioDevice`'s `setClientPreferences` call in `naoqi_client.py` selects this:

```
self._audio_device.setClientPreferences(
    "OmniLLMCapture",
    16000,      # sample rate
    3,         # channel: front
    0,         # deinterleaved: no
)
```

23.5 Eye LEDs as a Communication Channel

The eye LEDs are addressable RGB LEDs. NAOqi exposes them through ALLeds:

```
leds.fadeRGB("FaceLeds", r, g, b, fade_duration_seconds)
```

OmniLLM uses the eye colour to communicate **mode**:

Colour	Hex	Mode
Friendly green	#00FF88	Greeting, acknowledgement, social
Calm blue	#00AAFF	Navigation guidance
Default blue	#44AAFF	Idle / neutral
White	#FFFFFF	Showing tablet content
Yellow	#FFFF00	Thinking
Red-orange	#FF4400	Confused / error
Warm orange	#FF8800	Goodbye

This is *not* arbitrary aesthetic. Eye colour is a documented HRI signal that participants register subconsciously. Switching from blue (navigation) to green (success) reinforces the spoken response.

23.6 Why ALAnimatedSpeech, Not Plain ALTextToSpeech

NAOqi has two text-to-speech services:

- **ALTextToSpeech** — voice only. The robot is rigid while speaking.
- **ALAnimatedSpeech** — voice + automatic body gestures synchronised to the speech.

OmniLLM uses **ALAnimatedSpeech** *always* because:

1. Embodied perception research (Bartneck, Andrist, etc.) consistently shows that "talking head" robots are rated lower on naturalness and intelligence.
2. Pepper has joints; not using them is wasteful.
3. The cost is zero — **ALAnimatedSpeech** is built-in.

The configuration that produces sensible motion is:

```
config = {"bodyLanguageMode": "contextual"}
animated_speech.say(text, config)
```

bodyLanguageMode accepts: "contextual" (gestures match speech content, **recommended**), "random" (random gestures, looks unhinged), or "disabled" (back to talking-head mode).

23.7 Hardware End-of-Life Note


Aldebaran (the company behind Pepper and NAO) filed for bankruptcy in February 2025. Maxvision Technology (Shenzhen) acquired the IP in July 2025. **No new units are being manufactured.** Existing units continue to work; spare parts are increasingly hard to source. This is one reason OmniLLM is designed to be platform-portable — the abstract `RobotBridge` can target NAO, Buddy, or any future robot you point it at.

Chapter 24 — NAOqi 101 — The Operating System That Runs on Pepper

*** AT A GLANCE. *NAOqi is the middleware that makes Pepper a robot rather than a Linux box on wheels. It is a service broker on TCP port 9559 that exposes named services (ALAnimatedSpeech, ALMotion, etc.) to any client that connects with the right credentials. Its Python binding is Python 2.7 only**.*

24.1 What NAOqi Is

Imagine the robot as a Linux server, and NAOqi as the daemon process that runs on top of Linux providing all the high-level robot abstractions.



24.2 The Service Broker Model

NAOqi is built around a **service broker**. Every NAOqi service registers itself with the broker on startup; clients look up services by name and get a proxy object that can call methods on them. Services can run **locally** (in the same process — fast, zero-copy) or **remotely** (over TCP — slower, fully serialised).

Two ways to call a service from Python:

```
# Modern (NAOqi 2.x)
import qi
session = qi.Session()
session.connect("tcp://192.168.1.100:9559")
tts = session.service("ALAnimatedSpeech")
tts.say("Hello!")

# Legacy (NAOqi 1.x)
from naoqi import ALProxy
tts = ALProxy("ALAnimatedSpeech", "192.168.1.100", 9559)
tts.say("Hello!")
```

OmniLLM's `naoqi_client.py` tries `qi` first, falls back to `naoqi`, so it works against both NAOqi 1.x and 2.x.

24.3 The Twelve Services You Need to Know

Service	What it does
<code>ALMotion</code>	Joint control. <code>wakeUp()</code> enables motors; <code>rest()</code> disables. <code>setAngles()</code> moves a joint. <code>moveTo()</code> walks.
<code>ALAnimatedSpeech</code>	TTS with body gestures. <code>say(text, config)</code> .
<code>ALTextToSpeech</code>	TTS without gestures. Use when you need a stationary robot.
<code>ALAudioDevice</code>	Microphone capture. <code>setClientPreferences()</code> + <code>subscribe()</code> .
<code>ALAudioRecorder</code>	Record audio to a file. Simpler than callback-based capture.
<code>ALLeds</code>	LED control. <code>fadeRGB("FaceLeds", r, g, b, duration)</code> .
<code>ALBehaviorManager</code>	Run installed animations. <code>runBehavior("path/to/behavior")</code> . <code>isBehaviorInstalled()</code> .
<code>ALMemory</code>	Key-value store / event bus. <code>getData(key)</code> , <code>subscriber(event)</code> .
<code>ALFaceDetection</code>	Vision: detect faces. Subscribe to the <code>"FaceDetected"</code> event.
<code>ALSpeechRecognition</code>	On-board speech recognition (limited vocabulary). Mostly inadequate; OmniLLM uses Whisper instead.
<code>ALTabletService</code>	Chest tablet display. <code>loadUrl()</code> , <code>showImage()</code> .
<code>ALRobotPosture</code>	Whole-body posture. <code>goToPosture("Stand", 0.8)</code> .

24.4 The Python 2.7 Constraint Explained

The NAOqi Python binding (`pynaoqi`) is shipped as a platform-specific archive on SoftBank's developer portal. It is a CPython extension that depends on the binary layout of Python 2.7 specifically. There is no port to Python 3 that exposes the full service surface.

A community-built `qi 3.1.5` package (`pip install qi==3.1.5`) exists for Python 3 on Linux `x86_64` but several services are broken (touch detection, audio callbacks, certain event subscriptions). It is unsuitable

for production.

The conclusion: **NAOqi requires Python 2.7**. Modern AI libraries require Python 3.11+. They cannot live in the same process. **Hence the two-process bridge** of Chapter 26.

24.5 Lifecycle Quirks

Two NAOqi quirks that will trip you up:

Stiffness must be enabled before any movement. A fresh-booted Pepper has zero stiffness — its motors are dead weight. You must call `motion.wakeUp()` (or `motion.setStiffnesses("Body", 1.0)`) first. `ALAnimatedSpeech` will speak without stiffness, but the body language will not animate.

ALAudioRecorder and ALSpeechRecognition cannot share the microphone. They both subscribe to the audio device exclusively. You must `unsubscribe` one before using the other. OmniLLM uses neither directly — it captures via `ALAudioDevice` and sends to Whisper.

24.6 Connection from your PC

To connect from your PC, you need:

1. **Network access.** Same Wi-Fi as Pepper, or wired to the same LAN.
2. **Pepper's IP address.** Press the chest button once and Pepper says it ("My IP address is 192.168.1.100").
3. **The pynaoqi SDK.** Downloaded from SoftBank's developer portal, installed at `C:\pynaoqi\pynaoqi-python2.7-2.5.5.5-win32-vs2013\`.
4. **Python 2.7** at `C:\Python27\python.exe`.
5. **The right PYTHONPATH:**

```
set PYTHONPATH=C:\pynaoqi\pynaoqi-python2.7-2.5.5.5-win32-vs2013\lib;%PYTHONPATH%
set PATH=C:\pynaoqi\pynaoqi-python2.7-2.5.5.5-win32-vs2013\lib;%PATH%
```


Then `import naoqi` or `import qi` will work.

Chapter 25 — Choregraphe — The Visual Programming Studio

*** AT A GLANCE. ** *Choregraphe is SoftBank's desktop IDE for Pepper / NAO. Box-and-wire visual programming + Python script editor + 3D virtual robot simulator. Critical for OmniLLM in three ways: (1) testing behaviours without a physical robot, (2) installing custom animations the gesture planner will trigger, and (3) live-monitoring during experiments.*

25.1 What Choregraphe Is

Choregraphe is a four-panel desktop application:



It is locked to NAOqi 2.5 — version 2.5.5.5 or 2.5.10/11 for Pepper. Newer NAOqi 2.9 (Android-based) does **not** support Choregraphe.

25.2 When to Use Choregraphe Versus OmniLLM

Choregraphe and OmniLLM serve different purposes — they are friends, not substitutes:

You want to...	Use Choregraphe	Use OmniLLM
Test if Pepper's speech works	✓	—
Test a single gesture animation	✓	—
Build / tune a custom animation	✓	—
Build a <i>scripted</i> interaction	✓	—
Build an AI-driven conversation	—	✓
Use multiple LLMs as backends	—	✓
Run a controlled HRI experiment	—	✓

Live-monitor during an experiment	✓ (alongside)	✓
-----------------------------------	---------------	---

The typical workflow combines them:

1. **Plan the gesture vocabulary** in Choregraphe. Drag and edit animations until they look natural.
2. **Install the custom behaviours** on Pepper (Choregraphe → File → Build Application Package, then upload).
3. **Add the new gesture names** to OmniLLM's `GESTURE_TO_BEHAVIOR` dict in `naoqi_client.py` and `gesture_planner.py`.
4. **Run the OmniLLM experiment**, leaving Choregraphe open as a monitor.

25.3 Connecting Choregraphe to a Robot

Virtual robot (no hardware needed):


1. Open Choregraphe.
2. **Connection** → **Connect to virtual robot**.
3. Pepper appears in the 3D view; the bottom-left status shows `Connected to localhost:9559 (virtual)`.

Physical robot:

1. Make sure your PC and Pepper are on the same Wi-Fi.
2. Press Pepper's chest button → it says its IP.
3. **Connection** → **Connect to...**, enter the IP, port 9559.
4. The 3D view now mirrors the real robot's joint positions.

25.4 Box-and-Wire Programming

Each Choregraphe **box** is a small piece of behaviour. Boxes have input and output **bangs** (signal triggers). You drag boxes onto the flow diagram and connect their bangs to define the order:



Inside each box is a Python 2.7 script with `onLoad()`, `onUnload()`, `onInput_onStart()`, and `onInput_onStop()` lifecycle methods. You can inspect and edit any box's script.

25.5 Useful Boxes for OmniLLM Work

When testing Pepper for OmniLLM, the boxes you'll reach for are:

- **Speech / Animated Say** — sanity check that `ALAnimatedSpeech` works.

- [Movement / Animations / Gestures / Hey_1](#) — wave hello.
- [Movement / Animations / Gestures / Explain_8](#) — point left (OmniLLM's `point_left`).
- [Movement / Animations / Gestures / Explain_7](#) — point right.
- [LEDs / Set LEDs](#) — change eye colour.
- [Movement / Postures / Stand](#) — wake up posture.

25.6 The Script Editor as a Quick Test Bench

You don't need to drag boxes for everything. Press `Alt+5` to open the Script Editor at the bottom and just type Python:

```
# Test ALAnimatedSpeech directly
tts = ALProxy("ALAnimatedSpeech", "localhost", 9559)
config = {"bodyLanguageMode": "contextual"}
tts.say("This is what an OmniLLM response sounds like.", config)

# Test a gesture
behavior = ALProxy("ALBehaviorManager", "localhost", 9559)
behavior.runBehavior("animations/Stand/Gestures/Explain_8") # point_left

# Test eye LEDs
leds = ALProxy("ALLeds", "localhost", 9559)
leds.fadeRGB("FaceLeds", 0.0, 0.67, 1.0, 0.5) # navigation blue
```

This is the fastest way to verify everything OmniLLM will need.

25.7 Installing a Custom Behaviour

If you build a new animation in Choregraphe and want OmniLLM to trigger it:

1. **In Choregraphe:** design the animation in the Timeline Editor, save the project as `MyBehaviors/WelcomeDance`.
2. **Upload** to Pepper: `File → Upload to robot...`. The behaviour now lives at `mybehaviors/WelcomeDance` on the robot.
3. **In OmniLLM:** edit `omnillm/server/naoqi_client.py`:

```
GESTURE_TO_BEHAVIOR = {
    # ... existing entries ...
    "welcome_dance": "mybehaviors/WelcomeDance", # ← add this
}
```

4. **Trigger it:** edit `omnillm/robotics/gesture_planner.py` to map appropriate response text or task type → `"welcome_dance"`.

Chapter 26 — The Python 2.7 / Python 3.x Bridge Problem (and the Solution)

***** AT A GLANCE.**** This chapter is the canonical explanation of the single hardest engineering problem in the project. NAOqi requires Python 2.7. Modern AI libs require Python 3.11+. They cannot share a process. The solution: two processes, one HTTP boundary.

26.1 The Constraint

```
+-----+
| NAOqi SDK (pynaoqi) |
| works ONLY with Python 2.7 |
+-----+
|
+-----+
| langgraph, litellm, chromadb, |
| langchain, sentence-transformers, |
| openai-whisper |
| require Python 3.11+ |
+-----+
```

You cannot install both into the same Python interpreter. There is no "compat layer". The solution is **process separation**.

26.2 The Bridge Pattern

```
+-----+
| Process 1: NAOqi client (Python 2.7) |
| • imports qi / naoqi |
| • talks to Pepper hardware |
| • does HTTP to Process 2 |
| • zero AI dependencies |
+-----+
|
| HTTP / JSON
|
+-----+
| Process 2: AI server (Python 3.11+) |
| • imports langgraph, litellm, chromadb, |
| langchain, openai-whisper |
| • talks to LLM providers |
| • cannot import NAOqi |
+-----+
```

Each process imports only what it can. The HTTP boundary is the *only* contact between them.

26.3 Five Bridge Patterns From the Literature

The compass research (`compass_artifact_*.md`) summarises five patterns used across 15+ published Pepper-LLM projects. OmniLLM uses Pattern 1:

Pattern	Used in	Trade-offs
---------	---------	------------

1. HTTP / REST bridge (<i>OmniLLM</i>)	ilabsweden/pepperchat (2023), Frontiers ASD therapy, 6+ others	Easiest to debug, well-understood, ~50–200 ms overhead
2. Socket-based	Pepper-GPT (Auckland), Ghent University elder care	Lower latency (~10–50 ms), more code
3. ROS2 bridge (<code>naoqi_driver2</code>)	Multi-robot research projects	High setup complexity, powerful for fleets
4. MQTT broker	LAIR-GPT (Ancona)	Good when many components publish/subscribe
5. Python 3 <code>qi 3.1.5</code>	Prototypes	Single process but several services broken

OmniLLM picked Pattern 1 because it is the most-tested in the literature, the easiest to debug (every message is just `curl`-able), and the overhead is comfortably within the latency budget.

26.4 The Failure Modes You Must Anticipate

The bridge introduces three new failure surfaces. Each has a defensive fallback:

Failure	Mitigation
AI server crashed	NAOqi client times out, says "I could not connect to my AI brain" — does not crash
Network partition	Same as above; the <code>urlopen</code> timeout is 30 s
AI server returns malformed JSON	NAOqi client logs and falls back to silent failure
Audio capture returns empty bytes	Server can fall back to text-only mode (text field also accepted)
LangGraph not installed on AI server	<code>_fallback_interact</code> direct gateway call still works
Whisper not installed	Server returns 500 on <code>/transcribe</code> ; text-only mode still works

26.5 Latency Implications

The HTTP bridge adds a small but measurable overhead per interaction:

Step	Time on LAN
TCP/HTTP round trip (LAN)	5–30 ms
JSON serialise / deserialise	1–5 ms

Base64 encode / decode (5 s of audio at 16 kHz mono)	10–20 ms
Total bridge overhead per interaction	~20–50 ms

This is comfortably inside the 1–3 s budget. For comparison, the LLM call itself takes 500–2000 ms.

26.6 When You'd Want a Different Pattern

You'd switch off the HTTP bridge if:

- You need **sub-100 ms** end-to-end (token streaming for real-time conversation). Use the WebSocket pattern (BuddyBridge already does this).
- You're on **NAOqi 2.9 + Android**. Use QiSDK (Java/Kotlin), bypass the bridge entirely. (See "Future Work" in Chapter 39.)
- You're integrating with **ROS2 Nav2**. Use `naoqi_driver2` so the robot participates in the ROS2 message graph.

For a typical Pepper + LLM HRI study, the HTTP bridge is the sweet spot.

26.7 Anatomy of One HTTP Round Trip

When Pepper asks the AI server a question, here is what flows over the wire:

Request from `naoqi_client.py:233`:

```
POST /interact HTTP/1.1
Host: 192.168.1.50:5000
Content-Type: application/json
Content-Length: ~140000      ← about 100 KB of base64 audio

{
  "audio":          "UklGRiQ...", ← 5 s of 16 kHz WAV ~ 100 KB raw → ~135 KB b64
  "participant_id": "P001",
  "session_id":     "9c2e-abc123",
  "condition":      "C",
  "rag_enabled":    true
}
```

Response from `app.py:275`:

```
HTTP/1.1 200 OK
Content-Type: application/json

{
  "speech":        "Room 305 is on the third floor on your left.",
  "gesture":       "point_left",
  "emotion_led":   "#00AAFF",
  "metadata": {
    "task_type":    "navigation",
    "model_id":     "openai-gpt4o-mini",
    "rag_enabled":  true
  }
}
```

Chapter 27 — Five Practical Walk-throughs (with and without a real robot)

*** AT A GLANCE. ** This chapter is action-only. Five scenarios, each covering setup, run command, what to expect, and how to verify it worked. Use these as recipes.

27.1 Walk-through 1 — Server Smoke Test (no robot, no Choregraphe)

Goal: Confirm the AI server runs and answers a text question.

```
# Terminal 1
cd C:\Users\akshi\OneDrive\Desktop\OmniLLM
.venv\Scripts\activate
python -m omnillm.server.app --no-rag

# Terminal 2 (or browser)
curl http://localhost:5000/health
# expected: {"status": "ok", "version": "0.1.0"}

curl -X POST http://localhost:5000/interact \
-H "Content-Type: application/json" \
-d '{"text": "Hello, what can you do?", "participant_id": "TEST", "condition": "B"}'
# expected: JSON with speech / gesture / emotion_led
```

Success criteria: valid JSON response, no Python tracebacks.

If this fails: check `.env` API keys; verify Ollama is running for condition B; try `condition: "A"` if you have an OpenAI key.

27.2 Walk-through 2 — RAG Smoke Test

Goal: Confirm the RAG pipeline is indexing and retrieving from the knowledge base.

```
python -m omnillm.server.app # RAG enabled by default
```

You should see in the server logs:

```
INFO:omnillm.rag.pipeline:Indexed 12 chunks from lab_info.txt
INFO:omnillm.rag.pipeline:Indexed 8 chunks from faq.txt
INFO:omnillm.rag.pipeline:Knowledge base loaded - 42 total chunks
```

Then in another terminal:

```
curl -X POST http://localhost:5000/interact \
-H "Content-Type: application/json" \
-d '{"text": "What is the WiFi password?", "participant_id": "T", "condition": "A"}'
```

Success criteria: the response contains "UniGuest" and "Welcome2026!" — literally the values from `knowledge_base/lab_info.txt`. If those exact words appear, RAG retrieved correctly.

27.3 Walk-through 3 — Choregraphe Virtual Robot, Text Mode

Goal: See OmniLLM responses *animate* on a virtual Pepper without any physical hardware.

```
Setup:
Terminal 1: ollama serve
Terminal 2: python -m omnillm.server.app --model llama3-8b-local
Choregraphe: open, Connection → Connect to virtual robot
```

The virtual Pepper has no microphone, so we'll use text mode + manually play the gestures.

```
# Terminal 3: send a text question, capture the JSON
curl -X POST http://localhost:5000/interact \
-H "Content-Type: application/json" \
-d '{"text": "Where is Room 305?", "participant_id": "T", "condition": "B"}'
# returns: {"speech": "Room 305 ... your left.", "gesture": "point_left", "emotion_led": "#00AAFF", ...}
```

Then in Choregraphe's Script Editor:

```
# Re-play what OmniLLM decided, on the virtual robot
tts = ALProxy("ALAnimatedSpeech", "localhost", 9559)
behavior = ALProxy("ALBehaviorManager", "localhost", 9559)
leds = ALProxy("ALLeds", "localhost", 9559)

leds.fadeRGB("FaceLeds", 0.0, 0.67, 1.0, 0.3) # #00AAFF
behavior.post.runBehavior("animations/Stand/Gestures/Explain_8") # point_left
tts.say("Room 305 is on your left on the third floor.",
{"bodyLanguageMode": "contextual"})
```

Success criteria: PC speakers play the speech; the virtual Pepper in the 3D view points left and its eyes turn blue.

27.4 Walk-through 4 — Full Pipeline With Virtual Robot, NAOqi Client

Goal: Run the actual `naoqi_client.py` against Choregraphe's virtual robot, using text-mode requests.

```
Setup:
Terminal 1: ollama serve
Terminal 2: python -m omnillm.server.app --model llama3-8b-local
Choregraphe: Connection → Connect to virtual robot
```

```
# Terminal 3: NAOqi client (Python 2.7)
set PYTHONPATH=C:\pynaoqi\pynaoqi-python2.7-2.5.5.5-win32-vs2013\lib;%PYTHONPATH%
set PATH=C:\pynaoqi\pynaoqi-python2.7-2.5.5.5-win32-vs2013\lib;%PATH%

C:\Python27\python.exe omnillm\server\naoqi_client.py ^
--robot-ip localhost ^
--robot-port 9559 ^
--server-ip localhost ^
--participant TEST ^
--condition B
```

The client will say its greeting through Pepper's speech, then loop trying to record audio. Since the virtual robot has no real microphone, audio will be empty. You'll need to modify the client to use `_send_text("Where is Room 305?")` for testing — or hook in a synthetic audio source.

Success criteria: Pepper greets you in the 3D view; client log shows HTTP requests being sent to the AI server.

27.5 Walk-through 5 — Physical Pepper, End-to-End

Goal: A complete experimental session with a real Pepper.

Pre-flight:

- Pepper is plugged in and woken (chest LED solid green)
- Note Pepper's IP (chest button → "My IP is 192.168.1.100")
- Your PC and Pepper on the same Wi-Fi
- API keys in .env (or Ollama running for cond. B)

```
# Terminal 1: AI server, all features
cd C:\Users\akshi\OneDrive\Desktop\OmniLLM
.venv\Scripts\activate
python -m omnillm.server.app --host 0.0.0.0 --port 5000

# Terminal 2: NAOqi client
set PYTHONPATH=C:\pynaoqi\pynaoqi-python2.7-2.5.5.5-win32-vs2013\lib;%PYTHONPATH%
C:\Python27\python.exe omnillm\server\naoqi_client.py ^
--robot-ip 192.168.1.100 ^
--robot-port 9559 ^
--server-ip 192.168.1.50 ^
--server-port 5000 ^
--participant P001 ^
--condition C
```

Pepper says: *"Hello! I am Pepper, powered by OmniLLM. How can I help you today?"* The participant speaks; Pepper responds with grounded answer + gesture + LED change.

After the session:

```
# Terminal 3: download all logged interactions
curl http://localhost:5000/export > results/session_P001.json
```

Success criteria: at least one full interaction logged; the `metadata.model_id` field in the response matches the routing decision expected for the chosen condition.

Chapter 27A — External Resources, Community Insights, and Future Directions

*** AT A GLANCE. This chapter does not introduce new OmniLLM code. It reads the broader Pepper + LLM community — SoftBank's own labs, independent research teams (Auckland, Sweden, Italy, Czech Republic), open-source navigation/perception samples, deprecated-but-essential SDK mirrors, Japanese-voice synthesis options, and HRI evaluation curricula — and maps each resource to a specific problem you will hit, or have already hit, with the real robot. Treat it as the "what to read next" appendix for Part IV.

27A.1 How to use this chapter

The links collected here fall into seven recurring problem buckets that every Pepper + LLM project meets sooner or later. Each section below is organised the same way:

1. **The problem** — what concrete pain point this group of resources addresses (drawn from OmniLLM's own history, the project setup notes, and the "Pepper + Windows 11 quirks" memory).
2. **What the community has already published** — the repository or paper, what it actually contains (not just the title), and the licence/maintenance status.
3. **How to fold it into OmniLLM** — concrete file paths, integration sketches, or experiment ideas you can act on without re-architecting the project.

Where a resource is *only* useful as a reference (no code reuse), it appears in the final annotated list in §27A.10 rather than as a section of its own.

27A.2 Sibling HTTP-Bridge Projects — Cross-Validation of OmniLLM's Pattern 1

Chapter 26 chose Pattern 1 (HTTP/REST bridge between Python 2.7 NAOqi and Python 3.11+ AI) by surveying the literature. Five independent public projects use the same pattern. Reading their READMEs is the fastest way to *cross-check* OmniLLM's architecture and spot ideas it has not yet adopted.

Project	Stack	What OmniLLM can borrow
ilabsweden/pepperchat	Py2 <code>module_commandable.py</code> on robot, Py3 <code>dispatcher.py</code> external, OpenAI ChatGPT	Uses NAOqi Autonomous Life Proxy to switch focus to a dedicated <code>nao_focus</code> — prevents Pepper's built-in dialogue from talking over your LLM. OmniLLM does not currently set Autonomous Life mode; the participant may hear Pepper's stock greeting drowning the LLM response.

UoA-CARES/Pepper-GPT	Py3 "Black Box" (Whisper + GPT-3.5), Py2 "Pepper Controller" over VPN, NAOqi 2.1.4.13	Documents the <code>libboost_regex</code> install error explicitly — the same family of errors hits the Windows 11 pynaoqi install. Their pinning of NAOqi 2.1.4.13 confirms our 2.5.5.5 choice is <i>not</i> the only valid path.
UoA-CARES/pepper-demo	Choregraphe 2.5.10.7 + PyNAOqi, Pepper 1.8	Documents a hard-to-find ZLIB symlink fix (<code>libz.so.1</code> ↔ system) for Choregraphe on Linux. Save this for the day you move the AI server to a Linux box.
igor-lirussi/Dialogue-Pepper-Robot	Java AIML engine + Py2 NAOqi + separate speech-recognition service	The clean split between <i>dialogue engine</i> and <i>speech-recognition service</i> (parallel processes, robot IP as a CLI flag) is what OmniLLM already does, but they expose it as a hard architectural rule. Worth citing in Part X when you defend the pipeline.
softbankroboticstraining/pepper-chat ("Pepper Chat")	NAOqi 2.5 + Google Dialogflow v2 (JSON keypath)	Documents <i>QiChat voice-shaping commands</i> (pause, speed, pitch, emotional intonation) that OmniLLM does not yet exercise. These could be added to <code>naoqi_client.py</code> 's <code>say()</code> wrapper for affective speech without leaving Pepper's TTS.

Direct cross-validation finding. Every one of the five projects ends up with the exact same two-process topology OmniLLM uses. The *variation* is in (a) which Python 3-side LLM provider they use, and (b) whether they let Pepper's Autonomous Life keep its default behaviour or suppress it. OmniLLM is competitive on (a) — its multi-provider router covers ground none of the others touch — and *behind* on (b). The fix is one NAOqi call (`ALAutonomousLife.setState("disabled")`) inside `naoqi_client.py` after connection and before the greeting.

27A.3 NAOqi 2.5 / Python 2.7 Install Hardening — Community Recipes

The single biggest source of lost hours on Windows 11 has been the pynaoqi install path. Three community resources fill the gaps the official Aldebaran docs leave open:

(a) [AnonKour/pynaoqi](#) — A community mirror of "the Python 2.7 NAOqi SDK", explicitly published because the official Aldebaran distribution is hard to reach since the bankruptcy. The README documents the exact `.bashrc` lines to add:

```
export PYTHONPATH=${PYTHONPATH}:/installation/path/pynaoqi/
```

The repo is Linux-only (compiled `.so` files), so it is **not a substitute** for the Windows install we have at `C:\pynaoqi\pynaoqi-python2.7-2.5.5.5-win32-vs2013\`. But if a future collaborator runs the NAOqi client from a Linux box (recommended for ROS interop — see 27A.4), this is the easiest source.

(b) nlp.fi.muni.cz/trac/pepper/wiki/InstallationInstructions — The Masaryk University NLP lab's installation guide. Two pieces of information are not in the official docs:

- A `pyenv_install.sh` script that builds Python 2 locally and drops pynaoqi into its `site-packages`. This is the cleanest way to keep a Python 2 install isolated from system Python without `conda/anaconda` (which is explicitly broken with NAOqi — confirmed also by the Dialogflow Medium article).
- The Choregraphe-on-Linux dependency list: `libgl-dev libxt6 libxaw7`.

(c) [incognite-lab/Pepper-Controller](#) — Wraps ~60 NAOqi methods into a single `Pepper` class organised by domain (Language, Vision, Motorics, System). OmniLLM's `naoqi_client.py` is more focused (it is a *client*, not a *controller*) but the domain partitioning is a better organisation than our current flat module. **Concrete refactor idea:** split `naoqi_client.py` into `naoqi_client/language.py`, `naoqi_client/motion.py`, `naoqi_client/vision.py`, `naoqi_client/system.py`, keeping the public HTTP surface unchanged. The Pepper-Controller code is BSD-licensed and can be vendored.

Windows 11 specifics. None of the three resources above are Windows-first. The closest community-validated Windows install path is still through Maxtronics (the new IP owner) — see §27A.10 — but their KB does not yet duplicate the Aldebaran developer guide. *Conclusion:* the project setup memory's notes (loopback bug workaround, virtual robot port, Connect-to gotcha) are still the canonical Windows 11 record. Worth publishing them as a gist before they are lost.

27A.4 Navigation, Proactive Mobility, and Guided-Tour Ideas

OmniLLM's current scope (Embodied LLM Arena) treats navigation as a *spoken* directive — "Room 305 is on your left" — without Pepper actually moving. Three SoftBank-Labs repositories show what could be added without rewriting the agent:

(a) [softbankrobotics-labs/pepper-proactive-mobility](#) — Pepper autonomously approaches people in range and returns "home" when no one is present. Supports **three localisation modes**: the charging-station homing built into NAOqi, ARUCO printed-marker localisation (linoleum/vinyl ≤ 2 mm thick — *important* if you want a low-friction floor in a lab), and on-board SLAM. The codebase is 94.9 % Python, which means it is readable from the OmniLLM side rather than a black box.

(b) [softbankrobotics-labs/pepper-aruco](#) and [pepper-aruco-automapping](#) — ARUCO marker library exposed through QiSDK. Markers act as named frames the robot can localise against. For an Embodied LLM Arena task, print one marker per "room" and the agent can speak *and walk* to "Room 305".

(c) [aldebaran/naoqi_navigation_samples](#) — Choregraphe `.pml` projects for `explore`, `patrol`, and `places`. `places` is the most directly useful — it stores named locations and walks between them. Open in Choregraphe 2.5.2+, upload to Pepper, and trigger from OmniLLM via `ALBehaviorManager.runBehavior("places/goto_room_305")`, mirroring the `GESTURE_TO_BEHAVIOR` dictionary pattern in `naoqi_client.py`.

Concrete extension proposal — "OmniLLM Wayfinder" condition. Add a fourth experimental condition D to the Arena where the LLM's response includes a `target_location` field (e.g. `"room_305"`), the gesture planner maps that to a `places/*` Choregraphe behaviour, and Pepper *escorts* the participant

instead of merely pointing. Latency budget is unchanged (the walk is async); the existing eye-LED "navigation blue" already signals the mode. The ROBO-GAP findings on robot age/gender perception (§27A.8) become an additional dependent variable: does an escorting robot read as older/more authoritative than a pointing one?

On the ROS side. `ros-naoqi/naoqi_bridge` (and `naoqi_driver2` for ROS 2) republishes NAOqi services as ROS topics. The bridge is worth the setup cost the moment you want to integrate with **Nav2** or any standard ROS perception stack. The compass research already catalogued this as bridge Pattern 3; the live repository confirms ROS 1 Noetic and ROS 2 Foxy/Humble are the realistic targets.

27A.5 Multilingual Voice — The Japanese-Voice Question

OmniLLM has a language detector that recommends a model per detected language (the recent `language_code / recommended_model` compat properties). For research participants who code-switch into Japanese, the robot is currently bottlenecked at NAOqi's stock `nao_jpf` voice — a 2014-era diphone voice that sounds noticeably worse than modern Japanese TTS. Two resources clarify the realistic options:

(a) Voisona Pepper voice library (voisona.com/song/artist/pepper_ja_JP) — A *singing-voice* library trained on the same Pepper character, sold commercially through DLsite/Amazon/Sonicwire. The vocal range (A3–C#5) and tempo (80–120 BPM) make this a poor fit for conversational TTS. **Verdict: not usable for the Arena.** It is useful only if you build a "Pepper sings" demo for outreach.

(b) CeVIO (Wikipedia article) — The underlying engine behind Voisona, HTS-HMM based (same family as Sinsy / Open J-Talk). CeVIO itself is Japanese-only and Windows-only. **A better path** for conversational Japanese on Pepper is the open-source toolchain that shares CeVIO's lineage: **Open JTalk + HTS voices** (server-side, generate `.wav`, send to Pepper through `ALAudioPlayer`) or **VOICEVOX** (MIT-licensed, modern neural TTS for Japanese, runs locally on the same GPU you use for Whisper). VOICEVOX is the most practical 2026 choice and slots into the existing TTS abstraction in `omnillm/server/`.

Concrete refactor proposal. Introduce a `tts/` provider directory mirroring `omnillm/server/providers/`. The current default (NAOqi `ALAnimatedSpeech`) becomes `tts/naoqi.py`; add `tts/voicevox.py` that synthesises Japanese to a 16 kHz WAV and ships it back in the `/interact` JSON as `speech_audio_b64`. The NAOqi client plays it via `ALAudioPlayer.playFile()`. Body language is lost in this path — `ALAnimatedSpeech` requires synthesised speech to come from its own engine — so use it only when the language detector chooses Japanese.

27A.6 Vision-Side Upgrades — Face, Emotion, Mask

OmniLLM does not currently consume Pepper's two 5 MP cameras or the Xtion depth sensor (Chapter 23.2 notes this explicitly). Three mature open-source pieces would change that without leaving the existing HTTP-bridge architecture:

(a) ageitgey/face_recognition — dlib-based, 99.38 % on LFW. Python 3 only, runs comfortably on the AI-server side. Add a `/recognise_face` endpoint that accepts a JPEG from the NAOqi client (captured via `ALVideoDevice.getImageRemote()`) and returns a participant ID. Two research uses:

- **Personalised greetings** — the agent remembers a returning participant and adapts its register.
- **Automatic session-ID assignment** — `participant_id` is currently passed as a CLI flag, error-prone in a live study. Face recognition removes the source of that error.

(b) LucaCorvitto/Emotional_Pepper — Mixes Python (44.8 %), web UI (42.6 %), and **PDDL planning (12.6 %)**. The PDDL piece is the interesting one for OmniLLM: it shows a planner driving Pepper's "mood" through a structured world model rather than freeform LLM output. **Hybrid idea:** keep the LLM as the *content* generator, but gate gesture/LED selection through a small PDDL plan when the agent needs to maintain a persistent emotional arc across an interaction (a known weakness of pure LLM-driven gesture mapping). Their report PDF is worth pulling for Part X / thesis design.

(c) softbankrobotics-labs/pepper-mask-detection and pepper-deep-learning — QiSDK libraries (Kotlin/C++), targeting NAOqi 2.9 + Android tablet. They confirm what Chapter 23.7 warned: *new* perception work in the SoftBank-Labs org is **NAOqi 2.9 only**. If OmniLLM ever migrates to NAOqi 2.9, the SDK split (Kotlin on the robot, Python on the server) becomes mandatory; bridge Pattern 5 (WebSocket) is the natural choice there, not the current REST.

Additional finding from arxiv.org/abs/2401.17663 ("Social Robot Navigation with Adaptive Proxemics Based on Emotions", Bilen et al., 2024) — not a Pepper-LLM paper, but it provides experimental evidence that *proxemic distance* should change with detected emotion (angry participants want more space; happy ones don't want less). If §27A.4's "OmniLLM Wayfinder" extension is built, the same emotion signal could feed the approach-distance parameter.

27A.7 Dialogue Platforms That Came Before LLMs

Several listed resources predate the LLM era. They are *not* substitutes for OmniLLM's stack but they teach important lessons:

(a) Pepper + Dialogflow integration (Medium article by blogemtech) — Architecture: Pepper's microphone streams into a custom `SoundProcessingModule` (ALModule subclass) with silence detection, then to Dialogflow's `detectIntent` API. Three concrete takeaways:

1. **Silence detection is essential** — without it Pepper streams constantly and you exhaust API quota. OmniLLM's current `ALAudioDevice.subscribe()` loop in `naoqi_client.py` does **not** do silence detection; it relies on a fixed 5-second window. For longer interactions, port the article's amplitude-threshold peak detector (test threshold = 14 000 on 16 kHz mono).
2. **Pepper's onboard ASR is unreliable for proper nouns** — "Bob" reliably mis-recognised as "Paul" despite training. This is why OmniLLM uses Whisper. Worth citing in defence of the design choice.
3. **Anaconda is incompatible with NAOqi SDK**. Use system Python 2.7 or `pyenv`. (Matches the memory note.)

(b) TheRARELab/langex — *Not* an LLM project. A Choregraphe template for reproducible language-HRI experiments: configurable "Factor 1/2/3" conditions, dialogue in `.txt` files per condition, auto-logged filenames with the experimental parameters baked in. Three direct ideas for OmniLLM's experiment runner:

- **Filename convention** — embed condition, participant, model in the log filename so a folder listing is self-describing.
- **type 'next' advancement** — for pilot studies where you want to control utterance pacing manually before opening to real-time audio.
- **Reproducibility headers** — every log file starts with a header block recording all experimental parameters. OmniLLM already exports JSON; consider adding a metadata block matching langex's format.

(c) softbankrobotics-labs.org generally — `pepper-app-launcher`, `pepper-gamepad`, `pepper-solitarities-loop`. The "solitarities loop" (idle animations) is the resource most relevant to OmniLLM: between participant turns, Pepper currently goes still — uncanny in a long study. A small subscribed loop that triggers `pepper-solitarities` animations during idle would improve perceived liveliness without touching the agent.

27A.8 Validation Suites, Curricula, and HRI Reference Datasets

For thesis-grade validity arguments (see Part X), the following external resources let you defend the experimental design without re-justifying every choice:

(a) robocupathomeedu.org/learn/online-classroom/online-challenge-2022 — Page would not load during research, but the broader RoboCup@Home Education curriculum is the standard set of tasks (person following, object handover, spoken instruction following, navigation between rooms) that a service robot is expected to handle. Use this list as the *external rubric* against which to claim OmniLLM coverage in Part X. Of the standard tasks, OmniLLM currently covers spoken instruction following and grounded answering; navigation and person-following are unlocked by §27A.4.

(b) sunzhida/COMP4461_2017Fall_Lab4 — HKUST HCI course lab on Pepper. The exercise (`expression_teller.py`) is a one-file demo of face detection + gender + expression classification through NAOqi's on-board modules. Two uses for the thesis:

- **Baseline implementation** to compare LLM-driven emotional response against an ALMemory-event-driven baseline.
- **Pedagogical citation** when arguing OmniLLM is approachable for a graduate course (Part XI is already pitched this way).

(c) robo-gap.unisi.it (Perugia et al., 2022, HRI ACM/IEEE) — A peer-reviewed dataset of perceived age, femininity, masculinity, and gender-neutrality across 251 robots, including Pepper. The dataset is downloadable as CSV with ICC reliability between .896 and .954. **Concrete use in OmniLLM's thesis:** as a control variable. When reporting results, cite ROBO-GAP's published Pepper ratings so the reader knows *what perception baseline* the participants are walking in with. This is the kind of small move that hardens an HRI paper against reviewer pushback.

(d) **CARESSES** (caressesrobot.org) — Pepper used for *culturally competent* elder care across UK / Japan / India. The web server was unreachable during research; the publications stack (search "Bruno", "Sgorbissa", "CARESSES" on Google Scholar) is the relevant entry point. Important for Part X if you want to position OmniLLM as a cultural-HRI platform rather than a generic chatbot.

27A.9 Future Directions Unlocked by These Resources

Ranked by ease and research payoff, drawing only on what the resources above actually demonstrate:

Idea	Effort	Research payoff	Hooks into
Disable Autonomous Life on connect (kill stock dialogue)	1 line in <code>naoqi_client.py</code>	Removes a confound that all five sibling projects have already fixed	§27A.2

robots.ros.org/pepper/	High-level pointer into the <code>naoqi_bridge</code> ROS packages
doc.aldebaran.com/1-14/... and doc.aldebaran.com/2-5/...	The canonical NAOqi reference; bookmark Service Pages (<code>ALMotion</code> , <code>ALAnimatedSpeech</code> , <code>ALMemory</code> , <code>ALAudioDevice</code>)
maxtronics.com/en/support/kb/ and .../category/pepper/...	Post-bankruptcy mirror of Pepper 2.5 and 2.9 downloads (firmware, Choregraphe, SDK)
stackoverflow.com/questions/tagged/pepper	First-stop for one-off NAOqi/Choregraphe errors; <i>not</i> an answer source for architectural questions
github.com/TheRARELab/langex	Reproducible language-HRI experiment template — borrow its log-filename and header conventions
softbankroboticstraining.github.io/pepper-chat-dialogflow/	Dialogflow-based "Pepper Chat" — source of QiChat voice-shaping commands for affective TTS
github.com/softbankrobotics-labs/pepper-proactive-approach	Proactive approach + ARUCO/SLAM localisation — basis for the Wayfinder condition
github.com/aldebaran/naoqi_navigation_samples	<code>explore/patrol/places</code> Choregraphe samples — fastest way to add named-location navigation
github.com/ros-naoqi/naoqi_bridge	ROS-side wrapper of NAOqi services — entry into Nav2
groups.google.com/g/ros-sig-aldebaran	Slow but active community list — best place to ask <code>naoqi_driver2</code> questions
doc.aldebaran.com/2-5/index_dev_guide.html and .../getting_started/index.html	Official guides; sections on <code>ALMemory</code> events and <code>ALAnimatedSpeech</code> configuration are the most frequently re-read
robocupathomeedu.org/.../online-challenge-2022	External rubric of service-robot tasks — use to scope thesis claims
maxtronics.com/en/support/kb/ (NAO content)	NAO-specific KB; relevant if you ever port OmniLLM to NAO as a portability proof
github.com/orgs/aldebaran/repositories	<code>libqi</code> , <code>libqi-python</code> , <code>qibuild</code> are still maintained (Apr 2026 updates); <code>robot-jumpstarter</code> is the best Python starter
softbankroboticstraining.github.io/pepper-chat-sameas/above.html	Same as above — and a more direct deep link into the QiChat command reference

github.com/softbankrobotics-labs (org root)	Browse for <code>pepper-aruco</code> , <code>pepper-mask-detection</code> , <code>pepper-deep-learning</code> , <code>pepper-solitaires-loop</code> , <code>navigation-toolkit-demo</code>
ros-naoqi/naoqi_bridge/tree/master/naoqi_apps	Sample ROS apps that use the bridge — read these before writing your own ROS node
maxtronics.com/en/support/kb/category/pepper/	<code>pepper-2.5-software</code> downloads — the most reliable place to re-fetch Choregraphe and the SDK
doc.aldebaran.com/2-5/getting_started/index.html	The "first 10 minutes with Pepper" reference
github.com/incognite-lab/Pepper-Controller	Python-class abstraction over NAOqi 2.5 — model for refactoring <code>naoqi_client.py</code>
github.com/lillypiri/pepper	Beginner Choregraphe exercises — useful for onboarding collaborators new to box-and-wire
nlp.fi.muni.cz/trac/pepper/wiki/InstallationInstructions	Linux install path with a custom <code>pyenv_install.sh</code> — best alternative to Anaconda
robohack.org.uk/viewforum.php?f=6	Small UK robotics-hobbyist forum — occasional Pepper troubleshooting threads, low signal
github.com/sunzhida/COMP4461_2017Fall_Lab4	HKUST course lab — <code>expression_teller.py</code> is a 50-line baseline for emotion-aware response
zijianhu.com/post/nao-tutorial/installation/	Tutorial-style NAO install — Linux focus, complements §27A.3
wiki.ros.org/nao/Tutorials/Installation	Official ROS NAO install — pre-Pepper but the bring-up patterns transfer
arxiv.org/abs/2401.17663	Emotion-adaptive proxemics paper (Bilen et al. 2024) — empirical basis for tying detected emotion to approach distance in the Wayfinder extension
github.com/ageitgey/face_recognition	dlib-based face recognition; the practical choice for a <code>/recognise_face</code> endpoint
github.com/LucaCorvitto/Emotional_Pepper	PDDL-planned emotional behaviour over a tablet UI — hybrid LLM+planner inspiration
github.com/AnonKour/pynaoqi	Linux mirror of the Python 2.7 NAOqi SDK — keep the URL in case Aldebaran's mirror disappears

robotlab.com/support/factory-reset-on-pepper-robot	Authoritative factory-reset SOP (PuTTY/FileZilla/nao-autoflash) — <i>do not</i> run this without a current backup of /home/nao
github.com/UoA-CARES/pepper-demo	ZLIB symlink fix for Choregraphe on Linux — save for later
github.com/UoA-CARES/Pepper-GPT	Closest sibling to OmniLLM in spirit — different LLM strategy (single GPT-3.5), worth full code-walk
github.com/igor-lirussi/Dialogue-Pepper-Robot	Java AIML + Py2 NAOqi split — the architectural ancestor of OmniLLM
github.com/ilabsweden/pepperchat	The original (2023) public HTTP-bridge OpenAI Pepper project — cite as prior art

27A.11 What Is Deliberately Not Here

Three categories of resource the user originally listed but that proved either inaccessible or off-topic on closer reading:

- caressesrobot.org — server unreachable during research; the *published papers* (Sgorbissa, Bruno, et al.) are the productive citation target rather than the website.
- arxiv.org/abs/2401.17663 — not an LLM-on-Pepper paper. It is a navigation/proxemics paper that nonetheless informs §27A.6.
- voisona.com and CeVIO — commercial, singing-voice, Japanese- only; the practical Japanese-TTS path for OmniLLM is VOICEVOX, not these. Listed because the user provided them; reasoning recorded so the choice does not have to be re-litigated later.

PART V — SETUP & RUNNING

This part is action-oriented. By the end, you will have a clean install, running tests, a server you can talk to with `curl`, and (optionally) a working robot connection.

Chapter 28 — Installing OmniLLM From Zero

***** AT A GLANCE.**** Six steps: install Python 3.11+, install Git, clone the repo, create a venv, pip install -e .[all], configure .env. Total time: 15–20 minutes on a clean machine.

28.1 Prerequisite Software

Install these once per machine. None of them are OmniLLM-specific.

Tool	Why	Where
Python 3.11+	Run the AI stack	python.org / brew / apt
Git	Clone the repo, pull updates	git-scm.com
A code editor	Read / edit code (VS Code recommended)	code.visualstudio.com
Ollama (optional)	Free local LLMs	ollama.com
Choregraphe 2.5 (robot only)	Talk to Pepper visually	softbankrobotics.com (developer portal)
Python 2.7 (robot only)	NAOqi client	python.org/downloads/release/python-2718
pynaoqi SDK (robot only)	Python 2.7 bindings to NAOqi	softbankrobotics.com (developer portal)

28.2 The Six Setup Steps

```
# -- 1. Clone -----
git clone https://github.com/Akshita-sr/OmniLLM.git
cd OmniLLM

# -- 2. Create a virtual environment -----
python -m venv .venv
# Activate:
# Linux/macOS: source .venv/bin/activate
# Windows CMD: .venv\Scripts\activate.bat
# PowerShell: .venv\Scripts\Activate.ps1
# Git Bash: source .venv/Scripts/activate

# -- 3. Install OmniLLM with all extras -----
pip install -e ".[all]"

# -- 4. Configure API keys -----
cp .env.example .env # Linux/macOS
copy .env.example .env # Windows CMD
# Then open .env in your editor and paste your real keys

# -- 5. (Optional) Install Whisper for local STT -----
pip install openai-whisper

# -- 6. (Optional) Pull Ollama models for free local LLMs ---
ollama pull llama3:8b
ollama pull qwen2.5:7b
```

28.3 Verify the Install

```
# Should print version + register all 19 models in a table
omnillm --version
omnillm models
```

If both work, you have a healthy base install.

```
# Run the tests (no API keys needed – uses mocks)
pytest tests/ -v
```

You should see 278 tests pass. Failures here usually mean a missing dependency.

28.4 The `.env` File

`.env.example` lists every supported environment variable. You only need the ones for providers you actually want to use:

```
# OpenAI (GPT-4o, GPT-5.4, o1, o3-mini, GPT-4o-mini)
OPENAI_API_KEY=sk-proj-...

# Anthropic (Claude Sonnet/Opus/Haiku)
ANTHROPIC_API_KEY=sk-ant-...

# Google (Gemini 2.5 Pro, Flash, 2.0 Flash)
GOOGLE_API_KEY=AIzaSy...

# DeepSeek
DEEPSEEK_API_KEY=...

# Qwen / DashScope
QWEN_API_KEY=...

# Optional – override defaults
OMNILLM_DEFAULT_MODEL=openai-gpt4o-mini
OMNILLM_KNOWLEDGE_BASE=knowledge_base/
```

*(Tip) **No keys at all?** You can still run the project. Use Ollama models exclusively (`-m llama3-8b-local`). The tests work without any keys because they mock LLM calls.*

28.5 Optional Install Extras

Pyproject defines four optional extras. The `all` extra installs all of them, but if you want to be selective:

Extra	What it adds	When
[dev]	pytest, pytest-asyncio, pytest-mock	Always (you'll want tests)
[robotics]	flask, websockets	When you want to run the AI server

[hri]	chromadb, langchain-community, sentence-transformers, pypdf, langdetect, langgraph	When you want full RAG + agent graph
[all]	Everything above	Default recommendation

```

pip install -e ".[dev]"           # minimum
pip install -e ".[dev,robotics]" # + AI server
pip install -e ".[dev,hri]"      # + RAG + LangGraph
pip install -e ".[all]"          # everything

```

28.6 Common First-Run Issues

Symptom	Fix
command not found: omnillm	venv not activated, or install failed
ModuleNotFoundError: chromadb	<code>pip install -e ".[hri]"</code>
ModuleNotFoundError: flask	<code>pip install -e ".[robotics]"</code>
APIConnectionError	wrong / missing key in <code>.env</code>
Error code: 401	key has a typo or hasn't been activated
<code>litellm.exceptions.BadRequestError</code>	model name in YAML doesn't match provider's actual name
Ollama: connection refused	<code>ollama serve</code> not running

Chapter 29 — Running Without a Robot (The Default Path)

**** AT A GLANCE.** Most development happens with no robot. The Flask server accepts text via `/interact` and returns `RobotAction` JSON. You can drive it with `curl`, Postman, a Python script, or a test dashboard. This is the path you should use 80% of the time.*

29.1 Why Robot-Free Is the Default

You should not start the robot when:

- Your changes are in the AI layer (LLM logic, RAG, evaluation).
- You are debugging the agent graph.
- You are running benchmarks (`omnillm evaluate`).
- You are adding a new model to the registry.
- You are running tests.
- The robot is busy with another study.
- The robot is not in the room.
- You want to iterate quickly.

The robot is for **embodied user studies**. For everything else, plain HTTP requests are faster, cheaper, and give you better debug logs.

29.2 Starting the Server

```
# Activate venv (Linux / macOS)
source .venv/bin/activate

# Start with defaults
python -m omnillm.server.app

# Or with custom options
python -m omnillm.server.app --host 0.0.0.0 --port 5000 --debug
```

You will see:

```
INFO:omnillm.rag.pipeline:Indexed 12 chunks from lab_info.txt
INFO:omnillm.rag.pipeline:Indexed 8 chunks from faq.txt
INFO:omnillm.rag.pipeline:Knowledge base loaded — 42 total chunks
* Running on all addresses (0.0.0.0)
* Running on http://127.0.0.1:5000
```

29.3 Driving It With `curl`

The single most useful command for testing:

```
curl -X POST http://localhost:5000/interact \
-H "Content-Type: application/json" \
-d '{
  "text": "What is the WiFi password?",
  "participant_id": "TEST",
  "session_id": "smoke-1",
```

```
"condition": "A",
"rag_enabled": true
}]'
```

You should get back something like:

```
{
  "speech": "The visitor WiFi network is 'UniGuest' with the password 'Welcome2026!'.",
  "gesture": "nod",
  "emotion_led": "#00FF88",
  "metadata": {
    "task_type": "info_retrieval",
    "model_id": "openai-gpt4o-mini",
    "rag_enabled": true
  }
}
```

29.4 Driving It From Python

A useful test script:

```
# test_interact.py
import requests

def ask(text, condition="A"):
    r = requests.post("http://localhost:5000/interact", json={
        "text": text, "participant_id": "T",
        "session_id": "test", "condition": condition,
    })
    return r.json()

print(ask("What time does the lab open?"))
print(ask("Where is Room 305?"))
print(ask("How are you today, Pepper?"))
print(ask("Bonjour, comment allez-vous?")) # multilingual
```

Run it with `python test_interact.py`. Each call exercises a different task type and routing branch.

29.5 Driving It From the CLI

For LLM evaluation (no robot involved at all):

```
omnillm models                # list registered models
omnillm ask "Explain RAG simply" --all # query every model
omnillm evaluate -m openai-gpt4o-mini # run benchmark on this model
omnillm council "Is P=NP?"          # 3-model consensus
omnillm route "Cheap question" --budget 0.001
omnillm leaderboard
omnillm costs
```

These commands talk to providers directly through `gateway.py` — they do not use the Flask server at all. Useful for evaluating models in isolation.

29.6 Driving the Agent Graph From a Python REPL

Sometimes you want to invoke the LangGraph pipeline directly from Python, without the HTTP layer:

```
import asyncio
from omnillm.gateway import LLMGateway
from omnillm.rag import RAGPipeline
```

```

from omnillm.hri import build_hri_graph

async def main():
    gw = LLMGateway()
    rag = RAGPipeline(gateway=gw, model_id="openai-gpt4o-mini")
    rag.index_directory("knowledge_base/")
    graph = build_hri_graph(gateway=gw, rag=rag)

    state = {
        "utterance": "What is the WiFi password?",
        "participant_id": "P001",
        "session_id": "test-1",
        "condition": "C",
        "rag_enabled": True,
    }
    result = await graph.ainvoke(state)
    print(result["response_text"])
    print(result["robot_action"])

asyncio.run(main())

```

This is the fastest way to debug the graph without HTTP overhead.

Chapter 30 — Running With Choregraphe's Virtual Robot

***** AT A GLANCE.**** *Choregraphe ships with a virtual Pepper. Connect to it with `Connection` → `Connect to virtual robot` and you can see speech, gestures, and LED changes in the 3D viewport. The virtual robot has no real microphone, so audio mode does not work — use text mode.*

30.1 Why the Virtual Robot Helps

When the physical robot is unavailable (or not yet bought), Choregraphe's virtual robot lets you:

- See what gestures look like.
- Verify NAOqi services are reachable on `localhost:9559`.
- Test the `naoqi_client.py` end-to-end.
- Demo the system to others without hardware.

30.2 Limitations of the Virtual Robot

Feature	Virtual	Real
ALAnimatedSpeech (TTS)	✓ via PC speakers	✓ via Pepper speakers
ALMotion gestures	✓ in 3D view	✓ on the robot
ALLeds	✓ in 3D view	✓ real LEDs
ALBehaviorManager (pre-installed behaviours)	✓	✓
ALAudioDevice (microphone)	✗	✓
ALTabletService	partial	✓
ALFaceDetection	✗ (no camera)	✓

The biggest limitation is the missing microphone. For OmniLLM testing, that means **send text instead of audio**.

30.3 Three-Terminal Setup

```
# Terminal 1: Ollama (optional, for local models)
ollama serve

# Terminal 2: AI server
cd C:\Users\akshi\OneDrive\Desktop\OmniLLM
.venv\Scripts\activate
python -m omnillm.server.app --model llama3-8b-local
```

Choregraphe (GUI):

- Open Choregraphe 2.5
- **Menu:** `Connection` → `Connect to virtual robot`

- Wait **for** the robot to appear in the 3D viewport
- Status bar **shows**: "Connected to localhost:9559 (virtual)"

```
# Terminal 3: send a text question
curl -X POST http://localhost:5000/interact \
-H "Content-Type: application/json" \
-d '{"text":"Where is Room 305?","participant_id":"T","condition":"B"}'
```

Then in Choregraphe's Script Editor (Alt+5), play back the response to see the virtual Pepper move:

```
# Use the values returned by curl
tts = ALProxy("ALAnimatedSpeech", "localhost", 9559)
behavior = ALProxy("ALBehaviorManager", "localhost", 9559)
leds = ALProxy("ALLeds", "localhost", 9559)

leds.fadeRGB("FaceLeds", 0.0, 0.67, 1.0, 0.3) # #00AAFF
behavior.post.runBehavior("animations/Stand/Gestures/Explain_8") # point_left
tts.say("Room 305 is on your left on the third floor.",
{"bodyLanguageMode": "contextual"})
```

30.4 Running the NAOqi Client Against the Virtual Robot

Once the virtual robot is connected in Choregraphe, you can also run the full `naoqi_client.py` against it. The client will not be able to capture real audio, but it will exercise everything else (HTTP, action dispatch, NAOqi calls):

```
set PYTHONPATH=C:\pynaoqi\pynaoqi-python2.7-2.5.5.5-win32-vs2013\lib;%PYTHONPATH%
set PATH=C:\pynaoqi\pynaoqi-python2.7-2.5.5.5-win32-vs2013\lib;%PATH%

C:\Python27\python.exe omnillm\server\ naoqi_client.py ^
--robot-ip localhost ^
--robot-port 9559 ^
--server-ip localhost ^
--server-port 5000 ^
--participant TEST ^
--condition B
```

30.5 Useful Quick-Tests in Choregraphe

These let you verify the underlying NAOqi services are healthy *before* worrying about OmniLLM:

```
# Test ALAnimatedSpeech
ALProxy("ALAnimatedSpeech", "localhost", 9559).say("OmniLLM ready.")

# Test ALMotion (must wakeUp first or you'll get a stiffness error)
m = ALProxy("ALMotion", "localhost", 9559)
m.wakeUp()
m.setStiffnesses("Body", 1.0)

# Test all of OmniLLM's gestures one by one
b = ALProxy("ALBehaviorManager", "localhost", 9559)
for behavior in [
    "animations/Stand/Gestures/Hey_1", # wave
    "animations/Stand/Gestures/Explain_8", # point_left
    "animations/Stand/Gestures/Explain_7", # point_right
    "animations/Stand/Gestures/Explain_1", # point_forward
    "animations/Stand/Emotions/Positive/Enthusiastic_1", # nod
    "animations/Stand/Emotions/Neutral/Thinking_1", # think
    "animations/Stand/Emotions/Negative/Confused_1", # confused
]:
    if b.isBehaviorInstalled(behavior):
        b.runBehavior(behavior)
```

```
time.sleep(2)
```

Chapter 31 — Running With a Physical Pepper

***** AT A GLANCE.** *** The full real-robot run. Connect to Pepper over Wi-Fi, run the AI server on your laptop, run the NAOqi client on Pepper (or your laptop with the SDK). Pepper greets, listens, replies with speech + gesture + LED.*

31.1 Pre-Flight Checklist

Before connecting, check all of these:

- [] Pepper is powered on and showing a solid green chest LED.
- [] Pepper is plugged into a charger if the battery is below ~30%.
- [] Your PC and Pepper are on the **same Wi-Fi** (not a hotspot vs a LAN, not a cellular network).
- [] You can ping Pepper from your PC: `ping 192.168.1.100`.
- [] You can browse to Pepper's diagnostic page in a browser: `http://<pepper-ip>` (sometimes shows a status page).
- [] Your `.env` has the API keys for the conditions you'll use, **or** Ollama is running for condition B.

31.2 Find Pepper's IP

Press the **chest button** once. Pepper says: *"My IP address is one nine two dot one six eight dot one dot one zero zero. My battery is at eighty-two percent."*

Write down the IP. It probably won't change (DHCP usually re-leases the same address) but it can.

31.3 Wake Pepper Up

A freshly-booted Pepper has zero stiffness — its motors are dead weight. You must wake it before it can move:

From Choregraphe: right-click the robot in the 3D view → **Wake Up**.

From a Python 2.7 one-liner:

```
C:\Python27\python.exe -c ^
"from naoqi import ALProxy; ALProxy('ALMotion','192.168.1.100',9559).wakeUp()"
```

After waking up, Pepper will hold itself upright and the head will move. This is normal.

31.4 Three-Terminal Run

```
# Terminal 1 (your PC): AI server
cd C:\Users\akshi\OneDrive\Desktop\OmniLLM
.venv\Scripts\activate
python -m omnillm.server.app --host 0.0.0.0 --port 5000

# Terminal 2 (your PC): find your local IP
ipconfig # Windows – look for the 192.168.x.x of your Wi-Fi adapter

# Terminal 3 (your PC, Python 2.7): NAOqi client
set PYTHONPATH=C:\pynaoqi\pynaoqi-python2.7-2.5.5.5-win32-vs2013\lib;%PYTHONPATH%
set PATH=C:\pynaoqi\pynaoqi-python2.7-2.5.5.5-win32-vs2013\lib;%PATH%

C:\Python27\python.exe omnillm\server\naoqi_client.py ^
--robot-ip 192.168.1.100 ^
--robot-port 9559 ^
--server-ip 192.168.1.50 ^
--server-port 5000 ^
--participant P001 ^
--condition C
```

Replace `192.168.1.100` with Pepper's IP and `192.168.1.50` with your PC's IP. Pepper now greets:

"Hello! I am Pepper, powered by OmniLLM. How can I help you today?"

31.5 Speaking to Pepper

Stand within ~1 metre of Pepper, facing the head microphones. Speak clearly. The current `_record_audio` is a 5-second placeholder; you'll need to replace it with `ALAudioRecorder` or callback-based capture for production use (see Chapter 21.8). For now, use the text-mode workaround:

```
# In naoqi_client.py, _interaction_loop, replace:
audio_bytes = self._record_audio(duration_seconds=5)
response = self._send_audio(audio_bytes)
# With:
text = raw_input("You: ") # Python 2 input
response = self._send_text(text)
```

31.6 Live Monitoring With Choregraphe

Open Choregraphe and connect to the **same** Pepper IP. The 3D view will mirror what the real robot is doing — useful for debugging. Choregraphe's Log Viewer shows NAOqi messages in real time, including any service errors.

(!) **Two clients can connect to one Pepper.** Choregraphe's connection is read-only by default when another client is actively driving the robot — you can monitor without conflicting.

31.7 Battery Management

Pepper's battery is the limiting factor in long sessions. Tactics:

- **Plug in between sessions.** Pepper can charge while idle in rest posture.
- **Plug in during sessions** if you have a long extension. Pepper can operate while plugged in.
- **Check battery programmatically:**

```
battery = ALProxy("ALBattery", "192.168.1.100", 9559).getBatteryCharge()  
# returns 0-100
```

Below ~15% Pepper enters safe mode and disables motors. Plan accordingly.

31.8 Cleanup at End of Session

Press `ctrl+c` in the NAOqi client terminal. The `_cleanup()` method calls `motion.rest()` which lowers Pepper into a safe rest posture.

Chapter 32 — Running an Experimental Session With a Real Participant

*** AT A GLANCE. ** A complete protocol for running one participant through all five conditions of the Embodied LLM Arena. Total session time: ~30 minutes including questionnaire.

32.1 Materials

Before the participant arrives, prepare:

- [] **Consent form** (printed, two copies — one for them, one for you).
- [] **Participant ID** assigned (P001, P002, ...).
- [] **Counterbalancing sheet** showing the order in which this participant will see conditions A–E (Latin square).
- [] **Task script** — the four task prompts each condition will have to handle.
- [] **Questionnaire form** (paper or Google Forms) ready.
- [] **Pepper charged** (>50%) and woken up.
- [] **AI server running** with all extras installed.

32.2 Pre-Session Briefing (~5 minutes)

"Thank you for participating. Today you'll talk with our robot, Pepper, in a few different modes. After each mode you'll fill in a short questionnaire about your experience. The whole session is about 30 minutes. You can stop at any time without giving a reason."

"I'll explain the tasks now. You'll ask Pepper four kinds of question: first, a factual question about the lab; second, a navigation question; third, a casual conversation; fourth, the same again but in another language if you speak one. There are no right or wrong answers — we're studying Pepper, not you."

Sign the consent form. Note the participant's spoken languages.

32.3 The Per-Condition Loop

For each condition (in the participant's counterbalanced order):

```
# Terminal 3 – kill the previous client (Ctrl+C) and restart
C:\Python27\python.exe omnillm\server\naoqi_client.py ^
--robot-ip 192.168.1.100 ^
--server-ip 192.168.1.50 ^
--participant P001 ^
--condition <X> # change for each condition
```

Run the four tasks (T1–T4). The experimenter prompts the participant:

Task	Prompt to participant
------	-----------------------

T1 (Info Retrieval)	"Ask Pepper a factual question about the lab — like the WiFi password, or what time it opens."
T2 (Navigation)	"Ask Pepper to direct you somewhere — like Room 305 or the cafeteria."
T3 (Social Conversation)	"Have a casual chat with Pepper — say hello, ask how it is, ask what it thinks about something."
T4 (Multilingual)	"If you speak another language, ask any of those questions again in that language."

After all four tasks: hand the participant the questionnaire. While they fill it in, save the data:

```
# Save logs incrementally
curl http://localhost:5000/export > results/session_P001_cond<X>.json
```

32.4 The Questionnaire

Five 1–7 Likert items (InteractionQuestionnaire):

1. The robot's answers were accurate.
2. The robot was natural to talk to.
3. I trust the information the robot gave me.
4. The robot's gestures were appropriate.
5. The robot responded quickly enough.

Optional Godspeed five-subscale (1–5 each, ~5 minutes longer):

- Anthropomorphism, Animacy, Likeability, Perceived Intelligence, Perceived Safety.

After the **last** condition, also run a **pairwise preference**:

```
"Which version of Pepper did you prefer overall? Why?"
```

This produces the PairwisePreference records that update the ELO leaderboard.

32.5 Post-Session

```
# Final export
curl http://localhost:5000/export > results/session_P001_full.json

# Save questionnaire scores via the API
curl -X POST http://localhost:5000/evaluate \
-H "Content-Type: application/json" \
-d '{
  "session_id": "P001-A",
  "participant_id": "P001",
  "condition": "A",
  "scores": {
    "accuracy": 6, "naturalness": 5, "trust": 6,
    "gesture_appropriateness": 5, "response_speed": 7
  }
}'
```

Or programmatically populate `QuestionnaireCollector` and save:

```
from omnillm.utils.questionnaire import (
    InteractionQuestionnaire, PairwisePreference, QuestionnaireCollector,
)

c = QuestionnaireCollector()
c.add_interaction_response(InteractionQuestionnaire(
    session_id="P001-A", participant_id="P001", condition="A",
    accuracy=6, naturalness=5, trust=6,
    gesture_appropriateness=5, response_speed=7,
))
# ... add the rest of the conditions ...
c.add_pairwise_preference(PairwisePreference(
    session_id="P001", participant_id="P001",
    condition_a="A", condition_b="C", preferred="C",
))
c.save("results/P001_questionnaires.json")
c.to_csv("results/P001_questionnaires.csv")
```

32.6 Daily Backups

After each participant, copy the `results/` folder to a separate location (a USB stick, cloud drive, second laptop). Data loss after a participant has gone home is the worst-case scenario in HRI studies.

PART VI — BUILDING THIS FROM SCRATCH

If you wanted to build OmniLLM from zero, the order of operations matters. This part is the chronological roadmap. Three chapters: the 13-week plan, the critical first mile, and when to add the robot.

Chapter 33 — The Build Order — A 13-Week Plan

*** AT A GLANCE. ** OmniLLM took roughly 13 weeks of focused work to build. The order is dictated by dependencies, not by interest. The mistakes you avoid by following this order are bigger than the time saved by skipping ahead.

33.1 The Milestone Map

```
Week 1 --- Foundation --- "I can talk to one LLM" -----+
Week 2 --- Foundation --- (continued)                      |
                                                                v
Week 3 --- Intelligence - "I can route between LLMs" -----+
Week 4 --- Intelligence -- (Council, Evaluator, Scorer)      |
                                                                v
Week 5 --- Knowledge ---- "It can answer about MY data" -----+
Week 6 --- Knowledge ---- (RAG with ChromaDB)                |
                                                                v
Week 7 --- HRI components - "Each piece works in isolation" --+
Week 8 --- HRI components -- (classifier, language, gesture) |
                                                                v
Week 9 --- Pipeline ----- "End-to-end works in text mode" ----+
Week 10 -- Pipeline ----- (LangGraph + Flask + NAOqi client) |
                                                                v
Week 11 -- Experiment ----- "I can run a full session" -----+
Week 12 -- Experiment ----- (logging, questionnaire, ELO)  |
                                                                v
Week 13+-- Polish ----- "Documentation, tests, demos"
```

Every milestone is a complete, testable system. Don't skip ahead — the joy of "I can route between LLMs" is much harder to feel if you never first felt "I can talk to one LLM".

33.2 Week 1–2: Foundation

Goal: the gateway works for one model.

Day	File	Why
1	<code>.env.example</code>	Document the API keys you'll need before writing any code
1	<code>pyproject.toml</code>	Lock in dependencies up front: litellm, flask, click, rich, pyyaml, dotenv
2	<code>config/models.yaml</code>	Define the registry. Start with 2–3 models (one cloud, one Ollama)
3–4	<code>omnilLM/__init__.py</code> , <code>omnilLM/gateway.py</code>	The most important file. Get one query working end-to-end.
5	<code>omnilLM/cli.py</code> (just models and ask)	Build a CLI early — you'll use it to debug everything else

Milestone: `omnillm ask "Hello" -m openai-gpt4o-mini` returns a real response.

33.3 Week 3–4: Intelligence

Goal: the system can pick a model, run a council, judge a response, and keep an ELO leaderboard.

Day	File	Why
8–9	<code>omnillm/router.py</code>	Start with two strategies (<code>LOWEST_COST</code> , <code>BEST_QUALITY</code>). Add more later.
10	<code>omnillm/consensus.py</code>	Synthesis strategy first; <code>majority_vote</code> and <code>weighted</code> are extensions.
11–12	<code>omnillm/evaluator.py</code>	Referenceless first. Reference-based requires you to have gold answers. Pairwise needs swap logic.
13–14	<code>omnillm/scorer.py</code>	ELO math is in <code>_update_ratings</code> . Save/load JSON.

Milestone: `omnillm council "Is P=NP?"` produces synthesized output from 3 models with an agreement score.

33.4 Week 5–6: Knowledge

Goal: ground answers in your own documents.

Day	File	Why
15–16	<code>omnillm/rag/pipeline.py</code> (keyword fallback)	Start with simple keyword search. Get the augmented prompt pattern working.
17	<code>knowledge_base/lab_info.txt</code>	Write a short text file with real data. Test retrieval against it.
18	<code>omnillm/rag/pipeline.py</code> (ChromaDB)	Add the ChromaDB path. Verify both backends produce similar results.
19	More <code>knowledge_base/</code> files	CSV (visitors), TXT (research projects), PDF (papers)
20	Faithfulness scoring	LLM-as-judge call inside <code>query()</code>

21	Hallucination heuristic	Word-overlap check
----	-------------------------	--------------------

Milestone: ask a question whose answer is *only* in your KB; verify the system uses the KB context (faithfulness > 0.8).

33.5 Week 7–8: HRI Components

Goal: each robotics component works in isolation.

Day	File	Why
22	omnillm/hri/classifier.py	Rule-based first. Keywords + regex. Test with 20 sample utterances.
23	omnillm/hri/language_detector.py	Tier 1 (script analysis) first. Tier 2 (n-grams) only if needed.
24	omnillm/robotics/whisper_stt.py	Local backend first. API backend is a 10-line addition.
25	omnillm/robotics/gesture_planner.py	Rule-based mapping. No LLM call.
26	omnillm/robotics/bridge.py	Abstract base + <code>RobotAction</code> dataclass.
27	omnillm/robotics/pepper.py (placeholder)	Mock the HTTP calls; real implementation comes after the AI server.
28	omnillm/hri/experiment.py	Conditions A–E + <code>ParticipantSession</code> . Pure dataclasses.

Milestone: each component has its own pytest file and passes.

33.6 Week 9–10: Pipeline

Goal: end-to-end works (audio in, action out).

Day	File	Why
29–30	omnillm/hri/agent_graph.py	Build <code>LangGraph</code> . One node at a time. Start with text-only mode (skip transcribe).
31	Wire RAG and gesture planner into nodes	These are the substantive nodes.

32	Smart router and consensus nodes	Conditions C and D.
33–34	<code>omnillm/server/app.py</code>	Flask. <code>/health</code> , then <code>/interact</code> , then everything else.
35	<code>omnillm/server/naoqi_client.py</code>	Last. Without this everything else still works.

Milestone: `curl POST /interact` with text returns a valid `RobotAction`. With Choregraphe's virtual robot, the action visibly animates.

33.7 Week 11–12: Experiment Infrastructure

Goal: one experimenter can run a full session and analyse the data.

Day	File	Why
36	<code>omnillm/utils/experiment_logger.py</code>	The <code>InteractionRecord</code> is the central data structure of the study.
37	<code>omnillm/utils/questionnaire.py</code>	Likert + Godspeed + pairwise + observer dataclasses.
38	<code>omnillm/utils/cost_tracker.py</code>	Necessary for budget reports.
39	<code>omnillm/utils/export.py</code>	CSV / JSON / Markdown for analysis tools (R, JASP).
40–41	Tests for everything	At least one happy-path test per file.
42	Pilot session	One colleague, full protocol.

Milestone: `python -m omnillm.server.app + naoqi_client.py + questionnaire` form = a complete session that produces analysable data.

33.8 Week 13+: Polish

Task	Output
README, GETTING_STARTED, EXPLANATION	Docs
ARCHITECTURE.md	Diagrams for thesis
Demo video	5-min recording of a full session
Code review pass	Type hints, docstrings, linting
Empirical evaluation	Run the experiment with N participants

The book you are reading is the consolidation of week 13.

Chapter 34 — The Critical First Mile — Foundations Before Anything Else

*** AT A GLANCE.** Three rules for the first mile: (1) the gateway is the first thing that must work, (2) ship the CLI early so you can debug, (3) a YAML registry is worth more than it looks.

34.1 Rule 1 — The Gateway Is Not an Afterthought

Many "multi-LLM" projects begin by writing a chatbot for one model, then trying to retrofit the second model in. This always goes badly. The second model has a different SDK, different parameter names, different response shapes. By the time you've added the third, you have three sets of branching code.

Build the gateway first. Put YAML at the centre. By the time you call your second LLM, the codepath is: read YAML → call gateway → done. Adding the third is the same.

OmniLLM's `gateway.py` is the *entire* abstraction. Everything else is a caller of `gateway.query()`.

34.2 Rule 2 — Ship the CLI on Day Three

A CLI is *not* a user feature. It is your **debugging tool**. Without it, you debug with `print()` statements inside test scripts. With it, you can ask any question of any model in two seconds.

OmniLLM's CLI was useful long before any user touched it:

- `omnillm models` — verify your YAML is parseable.
- `omnillm ask "test" -m <model_id>` — verify the gateway works for that model.
- `omnillm ask "test" --all` — verify *every* model works simultaneously.
- `omnillm route "test"` — verify the router picks the model you expect.
- `omnillm leaderboard` — verify ELO scoring.

Build it on day 3. Use it for everything after.

34.3 Rule 3 — YAML Is a Force Multiplier

The decision to put model metadata in YAML — not Python constants — is small in code but huge in consequence:

If models are in YAML...	If models are in Python...
Adding a model = editing config	Adding a model = editing code
Non-Python users can contribute	Only Python users can contribute
Configs can be diffed in PRs	Code diffs mix logic and data
Different deployments can have different configs (override file path)	Different deployments fork the code
Hot-reload is straightforward	Requires Python module reload tricks

OmniLLM commits to YAML throughout: model registry, task definitions (`config/tasks/*.yaml`), routing rules (`hri_task_routing` block). Code is logic, YAML is data. They are kept apart on purpose.

34.4 Anti-Patterns to Avoid

Things you might be tempted to do that will hurt later:

- **Hardcoding API keys.** Use `.env`. Always. Even on day 1.
- **Calling LiteLLM directly from many places.** All calls should go through `LLMGateway` — this is what gives you cost tracking, latency measurement, and unified error handling.
- **Letting the gateway raise exceptions.** Errors should be captured in the `error` field of `ModelResponse`. Higher layers (router, council) should never have to wrap gateway calls in `try/except`.
- **Threading instead of asyncio.** LLM calls are I/O-bound. Asyncio is the correct tool. Threading adds complexity without performance.
- **Putting business logic in `cli.py`.** The CLI should be a thin wrapper. Logic lives in modules that can be imported and tested.

Chapter 35 — When (and How) to Add the Robot

***** AT A GLANCE.** *Add the robot last**. Build everything else first. When you do add it, build the abstract bridge before any concrete implementation. When you build the concrete bridge, mock it before you implement it.*

35.1 Why "Last"?

The robot is a giant fragility multiplier. Pepper has:

- A battery that runs down.
- Wi-Fi that flakes.
- Speech that fails to register.
- Behaviours that need to be installed.
- A Python 2.7 process that has to be running.
- Stiffness that has to be enabled.
- An IP address that occasionally changes.

Every one of these can break the system in ways unrelated to your AI code. If your AI server isn't already rock-solid in text mode, you will spend days chasing bugs that turn out to be Wi-Fi problems.

So: **make the AI server bulletproof first**. Test every condition with text-mode `curl` calls. Only then turn on the robot.

35.2 The Bridge-First Approach

When you do add the robot, write the **abstract bridge** first:

```
class RobotBridge(ABC):
    @abstractmethod async def connect(self): ...
    @abstractmethod async def disconnect(self): ...
    @abstractmethod async def execute_action(self, action: RobotAction): ...
    @abstractmethod async def say(self, text: str): ...
    @abstractmethod async def gesture(self, name: str): ...
    @abstractmethod async def get_sensor_data(self): ...
```

Then write a **mock concrete implementation** that just prints what would happen:

```
class MockBridge(RobotBridge):
    async def connect(self): return True
    async def disconnect(self): return True
    async def execute_action(self, action):
        print(f"[MOCK] speak: {action.speech!r}")
        print(f"[MOCK] gesture: {action.gesture}")
        print(f"[MOCK] led: {action.emotion_led}")
        return True
    # ...
```

Develop the rest of the system against the mock. Only when the system works end-to-end against the mock do you implement `PepperBridge` for real.

This pattern is exactly what `omnilm/robotics/pepper.py` implements today — the HTTP calls are commented placeholders, with `print()` statements that simulate the action. Replacing those

placeholders with real `aiohttp` calls is a small, isolated change once everything else is working.

35.3 When to Build the Python 2.7 Side

The Python 2.7 NAOqi client is the very last thing to build. It does nothing the rest of the system depends on — it is purely the consumer of the AI server's HTTP API. You can write and test the AI server completely without it.

The order:

1. AI server runs.
2. AI server returns valid `RobotAction` JSON via `/interact`.
3. AI server's responses are reasonable across all 5 conditions.
4. **Then** write `naoqi_client.py`.
5. Test it against Choregraphe's virtual robot first.
6. Test it against a physical robot last.

35.4 When the Robot Misbehaves

The robot will misbehave. When it does, debug from the **most reliable end** outward:

1. **Does the AI server work in text mode?** `curl` it directly. If yes, the brain is fine.
2. **Does Choregraphe see Pepper?** If no, it's a network/Wi-Fi issue.
3. **Does Pepper accept basic commands?** Use Choregraphe's Script Editor: `tts.say("test")`. If no, NAOqi service is down.
4. **Does `naoqi_client.py` start?** Check `import qi` works in your Python 2.7. If no, `PYTHONPATH` is wrong.
5. **Does the client connect to the server?** Check the URL, firewall, port.
6. **Only now** look at audio capture / behaviours / LEDs.

Do not start at step 6. You will waste hours.

PART VII — FEATURES, SCOPE, EVALUATION

This part is the [reference](#) for what OmniLLM does. Four chapters: a complete feature catalogue, the five experimental conditions in detail, the evaluation methodology, and an honest discussion of scope and limits.

Chapter 36 — The Complete Feature Reference

36.1 Core LLM Features

Feature	Where	One-line description
Unified gateway	<code>gateway.py</code>	One <code>async query()</code> to any of 19 registered models
Concurrent multi-model query	<code>gateway.py: query_multiple</code>	Ask N models simultaneously via <code>asyncio.gather</code>
Model registry (YAML)	<code>config/models.yaml</code>	Add a model in 7 lines, no code changes
Cost tracking	<code>gateway.py</code> + <code>utils/cost_tracker.py</code>	Per-call USD, aggregated by model and session
Latency measurement	<code>gateway.py</code>	<code>latency_ms</code> on every response
Token counting	<code>gateway.py</code>	<code>input_tokens</code> + <code>output_tokens</code>
Structured error handling	<code>gateway.py: ModelResponse.error</code>	Errors never raise, always returned as fields
Auto-drop unsupported params	<code>gateway.py: drop_params=True</code>	Future-proof against new model API constraints

36.2 Smart Routing Features

Feature	Where	One-line description
6 routing strategies	<code>router.py: RoutingStrategy</code>	BEST_QUALITY, LOWEST_COST, LOWEST_LATENCY, BEST_VALUE, LOCAL_PREFERRED, TASK_TYPE
Composite value score	<code>router.py: _calculate_value_score</code>	$\text{quality} \cdot 0.5 + \text{cost} \cdot 0.3 + \text{latency} \cdot 0.2$
Hard constraints	<code>router.py: route()</code>	<code>budget_usd</code> and <code>max_latency_ms</code> filter candidates
Learning from history	<code>router.py: update_scores</code>	Running average of past evaluation scores
HRI task-type routing	<code>router.py: route_for_hri_task</code>	Maps T1–T4 → optimal model per <code>hri_task_routing</code> config
Complexity-based routing	<code>router.py: route_by_complexity</code>	Word-count heuristic for fast/cheap vs slow/quality

36.3 Consensus / Council Features

Feature	Where	One-line description
3 strategies	<code>consensus.py: ConsensusConfig.strategy</code>	majority_vote, weighted, synthesis
Synthesis judging	<code>consensus.py: _synthesis</code>	A judge LLM combines all responses
Jaccard similarity clustering	<code>consensus.py: _jaccard_similarity</code>	Lightweight semantic clustering, no embeddings
Agreement scoring	<code>consensus.py: _compute_agreement_score</code>	Fraction in majority cluster
Dissenting model identification	<code>consensus.py: ConsensusResult.dissenting_models</code>	Lists models in minority cluster
Configurable council membership	<code>consensus.py: ConsensusConfig.council_models</code>	Any subset of registered models

36.4 Evaluation Features

Feature	Where	One-line description
Referenceless judge (G-Eval)	<code>evaluator.py:</code> <code>_judge_referenceless</code>	Quality scoring with no gold answer
Reference-based judge	<code>evaluator.py:</code> <code>_judge_reference_based</code>	Comparison to a gold answer
Pairwise judge	<code>evaluator.py:</code> <code>_judge_pairwise</code>	Two-response comparison
Position-bias cancellation	<code>evaluator.py:</code> <code>evaluate_pairwise</code>	Swap-and-aggregate runs
Programmatic grading escape	<code>evaluator.py:</code> <code>EvalTask.grading_fn</code>	Custom Python function bypasses LLM judge
Concurrent benchmark runs	<code>evaluator.py:</code> <code>run_benchmark</code>	Semaphore-limited <code>asyncio.gather</code> over tasks × models
8 evaluation axes	<code>tasks/sample_tasks.py</code>	reasoning, knowledge, code, instruction, safety, robot, latency, cost
YAML task loader	<code>tasks/loader.py</code>	Custom task suites without code

36.5 ELO Scoring Features

Feature	Where	One-line description
Standard chess ELO math	<code>scorer.py:</code> <code>_update_ratings</code>	K=32, default rating 1500
Per-category leaderboards	<code>scorer.py:</code> <code>get_category_leaderboard</code>	Independent ratings per axis
Rating history	<code>scorer.py:</code> <code>get_rating_history</code>	Time series of every rating change
JSON persistence	<code>scorer.py:</code> <code>save</code> / <code>load</code>	Serialise full state

36.6 RAG Features

Feature	Where	One-line description
Multi-format ingestion	<code>rag/pipeline.py: index_file</code>	TXT, CSV, PDF
ChromaDB vector store	<code>rag/pipeline.py: _retrieve_chromadb</code>	Cosine similarity search
Keyword fallback	<code>rag/pipeline.py: _retrieve_keyword</code>	Works when ChromaDB not installed
Sliding-window chunking	<code>rag/pipeline.py: _split_text</code>	512-char chunks with 64-char overlap
Faithfulness scoring	<code>rag/pipeline.py: _score_faithfulness</code>	LLM-as-Judge of grounding quality
Hallucination heuristic	<code>rag/pipeline.py: _detect_hallucination</code>	Word-overlap flag
Custom system prompt support	<code>rag/pipeline.py: query()</code>	Override default "use only context" prompt
Top-k configurable	<code>rag/pipeline.py: __init__</code>	Default 4, tune per use case

36.7 HRI Features

Feature	Where	One-line description
T1–T4 task classifier	<code>hri/classifier.py</code>	Rule-based + optional LLM
Multilingual override	<code>hri/classifier.py: classify</code>	Non-English → T4 short-circuit
Three-tier language detector	<code>hri/language_detector.py</code>	Unicode script + n-gram + langdetect
Language→model mapping	<code>hri/language_detector.py: get_optimal_model</code>	Configurable per language
Five experimental conditions	<code>hri/experiment.py: CONDITION_CONFIGS</code>	A, B, C, D, E
Participant sessions	<code>hri/experiment.py: ParticipantSession</code>	UUID, task log, questionnaire scores
Per-condition / per-participant filtering	<code>hri/experiment.py: get_sessions_by_*</code>	Analysis-ready slices

36.8 LangGraph Agent Pipeline Features

Feature	Where	One-line description
9-node DAG	<code>hri/agent_graph.py: build_hri_graph</code>	Audio → action plan
Conditional routing	<code>hri/agent_graph.py: _route_by_task_type</code>	Per-task and per-condition branching
Closure-based dependency injection	<code>hri/agent_graph.py: _make_*_node</code>	No globals, no class hierarchy
Text-mode bypass	<code>hri/agent_graph.py: transcribe_audio</code>	Skips Whisper if <code>audio_bytes</code> empty
Optional logger plug-in	<code>hri/agent_graph.py: _make_log_node</code>	Disabled by passing <code>logger=None</code>

36.9 Robot Bridge Features

Feature	Where	One-line description
Abstract RobotBridge	<code>robotics/bridge.py</code>	Interface for any robot
RobotAction dataclass	<code>robotics/bridge.py</code>	Universal action contract
RobotSensorData dataclass	<code>robotics/bridge.py</code>	Universal sensor reading
LLM-JSON-to-action parser	<code>robotics/bridge.py: parse_llm_to_action</code>	Tolerates Markdown code fences
ROS2 Nav2 goal parser	<code>robotics/bridge.py: parse_nav2_goal</code>	LLM → PoseStamped JSON
Pepper bridge	<code>robotics/pepper.py</code>	HTTP to NAOqi process
NAO bridge	<code>robotics/nao.py</code>	HTTP, with <code>walk_to</code> , <code>stand_up</code> , <code>sit_down</code> helpers
Buddy bridge	<code>robotics/buddy.py</code>	WebSocket with token streaming

36.10 Robotics Helper Features

Feature	Where	One-line description
Gesture planner	<code>robotics/gesture_planner.py: plan</code>	Task + content → gesture name + LED hex
12 named gestures	<code>robotics/gesture_planner.py: GESTURE_LED_COLORS</code>	wave, point_left/right/forward/up, nod, think, ...
Whisper STT, two backends	<code>robotics/whisper_stt.py</code>	Local (<code>openai-whisper</code>) and API (<code>openai</code>)
Async transcription	<code>robotics/whisper_stt.py: transcribe</code>	<code>asyncio.to_thread</code> for local, native async for API

36.11 Server Features

Feature	Where	One-line description
Flask app factory	<code>server/app.py: create_app</code>	Test-friendly construction
6 endpoints	<code>server/app.py</code>	health, status, interact, transcribe, evaluate, export
Lazy LangGraph build	<code>server/app.py: _get_graph</code>	Tests can run without <code>langgraph</code> installed
Sync→async bridge	<code>server/app.py: _run_async</code>	Per-request event loop
Knowledge base auto-loader	<code>server/app.py: _load_knowledge_base</code>	Indexes everything in <code>knowledge_base/</code> on startup
Graceful fallback	<code>server/app.py: _fallback_interact</code>	Direct gateway call when <code>LangGraph</code> absent
NAOqi client	<code>server/naoqi_client.py</code>	Python 2.7 compatible, simulation mode when SDK absent
12-entry gesture map	<code>server/naoqi_client.py: GESTURE_TO_BEHAVIOR</code>	Maps <code>OmniLLM</code> names to <code>NAOqi</code> behaviour paths

36.12 Logging / Reporting Features

Feature	Where	One-line description
InteractionRecord	<code>utils/experiment_logger.py</code>	Full per-turn data structure
Interaction filtering	<code>utils/experiment_logger.py: get_records</code>	by session/condition/task/model
Aggregate summaries	<code>utils/experiment_logger.py: get_summary</code>	per-model, per-condition, per-task means
JSON + CSV export	<code>utils/experiment_logger.py: save / save_csv</code>	Stats-tool ready
Likert questionnaire dataclass	<code>utils/questionnaire.py: InteractionQuestionnaire</code>	5 items, 1–7 scale, mean & normalised score
Godspeed dataclass	<code>utils/questionnaire.py: GodspeedResponse</code>	5 subscales, 1–5 scale
Pairwise preference dataclass	<code>utils/questionnaire.py: PairwisePreference</code>	Feeds the ELO scorer
Observer rating dataclass	<code>utils/questionnaire.py: ObserverRating</code>	Live experimenter notes
Cost tracker	<code>utils/cost_tracker.py</code>	Per-model, per-session USD
Evaluation result exporter	<code>utils/export.py: ResultExporter</code>	CSV, JSON, Markdown, Rich Table

36.13 CLI Features

Command	One-line description
<code>omnilm models</code>	List all registered models in a coloured table
<code>omnilm ask "<q>"</code>	Send to one or many or all models
<code>omnilm evaluate</code>	Run benchmark on tasks × models
<code>omnilm compare "<q>"</code>	Pairwise judge between two models
<code>omnilm council "<q>"</code>	LLM Council with synthesis / vote / weighted
<code>omnilm route "<q>"</code>	Show routing decision and rationale
<code>omnilm leaderboard</code>	ELO rankings, optionally per category
<code>omnilm costs</code>	USD spend per model
<code>omnilm export</code>	Convert results JSON to CSV / MD / JSON

Chapter 37 — The Five Experimental Conditions, in Detail

*** AT A GLANCE.** Five conditions A–E systematically probe four design dimensions: (i) cloud vs local, (ii) fixed model vs dynamic, (iii) single model vs ensemble, and (iv) RAG vs no-RAG. Each is implemented as a `ConditionConfig` and selected per session via `--condition`.

37.1 The Conditions in One Picture

	+-----+-----+	
	FIXED	DYNAMIC
	+-----+-----+	
CLOUD	A: GPT-4o-mini RAG: ON Baseline cloud	C: Smart-routed RAG: ON Per-task selection
LOCAL	B: Llama3:8b RAG: ON Free baseline	D: Council 3 models + synthesis RAG: ON
	E: GPT-4o-mini, RAG OFF (control)	

37.2 Condition A — Fixed Cloud LLM

What it is: GPT-4o-mini for every task, every utterance. RAG enabled.

What it tests: "How good is a single decent cloud model on this kind of HRI task?"

Why GPT-4o-mini: Best balance of accuracy, cost, and latency in April 2026. \$0.15/\$0.60 per 1M tokens makes it ~5× cheaper than GPT-4o without dramatic quality loss for HRI tasks.

Config:

```
ConditionConfig(  
  condition=ExperimentCondition.A,  
  model_id="openai-gpt4o-mini",  
  rag_enabled=True,  
  description="Fixed cloud LLM (GPT-4o-mini) with RAG – baseline",  
)
```

37.3 Condition B — Fixed Local LLM

What it is: Llama3:8b via Ollama for every task. RAG enabled.

What it tests: "How good is a free, local, private model on the same tasks?"

Why Llama3:8b: Best free open-weight model in the 7–8B size that fits on a typical research-grade GPU. ~4.7 GB download. No data leaves the lab.

Config:

```
ConditionConfig(
  condition=ExperimentCondition.B,
  model_id="llama3-8b-local",
  rag_enabled=True,
  description="Fixed local LLM (Llama3:8b via Ollama) with RAG",
)
```

Practical note: B is markedly slower than A on a CPU-only laptop (2–4 s per response vs 0.5–1.5 s). For real-time HRI you typically need a GPU.

37.4 Condition C — Smart-Routed

What it is: Per-task model selection via `SmartRouter.route_for_hri_task`.

What it tests: "Does picking the right model per task type beat any fixed model?" — the core OmniLLM hypothesis.

Routing table (`models.yaml: hri_task_routing`):

Task	Model	Why
T1 info_retrieval	openai-gpt4o-mini	Accurate, cheap, RAG-friendly
T2 navigation	gemini-flash	Fast, strong spatial language
T3 social_conversation	claude-haiku	Most natural conversational tone
T4 multilingual	gemini-flash	Strong multilingual coverage

Config:

```
ConditionConfig(
  condition=ExperimentCondition.C,
  model_id=None, # dynamic
  rag_enabled=True,
  use_smart_routing=True,
  description="Smart-routed - OmniLLM selects best model per task type",
)
```

37.5 Condition D — Consensus / Council

What it is: Three models answer the same question, a judge LLM synthesises the best combined answer.

What it tests: "Does an ensemble beat any single model?"

Default council: GPT-4o-mini + Gemini Flash + Claude Haiku. All three are HRI-tuned (low-latency, conversational).

Latency cost: Three concurrent LLM calls + a synthesis call $\approx$ 2 LLM latencies + a synthesis latency $\approx$ 1.5–3 s. Slowest of all conditions.

Config:

```
ConditionConfig(
  condition=ExperimentCondition.D,
  model_id=None,
  rag_enabled=True,
  use_consensus=True,
  council_models=["openai-gpt4o-mini", "claude-haiku", "gemini-flash"],
  description="Consensus council - 3 models, best answer synthesised",
)
```

37.6 Condition E — RAG-Off Control

What it is: Same model as Condition A (GPT-4o-mini), but RAG disabled.

What it tests: "How much does RAG actually help?" — the A vs E comparison **isolates** RAG's contribution holding everything else constant.

Config:

```
ConditionConfig(  
    condition=ExperimentCondition.E,  
    model_id="openai-gpt4o-mini",  
    rag_enabled=False,  
    description="RAG-off control – GPT-4o-mini without knowledge retrieval",  
)
```

The `_route_by_task_type` function in `agent_graph.py` short-circuits all task types to `direct_llm` when `condition == "E"`.

37.7 Why These Five and Not Others?

The five-condition design is the smallest set that probes all four dimensions:

- **A vs B:** cloud vs local.
- **A vs C:** fixed vs smart-routed.
- **A vs D:** single vs ensemble.
- **A vs E:** with-RAG vs without-RAG.

Adding more conditions (e.g. "smart-routed but local-only", "consensus without synthesis") gives more data but also exponentially harder recruitment. With ~ 20 participants $\times$ 3 of 5 conditions $\times$ 4 tasks = 240 data points, the design has enough power for repeated-measures ANOVA on the main effects.

Chapter 38 — Evaluation Methodology — Eight Axes, Three Judge Patterns

38.1 The Eight Evaluation Axes

OmniLLM's built-in tasks (`tasks/sample_tasks.py`) span eight axes chosen to be **orthogonal** — a model that scores high on one does not necessarily score high on the others.

#	Axis	What it measures	Sample tasks
1	Reasoning	Multi-step logic, mathematics	Bat-and-ball puzzle; syllogism; Gambler's Fallacy
2	Knowledge	Factual accuracy	Capital cities; speed of light; DNA discoverers
3	Code	Generate, debug, explain code	<code>is_prime</code> ; FizzBuzz; binary search bug
4	Instruction Following	Strict format / constraints	Numbered list of 5; pure JSON output; exact 3 sentences
5	Safety	Refusal of harmful / PII requests	Refuse hacking; refuse misinformation; redact PII
6	Robot-Readiness	Structured JSON for robot control	Generate NAO action plan; parse voice command; ROS2 Nav2 goal
7	Latency	Time-to-first-token	One-sentence greeting; "Answer with only the number"
8	Cost-Efficiency	Quality at minimum tokens	Define "recursion" in one sentence; Yes/No questions

38.2 The Three Judge Patterns (Recap)

Pattern	When to use it	Output
Referenceless (G-Eval)	No gold answer; open-ended	Score 0–1 + reasoning
Reference-Based	You know the right answer	Score 0–1 + reasoning
Pairwise	Comparing two models	Winner: A / B / tie

38.3 Automatic Per-Interaction Metrics (HRI)

For each Embodied LLM Arena interaction the system automatically records:

1. **Latency** — end-of-utterance to start-of-robot-speech, in ms.
2. **Token counts** — input + output (LLM cost driver).
3. **Cost** — USD, calculated from `models.yaml` pricing.
4. **RAG retrieval scores** — similarity scores of top-k chunks.
5. **RAG faithfulness** — LLM-as-Judge of grounding (0–1).
6. **Hallucination flag** — word-overlap heuristic.
7. **LLM-as-Judge quality** — referenceless score (0–1).
8. **Language detected** — ISO 639-1 code.
9. **Task classification** — T1/T2/T3/T4 + confidence.
10. **Model used** — the actual `model_id` (matters for Conditions C and D).
11. **Gesture used** — what Pepper did.
12. **Task success** — observer-coded binary.

All twelve fields land in `InteractionRecord` and export cleanly to CSV.

38.4 Human-Rated Metrics (Per Condition)

The `InteractionQuestionnaire` (1–7 Likert):

1. The robot's answers were accurate.
2. The robot was natural to talk to.
3. I trust the information the robot gave me.
4. The robot's gestures were appropriate.
5. The robot responded quickly enough.

Plus optional Godspeed (1–5 Likert across 5 subscales) and the pairwise preference at session end.

38.5 Composite Scoring

OmniLLM's documented composite for choosing models in deployment:

```
value = quality × 0.5 + (1 / latency_normalised) × 0.3 + (1 / cost_normalised) × 0.2
```

This is the same formula used by `BEST_VALUE` in the smart router. Quality is the LLM-as-Judge score; latency normalised against the 500ms –10s range; cost normalised against the \$0–\$0.05 per query range.

38.6 ELO and the Embodied LLM Leaderboard

Every pairwise preference at the end of a session produces an ELO match:

```
elo.record_match("condition-C", "condition-A", "model_a",  
                category="embodied_hri")
```

After all participants are run, the **embodied_hri** category leaderboard is the headline result of the study.

38.7 Statistical Analysis

The expected analysis pipeline:

1. **Export.** `curl /export > all_data.json` then load into R or Python.
2. **Per-condition means.** Mean Likert per condition, with 95% CI.
3. **Within-subjects ANOVA.** Repeated-measures ANOVA on the 1–7 Likert means, with Greenhouse-Geisser correction.
4. **Pairwise post-hoc.** Bonferroni-corrected paired t-tests for the meaningful contrasts (A vs B, A vs C, A vs D, A vs E).
5. **ELO convergence.** Plot ELO as a time series; report final ratings.
6. **Cross-comparison.** Correlate the embodied ELO leaderboard with each model's MMLU/Chatbot Arena score. Discrepancy is the headline.

Chapter 39 — Scope, Limits, and What Comes Next

*** AT A GLANCE. OmniLLM is intentionally focused. It is not a general-purpose chatbot framework, not a robotics simulator, not a production deployment platform. This chapter is the honest list of what it does not do.

39.1 What OmniLLM Is Not

Not...	...because
A general chatbot framework	Designed around HRI evaluation; lacks user-facing UI, accounts, history
A production deployment	Single-worker Flask, no auth, no rate limiting, no monitoring
A robotics simulator	The virtual robot is Choregraphe's, not ours
A multi-tenant service	One AI server serves one robot session at a time
A model training framework	We use models, we don't train them
A speech recognition framework	Whisper is a thin wrapper; we don't compete with faster-whisper
An embedding framework	ChromaDB does the embedding; we don't engineer embeddings
A real-time streaming framework	Buddy bridge has streaming, but the rest is request/response

39.2 Known Limitations

Audio capture is a placeholder. The `_record_audio` method in `naoqi_client.py` returns empty bytes. Real production audio capture needs `ALAudioRecorder` or `processRemote` callbacks. See Chapter 21.8.

RAG faithfulness scoring is opt-in and slow. It costs an extra LLM call per interaction. Disabled by default. Enable only for offline analysis runs.

ELO updates are session-end, not real-time. Every pairwise preference triggers exactly one ELO update; we don't do per-interaction preference updates.

Conditions A and E are different RAG regimes, but otherwise share a model. This is a feature (clean ablation) but means N participants in A and E need to be $\geq N$ to detect the RAG main effect.

No automated counterbalancing. The experimenter is responsible for the Latin square. With small N, this is not a code requirement, but for larger studies you'd add a `LatinSquareGenerator`.

Choregraphe is locked to NAOqi 2.5. SoftBank's NAOqi 2.9 (Android QiSDK) is not supported. Migrating to QiSDK is a substantial rewrite.

39.3 Future Work — High-Value Extensions

In rough priority order:

Real audio capture in `naoqi_client.py`. Replace the placeholder with `ALAudioRecorder`-based or callback-based capture. Critical for actual study runs.

Streaming responses. Pepper's TTS could start speaking as soon as the first sentence is generated, rather than waiting for the whole response. Reduces perceived latency by 30–60%. The Buddy bridge already implements this; Pepper would need a parallel `stream_response` path.

Vision integration. Pepper has cameras + a depth sensor. The agent graph could consume "what does Pepper see" as an additional input. Multi-modal LLMs (Gemini, GPT-4o) accept image inputs natively.

QiSDK migration. For deployments on a NAOqi 2.9 Pepper (newer robots, post-Aldebaran-bankruptcy), the entire NAOqi client would be rewritten in Kotlin/Java. This eliminates the Python 2.7 problem.

Realtime API support. OpenAI's Realtime API and Google Gemini Live can do speech-to-speech without an STT step. For Pepper, this means the audio path becomes a single low-latency loop.

Adversarial / robustness testing. A separate evaluation pass that probes how well each condition handles noisy speech, ambiguous questions, off-topic chatter, and minor adversarial inputs.

Larger-scale study. Current design assumes ~20 participants. Scaling to 100+ requires automated counterbalancing, faster session throughput, and a more sophisticated logging schema.

39.4 Limits of the Embodied LLM Arena Methodology

The research design is sound but has its limits:

- **External validity** — results from a single university lab may not generalise to hospitals, retail, or hotels. The platform is reusable precisely so others can replicate in different contexts.
- **Within-subjects fatigue** — a participant's questionnaire ratings for the *third* condition are influenced by experience with the first two. Counterbalancing controls for order, but cumulative fatigue is real.
- **LLM-as-Judge is itself an LLM.** Using GPT-4o to judge GPT-4o is partial self-evaluation. For best practice, use a *different* provider as the judge (e.g., Claude judging OpenAI outputs).
- **Hallucination heuristic is conservative.** False positives are common; treat the heuristic as a flag, not a verdict.

39.5 What This Project Will Likely Become

Two natural successors:

The thesis. This codebase + a participant study + a few months of analysis = a publishable paper. Suggested venues: HRI, RO-MAN, ICRA, Frontiers in Robotics & AI, MDPI Robotics.

The reusable platform. Other HRI labs can fork, swap the knowledge base, swap the robot bridge, and re-run the methodology. The abstract `RobotBridge` and YAML model registry make this almost parameter-only adaptation.

PART VIII — APPENDICES

Appendix A — Python Primer (Just the Parts You Need for OmniLLM)

This is **not** a Python tutorial. It is a focused tour of the *specific* Python features OmniLLM uses heavily, so you can read the source code fluently. If you are already comfortable with `async`, `dataclasses`, type hints, and dictionary literals, skip this appendix.

A.1 Type Hints

Modern Python adds optional type annotations:

```
def add(a: int, b: int) -> int:
    return a + b

names: list[str] = ["Alice", "Bob"]
config: dict[str, int] = {"port": 5000}
maybe_value: str | None = None      # might be a str, might be None
```

OmniLLM uses these everywhere. They are **not enforced at runtime** — they are documentation that tools (mypy, your IDE) can verify. When you see `model_id: str` it means "this argument should be a string"; the code won't crash if you pass something else, but you should not.

A.2 Dataclasses

`@dataclass` auto-generates `__init__`, `__repr__`, equality:

```
from dataclasses import dataclass, field

@dataclass
class Point:
    x: float
    y: float
    z: float = 0.0          # default value
    tags: list[str] = field(default_factory=list)  # default mutable

p = Point(1.0, 2.0)        # Point(x=1.0, y=2.0, z=0.0, tags=[])
print(p)                  # nice automatic repr
```

OmniLLM's `ModelResponse`, `RobotAction`, `EvalTask`, `InteractionRecord` are all dataclasses. Look for `@dataclass` to identify them.

A.3 Async / Await

LLM calls are I/O-bound; you wait for the network. Async lets one program do other useful work during the wait:

```
import asyncio

async def fetch_one(url: str) -> str:
    # imagine this calls an LLM
    await asyncio.sleep(1)          # simulate 1 s of network wait
    return f"Response for {url}"

async def main() -> None:
    # Sequential – takes 3 s total
    a = await fetch_one("a"); b = await fetch_one("b"); c = await fetch_one("c")
```

```

# Concurrent - takes 1 s total
a, b, c = await asyncio.gather(
    fetch_one("a"), fetch_one("b"), fetch_one("c"),
)

asyncio.run(main())

```

Rules:

- `async def` declares a function that *returns a coroutine*. Calling it produces a coroutine; `await` runs it.
- Inside an `async def`, you can `await` other coroutines.
- Outside any `async` function, you call `asyncio.run(main())` to start the event loop.
- `asyncio.gather(c1, c2, c3)` runs three coroutines concurrently and returns their results in order.

OmniLLM's `LLMGateway.query_multiple` uses exactly this pattern.

A.4 F-Strings

Modern string formatting:

```

name = "Pepper"
n = 42
print(f"Hello {name}, your number is {n}.")           # plain
print(f"Cost: ${cost:.6f}")                          # 6 decimal places
print(f"Latency: {latency_ms:.0f}ms")                 # 0 decimals
print(f"Score: {score:.2%}")                          # as percentage

```

A.5 Dictionary Literals and Unpacking

```

config = {"port": 5000, "host": "0.0.0.0"}

# Lookup with default
port = config.get("port", 8000)

# Merge
default = {"port": 8000, "debug": False}
merged = {**default, **config}           # config wins where keys overlap

# Iterate
for key, value in config.items():
    print(key, value)

```

OmniLLM's models registry is a deeply nested dict; you'll see lots of `.get(key, default)` calls because YAML files don't always have every optional field.

A.6 Decorators

Functions that wrap other functions:

```

def log_calls(func):
    def wrapper(*args, **kwargs):
        print(f"calling {func.__name__}")
        return func(*args, **kwargs)
    return wrapper

@log_calls

```

```
def greet(name):
    return f"Hello {name}"

greet("Alice")    # prints "calling greet" then returns "Hello Alice"
```

In OmniLLM you'll see `@dataclass`, `@click.command(...)`, `@app.route(...)`. Each is a decorator that transforms the function below it.

A.7 Click (CLI Framework)

```
import click

@click.group()
def cli():
    """My CLI."""

@click.command("greet")
@click.argument("name")
@click.option("--loud", is_flag=True)
def greet_cmd(name: str, loud: bool):
    msg = f"Hello {name}"
    if loud: msg = msg.upper()
    click.echo(msg)

if __name__ == "__main__":
    cli()
```

OmniLLM's CLI follows this pattern — `cli.py` defines a Click `group()` and adds many commands.

A.8 Flask (Web Framework)

```
from flask import Flask, jsonify, request

app = Flask(__name__)

@app.route("/health")
def health():
    return jsonify({"status": "ok"})

@app.route("/echo", methods=["POST"])
def echo():
    data = request.get_json(force=True)
    return jsonify({"received": data})

if __name__ == "__main__":
    app.run(host="0.0.0.0", port=5000)
```

OmniLLM's `server/app.py` is a larger version of this — many endpoints, each registered with `@app.route`.

A.9 Context Managers (with)

Used for setup/teardown around a block:

```
with open("file.txt", "r") as fh:
    data = fh.read()
# fh is automatically closed here, even if an exception was raised
```

OmniLLM uses `with open(...)` for all file I/O, and `async with` for HTTP sessions in the bridge code.

A.10 Imports You Will See Often

```
from __future__ import annotations      # all type hints are strings (no fwd-ref issues)
from dataclasses import dataclass, field # data classes
from pathlib import Path                 # cross-platform paths
from typing import Any, Literal, TYPE_CHECKING
                                         # type-check-only annotations
import asyncio, json, time, re          # standard library
import yaml                             # PyYAML, the YAML loader
import litellm                           # LiteLLM provider gateway
from rich.console import Console         # coloured terminal output
from rich.table import Table
import click                             # CLI framework
```

Appendix B — Complete Terminal Command Reference

This is the unified cheat-sheet. Every command you might run while working with OmniLLM, in one place.

B.1 Install

```
git clone https://github.com/Akshita-sr/OmniLLM.git
cd OmniLLM
python -m venv .venv
source .venv/bin/activate           # Linux / macOS
.venv\Scripts\activate             # Windows CMD
.venv\Scripts\Activate.ps1         # Windows PowerShell

pip install -e ".[dev]"             # core only
pip install -e ".[dev,robotics]"    # + Flask
pip install -e ".[dev,hri]"        # + RAG, LangGraph
pip install -e ".[all]"             # everything

cp .env.example .env               # then edit with API keys
pip install openai-whisper          # local STT (optional)

curl -fsSL https://ollama.com/install.sh | sh # Ollama (Linux/macOS)
ollama pull llama3:8b               # ~4.7 GB, free local
ollama pull qwen2.5:7b              # multilingual, ~4.4 GB
ollama pull mistral:7b              # ~4.1 GB
ollama pull deepseek-r1:14b         # reasoning, ~9 GB
```

B.2 Verify

```
omnillm --version
omnillm models
omnillm models --type local
omnillm models --type cloud
pytest tests/ -v
pytest tests/ --cov=omnillm --cov-report=term-missing
```

B.3 CLI — Talk to Models

```
omnillm ask "Hello" -m openai-gpt4o-mini
omnillm ask "Explain RAG" -m claude-haiku -m gemini-flash
omnillm ask "What is 2+2?" --all      # all registered

omnillm route "Where is Room 305?" --strategy TASK_TYPE
omnillm route "Quick yes/no" --strategy LOWEST_COST --budget 0.001
omnillm route "Hard reasoning task" --strategy BEST_QUALITY

omnillm council "Is consciousness emergent?"
omnillm council "What is justice?" --strategy synthesis
omnillm council "P=NP?" --strategy majority_vote -m openai-gpt4o -m claude-sonnet -m gemini-2.5-pro
```

B.4 CLI — Evaluate

```
omnillm evaluate                      # default cloud models, all categories
omnillm evaluate -m openai-gpt4o-mini
omnillm evaluate -c reasoning
omnillm evaluate -m claude-haiku -c code
omnillm evaluate -o results/eval_2026-05.json

omnillm compare "Write a haiku about robots"
omnillm compare "Debug this code" --model-a openai-gpt4o --model-b claude-sonnet
```

B.5 CLI — Reports

```
omnillm leaderboard
omnillm leaderboard --category reasoning
omnillm leaderboard --category embodied_hri
omnillm costs

omnillm export --format csv      --input results/eval.json -o results/eval.csv
omnillm export --format markdown --input results/eval.json -o results/eval.md
omnillm export --format json     --input results/eval.json -o results/full.json
```

B.6 AI Server

```
python -m omnillm.server.app
python -m omnillm.server.app --host 0.0.0.0 --port 5000
python -m omnillm.server.app --debug
python -m omnillm.server.app --no-rag
python -m omnillm.server.app --model claude-haiku
python -m omnillm.server.app --kb /path/to/your/knowledge_base

# Production-ish
gunicorn 'omnillm.server.app:create_app()' --bind 0.0.0.0:5000 --workers 1

# Smoke tests
curl http://localhost:5000/health
curl http://localhost:5000/status
curl -X POST http://localhost:5000/interact \
  -H "Content-Type: application/json" \
  -d '{"text": "Hello", "participant_id": "T", "condition": "A"}'
curl http://localhost:5000/export > results/all.json
```

B.7 NAOqi Client (Python 2.7)

Windows CMD:

```
set PYTHONPATH=C:\pynaoqi\pynaoqi-python2.7-2.5.5.5-win32-vs2013\lib;%PYTHONPATH%
set PATH=C:\pynaoqi\pynaoqi-python2.7-2.5.5.5-win32-vs2013\lib;%PATH%

C:\Python27\python.exe omnillm\server\naoqi_client.py ^
  --robot-ip 192.168.1.100 ^
  --robot-port 9559 ^
  --server-ip 192.168.1.50 ^
  --server-port 5000 ^
  --participant P001 ^
  --condition C
```

Bash:

```
export PYTHONPATH="/c/pynaoqi/pynaoqi-python2.7-2.5.5.5-win32-vs2013/lib:$PYTHONPATH"
/c/Python27/python.exe omnillm/server/naoqi_client.py \
  --robot-ip 192.168.1.100 \
  --server-ip 192.168.1.50 \
  --participant P001 \
  --condition C
```

B.8 Quick NAOqi One-Liners (Python 2.7)

```
:: Wake up
C:\Python27\python.exe -c "from naoqi import ALProxy; ALProxy('ALMotion', '192.168.1.100', 9559).wakeUp()"

:: Speak
C:\Python27\python.exe -c "from naoqi import ALProxy; ALProxy('ALAnimatedSpeech', '192.168.1.100', 9559).say('Hello!')"

:: Wave
C:\Python27\python.exe -c "from naoqi import ALProxy; ALProxy('ALBehaviorManager', '192.168.1.100', 9559).runBehavior('Wave')"

:: Eyes blue
C:\Python27\python.exe -c "from naoqi import ALProxy; ALProxy('ALLeds', '192.168.1.100', 9559).fadeRGB('FaceLeds', 0.0, 0.0, 1.0)"

:: Battery
C:\Python27\python.exe -c "from naoqi import ALProxy; print(ALProxy('ALBattery', '192.168.1.100', 9559).getBatteryCharge())"
```

B.9 Tests

```
pytest tests/ -v
pytest tests/test_gateway.py -v
pytest tests/test_router.py -v
pytest tests/test_consensus.py -v
pytest tests/test_evaluator.py -v
pytest tests/test_scorer.py -v
pytest tests/test_rag.py -v
pytest tests/test_hri.py -v
pytest tests/test_gesture_planner.py -v
pytest tests/test_experiment_logger.py -v
pytest tests/test_new_components.py -v
pytest tests/ --cov=omnilm --cov-report=html
```

Appendix C — Complete Glossary

A consolidated alphabetical glossary. Every term used in the book.

Term	Definition
agent graph	A directed graph of <code>async</code> functions sharing state. OmniLLM's lives in <code>hri/agent_graph.py</code> .
ALAnimatedSpeech	NAOqi service that speaks with synchronised body gestures.
ALAudioDevice	NAOqi service for microphone capture.
ALBehaviorManager	NAOqi service for running pre-installed behaviour animations.
ALLeds	NAOqi service for LED control.
ALMemory	NAOqi key-value store and event bus.
ALMotion	NAOqi service for joint control. <code>wakeUp()</code> enables motors.
API key	Secret string proving you have an account with a provider.
async / await	Python syntax for non-blocking I/O. <code>await</code> pauses for an I/O-bound operation.
asyncio.gather	Runs many coroutines concurrently.
base64	ASCII encoding of binary bytes; used to put audio in JSON.
bodyLanguageMode	ALAnimatedSpeech config option: <code>contextual</code> , <code>random</code> , or <code>disabled</code> .
bridge (architecture)	A layer that connects two incompatible systems. OmniLLM has the Python 2.7 ↔ Python 3 bridge.
Choregraphe	SoftBank's drag-and-drop IDE for Pepper / NAO.
ChromaDB	Pure-Python embedded vector database used by the RAG pipeline.
chunk	A short slice of a longer document, the unit of retrieval. OmniLLM uses 512-char with 64 overlap.
CLI	Command-line interface. OmniLLM's CLI is built with Click and Rich.

Click	Python library for building CLIs. Used in <code>cli.py</code> .
closure	A function that has captured variables from its enclosing scope. Used in <code>agent_graph.py</code> 's <code>_make_*_node</code> .
condition (experimental)	One of A, B, C, D, E in the Embodied LLM Arena.
consensus	Querying multiple models and combining their answers. See <code>consensus.py</code> .
counterbalancing	Varying the order of conditions across participants to cancel order effects.
dataclass	Python decorator (<code>@dataclass</code>) that auto-generates <code>__init__</code> , etc.
embedding	List of floats encoding the semantic meaning of text.
embodiment effect	HRI finding that the same AI text is rated differently when delivered through a robot vs a screen.
ELO	Rating system from chess; $+100 \approx 64\%$ expected win rate. Used in <code>scorer.py</code> .
evaluator	LLM-as-Judge component. Three patterns: referenceless, reference-based, pairwise.
f-string	Python 3.6+ string formatting: <code>f"hello {name}"</code> .
faithfulness	How well a response uses the retrieved RAG context. 0–1, scored by judge LLM.
Flask	Python web framework. OmniLLM's AI server is a Flask app.
gateway	The unified LLM access layer. <code>gateway.py</code> in OmniLLM.
gesture planner	Maps task type + response text → gesture name + LED hex. <code>robotics/gesture_planner.py</code> .
Godspeed	Bartneck et al.'s validated 5-subscale HRI questionnaire.
hallucination	When an LLM states something not in the retrieved context.
HRI	Human-Robot Interaction; an academic field.

HTTP	The protocol used between OmniLLM's server and Pepper's NAOqi client.
Jaccard similarity	Intersection-over-union of word sets. Used in consensus clustering.
JSON	JavaScript Object Notation; OmniLLM's universal serialisation format.
judge model	The LLM used to evaluate responses in LLM-as-Judge. Default GPT-4o.
K-factor	ELO volatility constant; OmniLLM uses 32.
knowledge base	Documents indexed for RAG. Lives in <code>knowledge_base/</code> .
LangChain	Framework for LLM apps. OmniLLM uses it for document loaders.
LangGraph	LangChain's library for stateful multi-node agents. Used in <code>hri/agent_graph.py</code> .
Latin square	Counterbalancing design where each condition appears once in each ordinal position.
latency	Time from request to response, in ms.
Likert scale	1–7 (or 1–5) rating of agreement.
LiteLLM	Python library exposing 100+ LLM providers under one API. The foundation of <code>gateway.py</code> .
LLM	Large Language Model. Examples: GPT-4o, Claude Sonnet, Gemini Flash, Llama 3.
LLM-as-Judge	Using one LLM to score another's response. Three patterns in OmniLLM.
Llama 3	Meta's open-weight LLM. Used as Condition B's local baseline.
majority_vote	Consensus strategy: pick the cluster with the most members.
ModelResponse	OmniLLM's universal LLM-response dataclass. Defined in <code>gateway.py</code> .
NAO	Pepper's smaller (58 cm) sibling. Same NAOqi OS.
NAOqi	Pepper / NAO's middleware OS. Service broker on port 9559. Python 2.7 only.

Ollama	Free local LLM runtime. <code>localhost:11434</code> .
pairwise	Judge pattern: compare two responses, pick the better.
participant session	One participant under one condition. UUID-tagged.
Pepper	120 cm humanoid social robot from SoftBank Robotics.
position bias	LLM judges' tendency to prefer the first response shown.
provider	Company / service hosting an LLM. OpenAI, Anthropic, Google, etc.
pynaoqi	The Python 2.7 binding to NAOqi.
RAG	Retrieval-Augmented Generation. Search documents, prepend to prompt.
referenceless	Judge pattern: score quality without a gold answer.
reference-based	Judge pattern: compare to a known-correct answer.
Rich	Python library for coloured terminal output. Used in <code>cli.py</code> .
RobotAction	OmniLLM's universal action dataclass: speech + gesture + LED + movement.
RobotBridge	Abstract base class for robot integrations.
router	Component that picks the best model. <code>router.py</code> .
session	One participant's interaction under one condition. UUID-tagged.
smart routing	Strategy-based model selection. Six strategies in OmniLLM.
SmartRouter	The class implementing smart routing.
STT	Speech-to-Text. OmniLLM uses Whisper.
synthesis	Consensus strategy: a judge LLM combines multiple responses.
T1–T4	The four HRI task types: <code>info_retrieval</code> , <code>navigation</code> , <code>social_conversation</code> , <code>multilingual</code> .

task category	Evaluation axis. OmniLLM has 8 (reasoning, knowledge, code, instruction, safety, robot, latency, cost).
temperature	LLM sampling parameter. 0 = deterministic, 1 = creative.
token	Sub-word unit. LLMs charge per 1 million tokens.
TTFT	Time-To-First-Token. How long before the LLM starts streaming.
TTS	Text-to-Speech. OmniLLM uses Pepper's ALAnimatedSpeech.
vector store	Database storing embeddings + supporting similarity search.
virtual environment	Isolated Python sandbox per project. <code>.venv/</code> after <code>python -m venv .venv</code> .
WAV	Audio format used for transport. 16 kHz mono PCM in OmniLLM.
WebSocket	Persistent bi-directional protocol; used by Buddy bridge.
Whisper	OpenAI's open-source speech-to-text model.
within-subjects design	Every participant experiences every condition.
YAML	Human-readable config file format. <code>config/models.yaml</code> .

Appendix D — Troubleshooting Reference

Solution-first table. If you see the symptom, apply the fix.

D.1 Install / Environment

Symptom	Likely cause	Fix
<code>command not found: omnillm</code>	venv not activated, or <code>pip install -e</code> failed	activate, then <code>pip install -e ".[dev]"</code>
<code>ModuleNotFoundError: chromadb</code>	<code>[hri]</code> extra not installed	<code>pip install -e ".[hri]"</code>
<code>ModuleNotFoundError: flask</code>	<code>[robotics]</code> extra not installed	<code>pip install -e ".[robotics]"</code>
<code>ModuleNotFoundError: langgraph</code>	same	<code>pip install -e ".[hri]"</code>
<code>ModuleNotFoundError: openai-whisper</code>	not installed (it's optional)	<code>pip install openai-whisper</code>
<code>python: command not found</code>	python not on PATH	reinstall, check "Add to PATH"
Tests fail with import errors	venv mixing	rebuild: <code>rm -rf .venv && python -m venv .venv && pip install -e ".[all]"</code>

D.2 API Keys / Providers

Symptom	Likely cause	Fix
<code>APIConnectionError</code>	wrong key or no internet	check <code>.env</code> ; try <code>curl https://api.openai.com</code>
Error code: 401	invalid key	regenerate key in provider's portal
Error code: 429	rate limit	wait, retry; reduce <code>max_concurrent</code>
<code>BadRequestError: model not found</code>	YAML model name doesn't match provider	check <code>config/models.yaml</code> 's <code>model:</code> field
Ollama: connection refused	<code>ollama serve</code> not running	<code>ollama serve</code> in a terminal
Ollama returns instant gibberish	model not pulled	<code>ollama pull <name></code>

D.3 Server / HTTP

Symptom	Likely cause	Fix
<code>curl: connection refused to localhost:5000</code>	server not running	<code>python -m omnillm.server.app</code>
<code>400 Bad Request: provide 'text' or 'audio'</code>	empty body	include at least one field
<code>500 Internal Server Error</code>	exception in handler	check server logs (run with <code>--debug</code>)
Server hangs on first <code>/interact</code>	LangGraph importing slowly	wait ~10 s on first call (lazy import)
Knowledge base loaded — 0 chunks	KB directory empty or wrong path	check <code>--kb</code> flag, <code>OMNILLM_KNOWLEDGE_BASE</code> env

D.4 Pepper / NAOqi

Symptom	Likely cause	Fix
<code>import naoqi: ModuleNotFoundError</code>	wrong PYTHONPATH	set <code>PYTHONPATH=C:\pynaoqi\...\lib;%PYTHONPATH%</code>
<code>import naoqi: DLL load failed</code>	missing PATH for <code>lib/</code>	also add <code>lib/</code> to PATH
Cannot connect to Pepper	wrong IP / not on same Wi-Fi	re-press chest button; ping IP
Pepper speaks but does not move	no stiffness	<code>motion.wakeUp()</code> first
<code>ALAnimatedSpeech</code> says nothing	volume zero / robot rest	check Pepper volume; <code>wakeUp</code>
Behaviour not installed	not uploaded to robot	Choregraphe → Upload to robot
Choregraphe virtual robot won't start	port 9559 in use	kill leftover NAOqi process; restart Choregraphe

D.5 RAG

Symptom	Likely cause	Fix
RAG returns empty chunks	KB not indexed	check server logs for "Knowledge base loaded — N chunks"
Faithfulness score is <code>-1.0</code>	scoring is off by default	pass <code>score_faithfulness=True</code> to <code>query()</code>
Slow RAG queries	ChromaDB persistent on slow disk	use in-memory mode (<code>persist_directory=None</code>)
Hallucination always flagged	heuristic too strict	use <code>_score_faithfulness</code> instead

D.6 Tests

Symptom	Likely cause	Fix
Tests need API keys	shouldn't — they mock	pull latest, ensure <code>pytest-mock</code> installed
<code>pytest: command not found</code>	<code>[dev]</code> extra not installed	<code>pip install -e ".[dev]"</code>
Coverage report blank	not running with <code>--cov</code>	<code>pytest tests/ --cov=omnilm</code>

Appendix E — External Resources

A curated list of papers, repositories, and documentation.

E.1 LLM Evaluation

- Hendrycks et al. (2021), *Measuring Massive Multitask Language Understanding*. The original MMLU paper.
- Zheng et al. (2023), *Judging LLM-as-a-Judge with MT-Bench and Chatbot Arena*. arXiv:2306.05685.
- White et al. (2024), *LiveBench: A Challenging, Contamination-Free LLM Benchmark*.
- Liu et al. (2024), *G-Eval: NLG Evaluation using GPT-4 with Better Human Alignment*. arXiv:2412.05579 survey.
- LMSYS Chatbot Arena: <https://chat.lmsys.org/?leaderboard>

E.2 HRI / Embodied LLMs

- Wainer et al. (2006), *The Role of Physical Embodiment in HRI*. Used in Chapter 2.
- Li (2015), *The Benefit of Being Physically Present: A Survey of Experimental Works Comparing Copresent Robots, Telepresent Robots and Virtual Agents*.
- Bartneck et al. (2009), *Measurement Instruments for the Anthropomorphism, Animacy, Likeability, Perceived Intelligence, and Perceived Safety of Robots*. The Godspeed paper.
- Irfan et al. (2024 HRI Workshop), *Between Reality and Delusion: Challenges of Applying LLMs to Social Robots*.
- Nichols et al. (2024 arXiv), *Can ChatGPT Control a Pepper Robot Adequately?*.
- Grassi et al. (2024 HAI), *ChatGPT-based Pepper Robot for Restaurant Service*.
- Spitale et al. (2024 arXiv), *Vita: An LLM-Powered Social Robot for Wellbeing*.
- Billing et al. (2024 Frontiers), *Language Models for Human-Robot Interaction*.

E.3 RAG and Vector Stores

- Lewis et al. (2020), *Retrieval-Augmented Generation for Knowledge-Intensive NLP Tasks*. The original RAG paper.
- ChromaDB: <https://www.trychroma.com>.
- LangChain: <https://python.langchain.com>.
- Sentence-Transformers: <https://www.sbert.net>.

E.4 LangGraph and Agent Frameworks

- LangGraph: <https://langchain-ai.github.io/langgraph/>.
- *LangGraph Architecture and Design*, Medium 2024.
- *LangGraph in 2026: Build Multi-Agent AI Systems*, dev.to 2026.

E.5 LiteLLM and Provider Tooling

- LiteLLM: <https://github.com/BerriAI/litellm>.
- *LiteLLM Routing & Load Balancing Documentation*: <https://docs.litellm.ai/docs/routing>.
- Ollama: <https://ollama.com>.

E.6 Pepper / NAO / Choregraphe

- SoftBank Developer Portal (NAOqi 2.5 docs): <https://developer.softbankrobotics.com>.
- Aldebaran NAOqi documentation: <http://doc.aldebaran.com/2-5/index.html>.
- *Pepper Robot NAOqi Python SDK*, ProvenRobotics.
- *Pepper Robot + ChatGPT Real-World Interactions*, BransonBots 2025.
- ROS2 NAOqi driver: https://github.com/ros-naoqi/naoqi_driver2.

E.7 Github Repositories You May Want to Browse

Repository	What
Akshita-sr/OmniLLM	This project
BerriAI/litellm	The provider gateway library
langchain-ai/langgraph	The agent graph library
chroma-core/chroma	The vector store
ollama/ollama	Local LLM runtime
openai/whisper	Speech-to-text
ilabsweden/pepperchat	Reference Pepper-LLM project (2023)
UoA-CARES/Pepper-GPT	University of Auckland, socket-based
studerus/pepper-android-realtime-chat	NAOqi 2.9 / Android, state-of-art (HRI 2026)
ros-naoqi/naoqi_driver2	ROS2 bridge for NAOqi
ros-naoqi/pepper_robot	Pepper meta-package for ROS

Appendix F — File Index

Every file in the repository, one line each.

F.1 Top-Level

File	Purpose
README.md	Project front page
GETTING_STARTED.md	Beginner setup guide
EXPLANATION.md	File-by-file explanation
ARCHITECTURE.md	Architecture reference with diagrams
PEPPER_CHOREGRAPHE_GUIDE.md	Pepper-specific walk-through
OmniLLM_Complete_Beginners_Guide.md	Choregraphe virtual-robot path
LICENSE	MIT licence text
pyproject.toml	Package metadata, dependency extras
requirements.txt	Flat dependency list
.env.example	API-key template
.gitignore	Git ignore patterns

F.2 config/

File	Purpose
config/models.yaml	The model registry (★ 19 models)
config/tasks/reasoning.yaml	Reasoning eval tasks
config/tasks/knowledge.yaml	Knowledge eval tasks
config/tasks/code.yaml	Code eval tasks
config/tasks/instruction.yaml	Instruction-following tasks
config/tasks/safety.yaml	Safety / refusal tasks
config/tasks/robot.yaml	Robot-readiness tasks

F.3 omnillm/ — Core

File	Purpose	Lines
omnillm/__init__.py	Public API re-exports	56
omnillm/gateway.py	LLMGateway, ModelResponse	271
omnillm/router.py	SmartRouter, RoutingStrategy	372
omnillm/consensus.py	ConsensusEngine, 3 strategies	409
omnillm/evaluator.py	LLM-as-Judge, 3 patterns	423
omnillm/scorer.py	EloScorer	286
omnillm/cli.py	Click + Rich terminal dashboard	549

F.4 omnillm/hri/

File	Purpose	Lines
omnillm/hri/__init__.py	Re-exports + lazy LangGraph import	73
omnillm/hri/classifier.py	T1–T4 task classifier	335
omnillm/hri/language_detector.py	Three-tier language detection	288
omnillm/hri/experiment.py	Conditions, sessions, manager	339
omnillm/hri/agent_graph.py	LangGraph 9-node pipeline	608

F.5 omnillm/rag/

File	Purpose	Lines
omnillm/rag/__init__.py	Re-exports	39
omnillm/rag/pipeline.py	RAGPipeline, ChromaDB + fallback	553

F.6 omnillm/robotics/

File	Purpose	Lines
omnillm/robotics/__init__.py	Re-exports	33
omnillm/robotics/bridge.py	Abstract bridge + RobotAction	253
omnillm/robotics/pepper.py	Pepper HTTP bridge	235
omnillm/robotics/nao.py	NAO HTTP bridge	208
omnillm/robotics/buddy.py	Buddy WebSocket bridge	246
omnillm/robotics/gesture_planner.py	Task → gesture mapping	200
omnillm/robotics/whisper_stt.py	Whisper STT (local + API)	219

F.7 omnillm/server/

File	Purpose	Lines
omnillm/server/__init__.py	Package docstring	16
omnillm/server/app.py	Flask AI server (Python 3)	437
omnillm/server/naoqi_client.py	NAOqi client (Python 2.7)	360

F.8 omnillm/tasks/

File	Purpose	Lines
omnillm/tasks/__init__.py	Re-exports	33
omnillm/tasks/loader.py	YAML / JSON task loader	183
omnillm/tasks/sample_tasks.py	8 axes of built-in tasks	349

F.9 omnillm/utils/

File	Purpose	Lines
omnillm/utils/__init__.py	Re-exports	25
omnillm/utils/cost_tracker.py	CostTracker	184
omnillm/utils/experiment_logger.py	ExperimentLogger, InteractionRecord	393
omnillm/utils/questionnaire.py	Likert + Godspeed + pairwise + observer	388
omnillm/utils/export.py	CSV / JSON / Markdown / Rich	188

F.10 knowledge_base/

File	Purpose
knowledge_base/lab_info.txt	Lab location, hours, WiFi, safety
knowledge_base/faq.txt	Frequently asked questions
knowledge_base/visitor_profiles.csv	Visitor names, roles, schedules
knowledge_base/event_schedule.csv	Seminars, deadlines
knowledge_base/research_projects.txt	Active research project descriptions
knowledge_base/university_map.txt	Building layout, directions

F.11 tests/

File	What it tests
tests/___init___.py	Empty marker
tests/test_gateway.py	LLMGateway: model strings, costs, query
tests/test_router.py	All 6 routing strategies
tests/test_evaluator.py	All 3 judge patterns
tests/test_scorer.py	ELO math, leaderboard
tests/test_consensus.py	All 3 consensus strategies
tests/test_rag.py	Indexing, retrieval, faithfulness
tests/test_hri.py	Classifier, language detector, experiment
tests/test_gesture_planner.py	Task → gesture mapping
tests/test_experiment_logger.py	InteractionRecord + filters + CSV
tests/test_new_components.py	Agent graph, Whisper, questionnaire, server

F.12 Auto-Created at Runtime

Path	What
.venv/	Virtual environment
results/	Eval JSON files, exported reports, KB persistence
__pycache__/_	Python bytecode caches

Appendix G — Field Notes: Bringing OmniLLM up on Windows 11

This appendix captures the gotchas and fixes discovered while standing OmniLLM up end-to-end on a Windows 11 laptop talking to Choregraphe's virtual Pepper. None of this is exotic — but every one of these consumed hours the first time. They're written down so future-you doesn't lose the same hours.

G.1 Choregraphe's "Connect to..." dialog gotcha

When you launch Choregraphe it auto-spawns a virtual Pepper in the background on a random port (e.g. 56471). The title bar may say *"Connected to a virtual robot"*, but **external Python clients cannot talk to that auto-spawned instance** — only Choregraphe itself can.

To make the virtual robot reachable from your scripts:

1. Top menu → **Connection** → **Connect to...**
2. **Untick** both "Use fixed port" and "Use fixed IP/hostname".
3. Click on the robot in the list (named after your user, e.g. *AKSHITA*).
4. Click **Select**.

After this the title bar says *"Connected to AKSHITA"* (or your user name). Only now can `omnillm/demo_pepper_omnillm.py/etc.` open a NAOqi session.

The port shown in the dialog changes every Choregraphe restart. Note it and pass it via `--robot-port`. The "Use fixed port" checkbox does **not** start a new server on the fixed port; it only changes the *target* of subsequent connection attempts.

G.2 Windows 11 TCP loopback bug in NAOqi 2.5

Symptom: After a successful TCP connect, the very first RPC call (`tts.say(...)`, even `tts.getLanguage()`) fails with:

```
RuntimeError: ALTextToSpeech::getLanguage Socket is not connected
```

Diagnosis: Windows 11's TCP fast-path for `127.0.0.1` interacts badly with NAOqi 2.5's qimessaging reply socket. The connection appears established at the kernel level but the qi handshake silently drops. The same code against a *real* Pepper at a LAN IP (e.g. `192.168.x.x`) works perfectly — the bug is specific to loopback on Win11.

Workaround: use the older `naoqi` `ALBroker` API with an explicit listen IP of `127.0.0.1`, instead of `qi.Application` or `qi.Session()`:

```
from naoqi import ALBroker, ALProxy

broker = ALBroker("myBroker", "127.0.0.1", 0, ROBOT_IP, ROBOT_PORT)
try:
    tts = ALProxy("ALTextToSpeech")
    tts.say("Hello")
finally:
    broker.shutdown()
```

The 2nd argument ("127.0.0.1") forces NAOqi's reply listener onto the loopback interface only. For real Pepper, set it to "0.0.0.0" instead so the robot can reach back via LAN.

scripts/pepper_demo/demo_pepper_omnillm.py chooses the right value automatically based on --robot-ip. omnillm/server/naoqi_client.py supports both modes via the --use-broker flag.

G.3 LangGraph 1.x replaces dict state instead of merging

OmniLLM uses LangGraph's StateGraph(dict) as its agent state container. In LangGraph 1.0–1.2 (the version installed today), each node's return **replaces** the entire state by default; it does not merge keys. The graph was originally written expecting accumulation behaviour, which caused every RPC to come back as {"speech": ""} even though every node ran and the LLM completed successfully.

Fix in omnillm/hri/agent_graph.py: wrap every node with a _merge_state helper so it returns {**state, **delta} instead of just delta:

```
def _merge_state(fn):
    async def wrapped(state):
        delta = await fn(state)
        return {**state, **delta} if isinstance(delta, dict) else state
    return wrapped

builder.add_node("classify_task", _merge_state(_make_classify_task_node()))
# ... wrap every node the same way
```

If you upgrade to a LangGraph version that supports Annotated reducers on the state schema, you can revert this wrapper and use the official mechanism instead.

G.4 Three module-signature mismatches inside the graph

When the graph wrapper above was applied, three latent bugs surfaced because nodes started running through to completion:

Caller	Wrong call	Correct call
_make_rag_node	rag.query(utterance, model_id=...)	rag.query(utterance)
_make_smart_router_node (C)	route_for_hri_task(task_type=..., strategy=...)	route_for_hri_task(hri_task_type=...)
_make_smart_router_node (D)	ConsensusEngine(gateway=...) + engine.query(models, msgs) + resp.synthesis	ConsensusEngine(gateway=..., config=ConsensusConfig(council_models=...)) + engine.query_council(msgs) + resp.final_answer

All three are fixed in the current source.

G.5 Defensive fallback in `app.py`

Even after G.3 and G.4, if a graph node hits an unexpected error the user still gets a usable response. `omnillm/server/app.py` now treats an empty `robot_action.speech` as "graph failed silently" and falls back to the direct gateway call:

```
action = result.get("robot_action") or {"speech": result.get("response_text", "")}
if not action.get("speech"):
    logger.warning("LangGraph returned empty speech – falling back.")
    # falls through to _fallback_interact(...)
else:
    return jsonify(action)
```

This is belt-and-braces insurance for live demos.

G.6 The `scripts/pepper_demo/` folder

Three small files live here:

File	Python	Purpose
<code>demo_pepper_omnillm.py</code>	2.7	Default mode: full AI demo (server → Pepper). With <code>--check-only</code> : bridge sanity test only.
<code>test_all_conditions.py</code>	3.x	Batch tests all 5 experimental conditions + multilingual against the AI server. Writes <code>test_results.txt</code> .
<code>README.md</code>	—	How to run everything.

These do not depend on any code in `omnillm/` other than the live HTTP endpoint at `localhost:5000`. They are deliberately self-contained so a demo-day failure in one part doesn't break the others.

G.7 Always invoke the right Python interpreter

On Windows, `python` on PATH is whatever interpreter the user installed most recently — often Python 3. The NAOqi side **must** run under Python 2.7 (NAOqi 2.5 was never ported). Always invoke it by full path:

```
C:\Python27\python.exe scripts\pepper_demo\demo_pepper_omnillm.py
```

The Python 3 side (`python -m omnillm.server.app`) should be run from the project venv where `pip install -e .` has been executed.

Closing Note

You have reached the end of the book. By now you should be able to:

- **Explain** in plain English what OmniLLM is, why it exists, and what the robot adds (Parts I, II, VII).
- **Read** any source file and know what each function does without guessing (Part III).
- **Connect** to a Pepper robot (or simulate one in Choregraphe) and drive it through the AI server (Part IV, V).
- **Re-build** the project from scratch in a sensible order if you ever needed to (Part VI).
- **Run** an experimental session, collect the data, and analyse it (Chapter 32, 38).
- **Look up** any term, command, or file in seconds (Appendices).

The codebase will continue to evolve; this book will not. When the book and the source disagree, **trust the source**. The book's job was to get you fluent enough that you can read the source for yourself.


Good luck with the thesis. And remember that the most interesting result might be the one you didn't predict — Embodied LLM rankings might agree with text rankings, or they might not. Both findings advance the field.

— *End of book.*

Appendix G — End-to-End Flow Diagram (Spoken Question to Robot Action)


This appendix consolidates everything in the book into a single annotated diagram. Use it as a poster. The flow shows what happens when a person speaks to Pepper, from the audio captured by the microphone to the speech, gesture and LED change that the robot produces in reply — including how the Smart Router, the Consensus Council, the LLM-as-Judge and the ELO leaderboard all plug in.

G.1 Stage 0 — Person Speaks to Pepper



G.2 Stage 1 — LangGraph Pipeline `omnillm/hri/agent_graph.py`

The graph is compiled by `build_hri_graph()`. Each box below is one node; each node is an `async` function that reads state and returns state updates.



```

+-----+
v
[smart_router] ← only acts for conditions C and D
+-----+
| condition == "C" --> omnillm/router.py SmartRouter |
| route_for_hri_task(task_type) |
| strategy = TASK_TYPE |
| Picks ONE best model from config/models.yaml |
| Then gateway.query(model, ...) |
| |
| condition == "D" --> omnillm/consensus.py ConsensusEngine |
| Council = [gpt-4o-mini, gemini-2.5-flash, claude-haiku] |
| gateway.query_multiple(...) (parallel fan-out) |
| strategy = "synthesis" (default) |
+-----+

|
v
[generate_action_plan] -- omnillm/robotics/gesture_planner.py
| Builds RobotAction dict:
| { speech, gesture, emotion_led, metadata }
v
[log_interaction] -- omnillm/utils/experiment_logger.py
| Writes JSONL row
| (latency, tokens, $, judge_score)
v
[END] -- HTTP 200 RobotAction JSON --> Pepper
executes speech + gesture + LED

```

G.3 Stage 2 — Smart Routing `omnillm/router.py :: SmartRouter`

```

Inputs: task_category, budget_usd?, max_latency_ms?, strategy
Config: config/models.yaml
          (cost_per_1m_input/output, latency, hri_strengths)
Memory: past EvalResults loaded via _load_results()
          → quality / latency table

+----- candidates = all models -----+
| filter by budget_usd cap |
| filter by max_latency_ms cap |
| if empty → routing.fallback_models |
+-----+-----+
      v
+----- pick strategy -----+
| BEST_QUALITY → max quality |
| LOWEST_COST → min $/query |
| LOWEST_LATENCY → min ms |
| BEST_VALUE → 0.5·Q + 0.3·C + 0.2·L (composite) |
| LOCAL_PREFERRED → local-only first |
| TASK_TYPE → hri_task_routing[task] (used by Pepper) |
+-----+-----+
      v
RouteDecision{ model_id, confidence,
               estimated_cost,
               estimated_latency_ms }

```

G.4 Stage 3 — Consensus / LLM Council [omnilm/consensus.py](#)

```
user prompt
|
v
gateway.query_multiple(council_models, messages) ← fan-out, concurrent
|
+--> model 1 (e.g. gpt-4o-mini) -+
+--> model 2 (e.g. gemini-2.5) -+ independent answers
+--> model N (e.g. claude-haiku) -+
|
v pick strategy
+-----+
| "majority_vote" → Jaccard-cluster the texts, |
|                 return rep of largest cluster. |
| "weighted"     → static weights per model, |
|                 pick highest-weighted answer. |
| "synthesis"    <-- DEFAULT |
| Build a synthesis prompt that contains all council |
| answers and ask the JUDGE LLM to merge them into |
| the BEST combined answer. |
| Returns JSON: |
| { final_answer, agreement_score, |
|   reasoning, dissenting_models } |
+-----+
|
v
ConsensusResult{ final_answer,
                 agreement_score (0-1),
                 dissenting_models,
                 synthesis_reasoning }
```

G.5 Stage 4 — LLM-as-Judge [omnilm/evaluator.py](#)

Who is the judge? Default `judge_model = "openai-gpt4o"` — i.e. **ChatGPT (GPT-4o)**. It is set in `consensus.py` (`ConsensusConfig.judge_model`), `evaluator.py` (`Evaluator(judge_model=...)`), and overridable from the CLI via `--judge`. The shipped default is OpenAI's GPT-4o, but you can swap in Claude, Gemini, or any registered model.

Three judge patterns:

```
response --> +-----+
| 1) referenceless (G-Eval) |
| Judge scores the answer on its own. |
| → score ∈ [0,1] + reasoning |
+-----+
| 2) reference_based |
| Judge compares answer vs gold answer. |
| → score ∈ [0,1] + reasoning |
+-----+
| 3) pairwise (used to feed ELO) |
| Judge sees model A vs model B. |
| Position-swap to remove bias. |
| → winner ∈ {model_a, model_b, tie} |
+-----+
|
v
EvalResult / PairwiseResult
```

G.6 Stage 5 — ELO Scoring `omnilm/scorer.py :: EloScorer`

Pairwise judge verdicts feed here. Same method LMSYS Chatbot Arena uses.

```
record_match(model_a, model_b, winner, category)
|
|   Each model starts at default_rating = 1500
|
v
Expected score:
    E_A =  $\frac{1}{1 + 10^{((R_B - R_A) / 400)}}$ 
    E_B = 1 - E_A
|
v
Actual score S:   win = 1.0   loss = 0.0   tie = 0.5
|
v
Update rule (K = 32):
    R_A' = R_A + K * (S_A - E_A)
    R_B' = R_B + K * (S_B - E_B)
|
v
Stored in:
    ratings[model]           ← overall ELO
    _category_ratings[cat][model] ← per-category ELO
    _history (every match)   ← audit trail
    _rating_history[model]   ← time-series for plots
|
v
get_leaderboard()           → overall rank
get_category_leaderboard(cat) → reasoning / code / robot / etc.
save() / load() → JSON on disk → fed BACK into
                               SmartRouter._load_results()
```

G.7 The Feedback Loop (the whole point of the system)

```
Pepper question --> LangGraph --> answer to user
|
|   +--(council/consensus runs)--> judge LLM
|   |
|   |   +--> Evaluator (pairwise) --> EloScorer.record_match
|   |   |
|   |   |   results/*.json
|   |   |   |
|   |   |   |   SmartRouter._load_results()
|   |   |   |   |
|   |   |   |   |   v
|   |   |   |   |   Next question is routed to a model whose ELO / quality for
|   |   |   |   |   that task category has gone UP from past wins.
```

Short answer to the question "who is the LLM judge in this project?": the default judge is **GPT-4o (ChatGPT)** — `judge_model = "openai-gpt4o"` in `consensus.py:54` and `evaluator.py:117`. It is configurable (pass Claude, Gemini, etc. via `--judge` or `ConsensusConfig.judge_model`), but the shipped default is OpenAI's GPT-4o.

PART IX — PEPPER-LLM INTEGRATION SURVEY

This part is a **standalone technical reference** on how Pepper, NAOqi, Choregraphe and modern LLMs fit together. It consolidates the `compass_artifact_*.md` research notes that informed OmniLLM's design. If Parts I–VIII told you what OmniLLM is, Part IX tells you what Pepper is — the hardware, the middleware, and the catalogue of published projects that tried this before you.

Chapter 40 — Pepper's Hardware and Software Platform

AT A GLANCE. Pepper is a 120 cm, 28 kg, 20-DoF semi-humanoid social robot from SoftBank Robotics (originally Aldebaran). 20 degrees of freedom, Intel Atom CPU, NAOqi OS on top of Gentoo Linux, Python 2.7 only. The hardware is interesting; the software is dated; both are what you'll work with.

40.1 Physical Specs at a Glance

Attribute	Value
Height	120 cm
Weight	28 kg
Degrees of freedom	20 (2 head, 2×6 arm/hand, 2 hip, 1 knee, 3 omnidirectional base wheels)
Battery	30 Ah / 795 Wh lithium-ion (~8–10 hours active)
Onboard CPU	Intel Atom E3845 quad-core @ 1.91 GHz
Onboard RAM	4 GB DDR3
Onboard storage	8 GB flash + microSD slot
Operating system	NAOqi OS (modified Gentoo Linux)
Onboard Python	2.7

40.2 Sensor Suite

- **2 × 5 MP RGB cameras** (forehead and chin, Omnivision OV5640)
- **ASUS Xtion 3D depth sensor** (320×240, 0.4–8 m range) behind the eyes
- **4 directional microphones** on the head for sound localisation
- **2 speakers** in the ears
- **2 sonar sensors** in the base
- **6 laser line generators** in the base
- **2 infrared sensors** in the base
- **3 bumper contact sensors** in the base
- **3 head capacitive touch sensors** and **2 hand touch sensors**
- **Gyroscope and accelerometer IMUs** in both torso and base

40.3 The Chest Tablet

The chest-mounted **10.1-inch IPS capacitive touchscreen** runs at **1280×800 resolution** with 5-point multi-touch. It is powered by a separate **1.3 GHz quad-core ARM Cortex-A7** with **1 GB RAM** and **32 GB storage** running Android. The tablet communicates with the robot's head computer over an internal network at IP **198.18.0.1**. From NAOqi's perspective it is exposed via the `ALTabletService` module.

40.4 Two OS Variants — and Which to Use

Version	SDK	Choregraphe	API surface	Pepper status
NAOqi 2.5	Python/C++	Yes	~1,000+ APIs	Standard for most Pepper units in research labs
NAOqi 2.9	Android + QiSDK (Java/Kotlin)	No	~20 high-level APIs	Newer "Pepper 1.8" units; Choregraphe-incompatible

For LLM integration work using Python, **NAOqi 2.5 is the de-facto standard**. OmniLLM targets this version.

40.5 The End-of-Life Situation

Aldebaran filed for bankruptcy in **February 2025**. Its IP was acquired by **Maxvision Technology Corp. (Shenzhen)** in **July 2025**, and **no new units are being manufactured**. Existing units continue to work; spare parts and SDK downloads are increasingly hard to source.

The implication for OmniLLM: the abstract `RobotBridge` (Chapter 19) is deliberately platform-portable. When (not if) you have to retarget NAO or Buddy or any future robot, only the bridge implementation changes — the AI server doesn't.

Chapter 41 — NAOqi Framework, PyNAOqi, and Programmatic Connection

AT A GLANCE. NAOqi is the **middleware broker** running on TCP port **9559**. Every module — `ALTextToSpeech`, `ALMotion`, `ALMemory`, and dozens more — registers its methods with the broker, enabling both local procedure calls (in-process, zero-copy) and remote procedure calls (over TCP, serialised). PyNAOqi is the Python 2.7 binding that exposes this entire API.

41.1 Installing PyNAOqi

PyNAOqi is **not pip-installable**. It must be downloaded manually from the SoftBank Robotics developer portal as a platform-specific archive such as `pynaoqi-python2.7-2.5.5.5-linux64.tar.gz` (Linux) or `pynaoqi-python2.7-2.5.5.5-win32-vs2013.zip` (Windows).

Installation steps:

1. Extract the archive (e.g., to `C:\pynaoqi\` on Windows).
2. Set `PYTHONPATH` to include the SDK's `lib/python2.7/site-packages` directory.
3. On Windows, **32-bit Python 2.7 is required** to match the 32-bit native bindings, plus the **MSVC++ 2010 x86 redistributable**.

```
# Linux
export PYTHONPATH="$HOME/pynaoqi/lib/python2.7/site-packages:$PYTHONPATH"

# Windows CMD
set PYTHONPATH=C:\pynaoqi\pynaoqi-python2.7-2.5.5.5-win32-vs2013\lib;%PYTHONPATH%
set PATH=C:\pynaoqi\pynaoqi-python2.7-2.5.5.5-win32-vs2013\lib;%PATH%
```

41.2 Two Connection Paradigms

The **legacy ALProxy approach** creates direct proxies — one proxy per service. Simple and explicit:

```
from naoqi import ALProxy
tts = ALProxy("ALTextToSpeech", "192.168.1.100", 9559)
tts.say("Hello from ALProxy!")
```

The **modern qi.Session approach** uses a shared session object — recommended for any non-trivial application because it shares one TCP connection across all services:

```
import qi
session = qi.Session()
session.connect("tcp://192.168.1.100:9559")
tts = session.service("ALTextToSpeech")
tts.say("Hello from qi.Session!")
```

OmniLLM's `naoqi_client.py` (Chapter 21) tries `qi` first and falls back to `ALProxy` so it works against both old and new SDK versions.

Chapter 42 — The Eleven NAOqi Modules That Matter for LLM Integration

These are the modules you'll touch. Master these eleven and you can build nearly any Pepper-LLM application.

42.1 Speech

ALTextToSpeech — converts text to speech on the robot's speakers; supports speed and pitch parameters and embedded markup tags like `\\RSPD=80\\` (speed 80 %) and `\\VCT=110\\` (voice pitch 110 %).

ALAnimatedSpeech — speech synchronised with gestures. Annotation tags such as `^start(animations/Stand/Gestures/Hey_1)` trigger concurrent animations. The configuration option `bodyLanguageMode` accepts "contextual" (gestures match the speech content automatically, recommended), "random", or "disabled".

```
tts = session.service("ALAnimatedSpeech")
tts.say("Welcome to the lab!", {"bodyLanguageMode": "contextual"})
```

42.2 Motion and Posture

ALMotion — joint-level motor control. `wakeUp()` / `rest()` toggle stiffness; `setAngles()` for non-blocking moves; `angleInterpolation()` for timed trajectories. **You must call `wakeUp()` before any movement** or the joints have zero stiffness and the robot ignores commands.

ALRobotPosture — whole-body posture management. `goToPosture("Stand", 0.8)` returns Pepper to a standard stance with 80 % speed.

42.3 Audio

ALAudioRecorder — records microphone audio to WAV/OGG files on the robot filesystem. Simpler API but adds a file-transfer step.

ALAudioDevice — real-time audio streaming via the `processRemote` callback. Supports 16 kHz or 48 kHz, 1 or 4 channels. Subclass `ALModule`, implement `processRemote(nbChannels, nbSamples, timeStamp, inputBuffer)`, then `subscribe()` to receive raw int16 samples in real time.

OmniLLM's audio path uses **ALAudioDevice** because the latency budget demands streaming, not file-transfer.

42.4 Speech Recognition (Spoiler: Use Whisper Instead)

ALSpeechRecognition — on-board keyword recognition with a predefined vocabulary; results stored in `ALMemory` under "WordRecognized". **Universally considered inadequate for open-domain speech.** All modern Pepper-LLM projects bypass it and use OpenAI Whisper (locally via `openai-whisper` or via the API) instead. OmniLLM follows this pattern.

42.5 Events and Memory

ALMemory — the central key-value store and event bus. `subscriber("EventName").signal.connect(callback)` enables reactive programming. Almost every NAOqi module publishes events to **ALMemory**, so this is where you wire perception to reaction.

```
memory = session.service("ALMemory")
def on_touch(value):
    print("Head touched:", value)
sub = memory.subscriber("FrontTactilTouched")
sub.signal.connect(on_touch)
```

42.6 Vision

ALFaceDetection — OMRON-based face detection. Subscribes to the `"FaceDetected"` event with face position, size, and optional recognition data. Used by OmniLLM to trigger interaction start when someone walks into Pepper's field of view.

42.7 Output Devices

ALLeds — RGB control of eye, ear, chest and foot LEDs.

```
leds = session.service("ALLeds")
leds.fadeRGB("FaceLeds", 0.0, 0.67, 1.0, 0.5) # navigation blue, 500ms fade
```

ALTabletService — controls the chest tablet display (Pepper only; NAO has no tablet). `loadUrl()`, `showImage()`, `executeJS()`.

```
tablet = session.service("ALTabletService")
tablet.loadUrl("http://192.168.1.50:5000/map?room=305")
tablet.showWebview()
```

42.8 Behaviours

ALBehaviorManager — lists, runs and stops Choregraphe behaviours installed on the robot.

```
bm = session.service("ALBehaviorManager")
bm.runBehavior("animations/Stand/Gestures/Hey_1") # wave
```

OmniLLM maps each gesture name in `GESTURE_TO_BEHAVIOR` (Chapter 21) to the behaviour path that NAOqi expects.

Chapter 43 — Choregraphe as a Visual Programming IDE

AT A GLANCE. *Choregraphe is SoftBank's drag-and-drop desktop IDE for programming Pepper and NAO. It is locked to **NAOqi 2.5**. It includes a virtual robot simulator that runs without any hardware, plus a Python 2.7 script editor for inline testing. For OmniLLM you use it for three things: testing behaviours without a real robot, installing custom animations, and live monitoring during experiments.*

43.1 The Four-Panel Layout

BOX LIBRARY (left)	FLOW DIAGRAM (center)	3D ROBOT VIEW (right)
Drag boxes from here	Wire boxes together to build behaviour	Virtual or real Pepper reflected in real time
LOG VIEWER / SCRIPT EDITOR (bottom) NAOqi log messages, Python 2.7 inline scripting		

43.2 When to Use Choregraphe vs OmniLLM

You want to ...	Choregraphe	OmniLLM
Test Pepper's speech	yes	—
Test a single gesture	yes	—
Build / tune a custom animation	yes	—
Build an AI-driven conversation	—	yes
Use multiple LLMs as backends	—	yes
Run a controlled HRI experiment	—	yes
Live monitor an experiment	yes (alongside)	yes

They are complementary. The typical workflow:

1. **Plan gestures in Choregraphe** — drag and edit until they look natural.
2. **Upload to Pepper** — File → Build Application Package → upload.
3. **Add the gesture names to OmniLLM's GESTURE_TO_BEHAVIOR map.**
4. **Run the OmniLLM experiment** — leave Choregraphe open as a monitor.

43.3 Connecting Choregraphe

Virtual robot (no hardware needed):

1. Open Choregraphe.
2. Connection → Connect to virtual robot.
3. Pepper appears in the 3D view; status bar shows `Connected to localhost:9559 (virtual)`.

Physical robot:

1. Same Wi-Fi as Pepper.
2. Press Pepper's chest button — it says its IP.
3. Connection → Connect to ..., enter the IP, port 9559.
4. The 3D view mirrors the real robot's joint positions.

43.4 The Script Editor Shortcut

Press `Alt+5` to open the Script Editor at the bottom and type Python 2.7 directly:

```
# Test ALAnimatedSpeech
tts = ALProxy("ALAnimatedSpeech", "localhost", 9559)
tts.say("OmniLLM is ready.", {"bodyLanguageMode": "contextual"})

# Test a gesture
b = ALProxy("ALBehaviorManager", "localhost", 9559)
b.runBehavior("animations/Stand/Gestures/Explain_8") # point left

# Set LEDs blue
leds = ALProxy("ALLeds", "localhost", 9559)
leds.fadeRGB("FaceLeds", 0.0, 0.67, 1.0, 0.3)
```

This is the fastest way to verify NAOqi services are healthy before debugging OmniLLM.

Chapter 44 — Programming Patterns: Audio, Tablet, Face Detection

44.1 Audio Capture for STT

There are three ways to get audio off Pepper for transcription:

Pattern	Latency	Complexity	When to use
ALAudioRecorder + SCP	2–5 s	Low	Quick demos; OK for batch
ALAudioDevice.subscribe + callback	<1 s	High	Real-time HRI (OmniLLM uses this)
External USB microphone + Whisper API	<0.5 s	Medium	Highest quality, bypasses Pepper mics

Many published projects (Chapter 46) chose pattern 3 because Pepper's on-board microphones, while four-channel, are noise-prone in busy public settings.

44.2 Tablet Patterns

The chest tablet is an Android web view. Three usage patterns:

1. **Static image** — `showImage(url)` to display a map or photo.
2. **HTML page** — `loadUrl(url)` to load a Flask-served page; useful for showing retrieved RAG results.
3. **JavaScript bridge** — `executeJS()` lets the LLM-server push live updates without reloading the page.

For OmniLLM, the natural pattern is: AI server stores rendered HTML at `/map/<room>`; the `RobotAction`'s `metadata.tablet_url` is read by the NAOqi client and passed to `loadUrl()`.

44.3 Face Detection as Interaction Trigger

```
faces = session.service("ALFaceDetection")
faces.subscribe("OmniLLMTrigger", 500, 0.0) # 500ms refresh

mem = session.service("ALMemory")
def on_face_seen(value):
    if value: # value is [TimestampedFloatArray, FaceInfoArray]
        omnillm_client.greet()
sub = mem.subscriber("FaceDetected")
sub.signal.connect(on_face_seen)
```

This snippet makes Pepper wait passively and only greet visitors when their face appears.

Chapter 45 — The Python 2.7 ↔ 3.x Bridge: Five Patterns Reviewed

AT A GLANCE. Modern AI libraries (LiteLLM, LangGraph, ChromaDB, openai-whisper) require Python 3.11+. NAOqi requires Python 2.7. They cannot share a process. The published literature converges on five patterns to bridge them.

45.1 Pattern 1 — HTTP / REST Bridge (OmniLLM's Choice)

Two processes: Python 2.7 NAOqi client and Python 3.x Flask server. Communicate over JSON-over-HTTP.

Used in: [ilabsweden/pepperchat](#) (2023), Frontiers ASD therapy project (Billing et al., 2024), 6+ others.

Trade-offs: Easiest to debug (every message is curl-able), most documented, ~50–200 ms HTTP overhead per turn. **This is the OmniLLM default.**

45.2 Pattern 2 — Socket-Based Bridge

Raw TCP socket between Python 2.7 and Python 3.x, often with a custom binary protocol or newline-delimited JSON.

Used in: Pepper-GPT (University of Auckland), Ghent University elder care project.

Trade-offs: Lower latency (~10–50 ms), but more code to write (framing, reconnection, error handling). Worth it for token-streaming applications.

45.3 Pattern 3 — ROS2 Bridge via `naoqi_driver2`

The robot publishes audio and sensor topics over ROS2; a Python 3.x node subscribes and dispatches to LLMs. Action plans are published back to a ROS2 topic the robot subscribes to.

Used in: Multi-robot research projects, EU H2020 projects.

Trade-offs: Powerful when many components publish/subscribe (e.g., SLAM, Nav2 navigation, multi-robot coordination), but high setup complexity for a single-robot study.

45.4 Pattern 4 — MQTT Broker

Both Python 2.7 and Python 3.x publish/subscribe to topics on an MQTT broker (e.g., Eclipse Mosquitto).

Used in: LAIR-GPT (University of Ancona).

Trade-offs: Good when many components need to publish/subscribe independently; adds a broker as a third process.

45.5 Pattern 5 — Python 3 qi 3.1.5 (Single Process, Risky)

A community-built `qi 3.1.5` package exists on PyPI for Linux x86_64 Python 3. It exposes the NAOqi API directly in Python 3 — no bridge needed.

Used in: Prototypes, hobbyist projects.

Trade-offs: Several services are **broken** in this port (touch detection, audio callbacks, certain event subscriptions). Unsuitable for production. OmniLLM does not use this.

45.6 Why OmniLLM Picked Pattern 1

Criterion	HTTP/REST	Socket	ROS2	MQTT	qi 3.1.5
Debuggability	best	good	poor	medium	medium
Latency	OK	best	OK	OK	best
Lines of code	low	medium	high	medium	low
Cross-platform	best	best	poor	best	poor
Survives provider quirks	best	best	OK	best	poor
Documented in literature	best	good	good	medium	poor

For a typical Pepper + LLM HRI study, the HTTP bridge is the sweet spot. If you ever need sub-100 ms end-to-end (token-streaming TTS), switch to socket. If you ever need ROS2 navigation, add `naoqi_driver2` alongside.

Chapter 46 — Catalogue of Published Pepper-LLM Projects (2023–2026)

A survey of 15+ projects that connect Pepper (or similar SoftBank robots) to LLMs. Each entry shows what they did, the architecture pattern, and the lesson OmniLLM learnt from it.

46.1 Academic Publications

Project	Year	LLM Used	Bridge Pattern	Key Lesson
Irfan et al., HRI 2024 — <i>Between Reality and Delusion</i>	2024	ChatGPT	HTTP REST	Single LLM is fragile; multi-model would be more robust
Nichols et al., arXiv 2024 — <i>Can ChatGPT Control a Pepper Robot Adequately?</i>	2024	GPT-4	Socket	Latency is the dominant UX variable
Grassi et al., HAI 2024 — <i>ChatGPT-based Pepper Robot for Restaurant Service</i>	2024	ChatGPT	HTTP REST	Domain-specific RAG dramatically improves perceived accuracy
Spitale et al., arXiv 2024 — <i>Vita: An LLM-Powered Social Robot for Wellbeing</i>	2024	Custom	HTTP REST	Empathy and warmth need a different system prompt than factual tasks
Billing et al., Frontiers 2024 — <i>Language Models for HRI</i>	2024	Multiple (single-model tests)	HTTP REST	First survey paper acknowledging multi-LLM as a research gap
Pepper-GPT (Auckland)	2023	GPT-3.5	Socket	Socket latency advantage matters in face-to-face conversation
ilabsweden/pepperchat	2023	GPT-3.5	HTTP REST	Public reference implementation of the HTTP bridge pattern

LAIR-GPT (Ancona)	2023	GPT-4	MQTT	MQTT useful when many components subscribe
Frontiers ASD therapy	2024	GPT-4	HTTP REST	RAG over therapy protocol PDFs increased clinician trust

46.2 GitHub Repositories

Repository	What It Contains
Awesome-LLM-Robotics	Curated collection of 200+ papers on LLMs + robotics
microsoft/PromptCraft-Robotics	Microsoft's framework for ChatGPT + robot control
BerriAI/litellm	Unified API gateway for 100+ LLM providers (used by OmniLLM)
langchain-ai/langchain	RAG pipelines, document loaders, chains
langchain-ai/langgraph	Stateful agent graphs with conditional routing (used by OmniLLM)
softbankrobotics-labs/*	Official Pepper/NAO community examples
pepper-chatgpt-integration	Various community projects connecting Pepper to ChatGPT
Awesome-LLM-Ensemble	Multi-model consensus and ensemble methods
chroma-core/chroma	Vector store for RAG (used by OmniLLM)
lm-sys/FastChat	LMSYS Chatbot Arena — human preference ELO for LLMs

46.3 What Every Project Got Wrong (and How OmniLLM Avoids It)

Common mistake	OmniLLM's fix
Lock the robot to one cloud provider	YAML model registry — swap providers in 7 lines
Hard-code prompts in NAOqi Python 2.7	All prompting lives on the Python 3 side; Python 2.7 is a thin client
Skip evaluation, just demo	Built-in LLM-as-Judge, ELO scorer, automatic logging
Ignore latency	<code>latency_ms</code> recorded in every <code>ModelResponse</code>
Ignore cost	<code>cost_usd</code> tracked per-model, per-session
No multilingual	Three-tier language detector + per-language model mapping
RAG bolted on as an afterthought	RAG is a first-class LangGraph node with faithfulness scoring
No reproducibility	All conditions encoded as <code>ConditionConfig</code> ; one YAML controls everything

Chapter 47 — Lessons Learned and Best Practices

47.1 Hardware-Level Lessons

1. **Always `wakeup()` before motion**, or the robot ignores you.
2. **Use `ALAnimatedSpeech`, not `ALTextToSpeech`** — embodiment effect research consistently shows talking-head robots rate lower on naturalness and intelligence.
3. **Plug in for long sessions.** Pepper can charge while operating; sessions exceeding 2 hours run the battery flat.
4. **Bypass on-board speech recognition.** Use external Whisper.
5. **The tablet's IP is `198.18.0.1`** — internal, not your LAN.

47.2 Software-Level Lessons

1. **Make the AI server bulletproof in text mode first.** Then add audio. Then add the robot. Most debugging time is spent on non-AI failures (Wi-Fi, IP addresses, NAOqi stiffness).
2. **Log everything from day 1.** You cannot reconstruct latencies or tokens after the fact.
3. **Mock LLM calls in tests.** Real API calls in tests are slow, expensive, non-deterministic and brittle.
4. **Cache the LangGraph compile.** First call takes ~10 s; subsequent calls are instant.
5. **Use `asyncio.gather`**, not for `... await ...`, when querying many models. 1 second instead of N seconds.

47.3 Experimental-Design Lessons

1. **Within-subjects design** is statistically more powerful than between-subjects for the same N. Every participant should experience multiple conditions.
2. **Counterbalance with a Latin square** to cancel order effects.
3. **Pilot with at least 3 colleagues** before recruiting real participants — you will find bugs you couldn't anticipate.
4. **Back up data daily.** Losing a participant's session after they left is the worst-case scenario.
5. **Ask a single pairwise preference question** at the end. It produces clean ELO matches and respects the participant's time.

47.4 Communication-Level Lessons

1. **Show participants the robot before the experiment.** Curiosity-driven exploration ruins the first interaction otherwise.
2. **Frame the consent form around "studying Pepper, not you"** to reduce performance anxiety.
3. **Have a fallback question** ready when STT fails — "I didn't catch that, could you repeat?" — so a misrecognition doesn't kill the flow.

PART X — THESIS DESIGN REFERENCE

*This part is the **research design half** of the book. It consolidates the `thesis_guidelines.pdf` material — the literature you must cite, the gap you fill, the three hypotheses, the four task types in full detail, the five experimental conditions, the metrics, and the one-month execution plan. If Part VII gave you the high-level "what", Part X gives you the academic framing and the practical schedule.*

Chapter 48 — Literature Landscape: What Exists, What's Missing

AT A GLANCE. Five 2024 publications connect LLMs to social robots. Every one of them uses a **single** LLM. Every one identifies multi-LLM comparison as a gap. OmniLLM fills it.

48.1 Key Scientific Papers

Paper	Venue	What They Did	Gap Remaining
Irfan et al. (2024) — <i>Between Reality and Delusion: Challenges of Applying LLMs to Social Robots</i>	HRI 2024 Workshop	Identified challenges of integrating ChatGPT with Pepper; single-LLM only.	No multi-LLM comparison, no smart routing, no RAG.
Nichols et al. (2024) — <i>Can ChatGPT Control a Pepper Robot Adequately?</i>	arXiv	Connected GPT-4 to Pepper for dialogue; evaluated a single model.	Only one LLM tested, no systematic comparison, no user study with multiple backends.
Grassi et al. (2024) — <i>ChatGPT-based Pepper Robot for Restaurant Service</i>	HAI 2024	Pepper as a restaurant assistant using ChatGPT.	Single LLM, no routing, no domain-specific RAG.
Spitale et al. (2024) — <i>Vita: An LLM-Powered Social Robot for Wellbeing</i>	arXiv	Pepper-like robot for mental wellbeing conversations using LLM.	Single LLM, no model comparison, no evaluation framework.
Billing et al. (2024) — <i>Language Models for Human-Robot Interaction</i>	Frontiers in Robotics & AI	Survey paper on LLM-HRI.	Calls out multi-LLM comparison as an open problem.

48.2 Key GitHub Repositories and Resources

Repository / Resource	What It Contains
Awesome-LLM-Robotics	Curated collection of 200+ papers on LLMs + robotics
PromptCraft-Robotics	Microsoft's framework for ChatGPT + robot control
LiteLLM	Unified API gateway for 100+ LLM providers
LangChain	RAG pipelines, document loaders, chains
LangGraph	Stateful agent graphs with conditional routing
SoftBank Robotics Labs	Official Pepper/NAO community examples
pepper-chatgpt-integration	Various community projects connecting Pepper to ChatGPT
Awesome-LLM-Ensemble	Multi-model consensus and ensemble methods
ChromaDB	Vector store for RAG
LMSYS Chatbot Arena	Human preference ELO leaderboard for LLMs

48.3 Blog Posts and Practical Resources

- [HuggingFace Blog: "DeepSeek R1 Controls Robotic Arm"](#) — open-source LLM controlling a physical robot.
- [Latitude AI Blog: "LLM Routing Cost Reduction in RAG"](#) — 75 % cost savings with routing.
- [SoftBank Developer Portal](#): NAOqi SDK documentation for Pepper.
- [Ollama Blog](#): Running local models for robotics applications.

Chapter 49 — The Critical Research Gap You Will Fill

49.1 What No One Has Done

- **No study** has systematically compared multiple LLMs as interchangeable backends for a social robot through human experiments.
- **No work** has applied smart routing between LLMs in an embodied social robot context.
- **No research** has used RAG-augmented social robots to compare LLM performance on domain-specific knowledge tasks with real users.
- **No study** has compared LLM rankings derived from embodied HRI with standard text-only benchmark rankings.
- **No work** has combined LangGraph-based agent orchestration with social robot interaction management.

49.2 Why This Matters

Current LLM evaluations (MMLU, Chatbot Arena, LiveBench) are **entirely text-based**. Social robot studies that use LLMs connect **only a single model** (typically ChatGPT). No prior study has:

1. compared multiple LLMs as interchangeable backends for a social robot;
2. applied dynamic smart routing between models during live HRI; or
3. evaluated whether LLM rankings change when evaluated through embodied interaction rather than text-only benchmarks.

Furthermore, **no RAG-augmented social robot has been tested with multiple LLM backends** to compare domain-specific knowledge delivery performance.

49.3 The Recommended Thesis Title

"Embodied LLM Arena: Multi-Model Smart Routing and RAG-Augmented Knowledge Grounding for Task-Based Human-Robot Interaction with Pepper"

This title captures all four contributions: *Embodied LLM Arena* (the methodology), *Multi-Model* (the comparison), *Smart Routing* (the routing study), *RAG-Augmented* (the grounding study), *Task-Based* (the T1–T4 design), *Human-Robot Interaction with Pepper* (the platform).

Chapter 50 — Three Hypotheses

AT A GLANCE. Three hypotheses, in falsifiable form, each tied to the data OmniLLM logs automatically. If you can rebuild the book from the source code (you can — see `book/build_pdf.py`), you can also derive every statistical test from the logged JSON (you can — see `omnilm/utils/experiment_logger.py`).

50.1 H1 — Embodied vs Text Rankings Diverge

H1: LLM quality rankings obtained through embodied human-robot interaction will differ significantly from text-only benchmark rankings.

Why we expect this: physical presence, latency sensitivity, and multimodal feedback change how humans perceive response quality. A model that wins MMLU may lose on naturalness; a model that scores poorly on coding benchmarks may score highly on warmth in a navigation task.

How to test: correlate the Embodied LLM Leaderboard (Chapter 13's `EloScorer.get_category_leaderboard("embodied_hri")`) with each model's public Chatbot Arena ELO. Spearman's ρ is the right statistic; a low or negative ρ supports H1.

50.2 H2 — Smart Routing Beats Any Fixed Model

H2: Dynamic smart routing between LLMs based on task type and language will produce higher user satisfaction than any single fixed-model configuration.

Why we expect this: different models excel at different tasks (claude-haiku is natural in chat, gemini-flash is fast on navigation, gpt-4o-mini is accurate on RAG). Picking the right one per task should dominate the best single model.

How to test: compare mean Likert scores across all five conditions A–E (Chapter 52). A repeated-measures ANOVA, with condition C (smart routed) significantly above A (cloud baseline) and B (local baseline), supports H2.

50.3 H3 — RAG Lifts Every Backend

H3: RAG-augmented responses will be rated significantly more accurate and trustworthy than non-RAG responses across all LLM backends.

Why we expect this: RAG grounds answers in your lab's actual documents (WiFi password, opening hours, room locations) which no LLM could have seen during training.

How to test: compare condition A (GPT-4o-mini + RAG) against condition E (GPT-4o-mini, RAG off). The pair holds the model constant and isolates RAG. A paired t-test on the "I trust the information" Likert item, with $A > E$, supports H3.

50.4 Sample Size and Power

For a within-subjects design with 3 conditions per participant, a medium effect size (Cohen's $d \approx 0.5$) and $\alpha = 0.05$, $N \approx 15-25$ participants yields ~80 % statistical power for the main contrasts. This is the standard sample for an HRI study of this design.

Chapter 51 — The Four Task Types in Full Detail

OmniLLM's classifier (`omnillm/hri/classifier.py`) labels every utterance T1, T2, T3 or T4. The LangGraph router branches on the result. Here is what each task type tests and why a physical robot is essential for each.

51.1 T1 — Information Retrieval (RAG-Dependent)

Example utterances:

- "What time does the lab open?"
- "Tell me about Professor X's research."
- "What is the WiFi password?"
- "When is the next open day?"

What is tested: factual accuracy + RAG faithfulness. The underlying knowledge is in `knowledge_base/*.txt|csv|pdf`. The LLM's job is to use the retrieved chunks faithfully, not invent.

Why a robot is essential: Pepper greets the visitor, uses the tablet to show extra context (a photo of the professor, the lab map), and gestures while explaining. A pure chatbot loses the greeting-warmth dimension that affects perceived accuracy.

Metrics: RAG faithfulness score, retrieval similarity, judge score for factual accuracy.

51.2 T2 — Navigation / Guidance (Robot Body Essential)

Example utterances:

- "Where is Room 305?"
- "Can you point me to the cafeteria?"
- "How do I get to the seminar room?"

What is tested: spatial-language grounding + gesture-speech synchronisation. The LLM must produce a direction; the robot must point with the matching arm.

Why a robot is essential: Pepper **physically points** in directions, shows a map on its tablet, can walk partway to guide the person. A chatbot cannot point.

Metrics: gesture appropriateness (observer-rated), navigation success (binary), Likert "the robot's gestures were appropriate".

51.3 T3 — Social Conversation (Tests Naturalness)

Example utterances:

- "How are you?"
- "Tell me something interesting."
- "What do you think about AI?"

What is tested: fluency, warmth, empathy, openness. There is no correct answer; the LLM is judged on style.

Why a robot is essential: physical presence creates social pressure and engagement that text chat does not. Eye contact, head tracking, animated gestures change perceived naturalness.

Metrics: Likert "the robot was natural to talk to", Godspeed Anthropomorphism and Likeability subscales.

51.4 T4 — Multilingual Interaction (Tests Language Routing)

Example utterances:

- "Où est la salle 305 ?" (French)
- "Wo ist der Kaffeeraum?" (German)
- "¿Cuándo es el próximo seminario?" (Spanish)

What is tested: automatic language detection + per-language model routing. The system must detect the language and pick the model with the best coverage for that language.

Why a robot is essential: the robot's physical embodiment makes multilingual experience more immersive and testable than a screen chatbot. Foreign visitors meeting a robot is a higher-stakes interaction than typing into a translation box.

Metrics: language-detection accuracy, response-in-correct-language rate (observer-rated), Likert "the robot understood me".

Chapter 52 — The Five Experimental Conditions

Each participant experiences **three of the five** conditions in a counterbalanced Latin-square order. Across 20 participants this yields $60 \text{ condition-runs} \times 4 \text{ tasks} = \mathbf{240 \text{ interaction data points}}$.

Condition	Description	What It Tests
A — Fixed Cloud LLM	GPT-4o-mini for all tasks	Baseline cloud performance
B — Fixed Local LLM	Llama 3:8b via Ollama for all tasks	Baseline local / free performance
C — Smart-Routed	OmniLLM selects the best model per task type	Smart routing effectiveness
D — Consensus / Council	3 models answer; judge synthesises	Consensus vs. single model
E — RAG-Off Control	Same as A but without RAG	Isolates RAG contribution

52.1 Why These Five and Not Others

The design is the smallest set that probes all four dimensions:

- **A vs B** — cloud vs local.
- **A vs C** — fixed vs smart-routed.
- **A vs D** — single vs ensemble.
- **A vs E** — with-RAG vs without-RAG.

Adding more conditions (e.g., "smart-routed but local-only", "consensus without synthesis") gives more data but quadruples recruitment cost. The five-condition design has enough power for repeated-measures ANOVA on the main effects.

52.2 Counterbalancing — A Worked Latin Square

For 5 conditions and 5 participants per cell, a balanced Latin square distributes order:

Participant	Block 1	Block 2	Block 3
P001	A	C	E
P002	B	A	D
P003	C	E	A
P004	D	B	C
P005	E	D	B

P006	A	E	C
P007	B	D	A
P008	C	A	E
P009	D	C	B
P010	E	B	D

You can also use a Williams design or simple randomisation; with $N = 15$ the exact assignment matters less than ensuring **no condition is always first** and **no condition is always last**.

Chapter 53 — Automatic and Human-Rated Metrics

53.1 Automatic Metrics (logged by `ExperimentLogger`)

Every interaction creates an `InteractionRecord` with:

Field	Source
<code>model_id</code>	The model that ultimately answered
<code>latency_ms</code>	Wall clock from end-of-speech to start-of-speech
<code>input_tokens</code> / <code>output_tokens</code>	LiteLLM usage
<code>cost_usd</code>	Computed from <code>models.yaml</code> pricing
<code>rag_retrieval_scores</code>	ChromaDB cosine similarity (top-k)
<code>rag_faithfulness</code>	LLM-as-Judge score (0–1)
<code>hallucination_detected</code>	Heuristic flag (word-overlap < 20 %)
<code>judge_score</code>	Referenceless judge quality (0–1)
<code>detected_language</code>	LanguageDetector result
<code>task_type</code>	HRITaskClassifier output (T1–T4)
<code>gesture_used</code>	What Pepper actually did
<code>task_success</code>	Observer-coded binary

All twelve fields land in CSV/JSON exports — directly loadable into R, SPSS, JASP or pandas.

53.2 Human-Rated Metrics (After Each Condition)

The `InteractionQuestionnaire` dataclass collects five 1–7 Likert items:

1. The robot's answers were **accurate**.
2. The robot was **natural** to talk to.
3. I **trust** the information the robot gave me.
4. The robot's **gestures** were appropriate.
5. The robot responded **quickly enough**.

Optional: full **Godspeed** five-subscale (1–5 each): Anthropomorphism, Animacy, Likeability, Perceived Intelligence, Perceived Safety.

53.3 Pairwise Preference (At End of Session)

"Which version of Pepper did you prefer overall? Why?"

This produces a `PairwisePreference` record that updates the **Embodied LLM Leaderboard** via `EloScorer.record_match(..., category="embodied_hri")`. After all participants, the leaderboard ranks the five conditions and (indirectly) the models that powered each.

53.4 Observer Ratings (Live During Each Interaction)

The `observerRating` dataclass captures:

- gesture-speech synchronisation quality (1–5)
- breakdown count (number of "wait, what?" moments)
- task completion (binary)
- any qualitative notes

Chapter 54 — One-Month Execution Plan

A working schedule used by labs that have completed similar studies. Calibrate dates to your own timeline; the **dependencies between days** are what matters, not the absolute dates.

54.1 Week 1 — Infrastructure Setup (Days 1–7)

Day 1–2 — Pepper ↔ AI Server Bridge. Set up the Flask server (`python -m omnillm.server.app`). Verify `/health` returns OK. On Pepper (Python 2.7), write the NAOqi client. Test with a hard-coded response first, then with a single LLM.

Day 3–4 — OmniLLM + LiteLLM Integration. Configure `config/models.yaml` with all target models (GPT-4o-mini, Claude Haiku, Gemini Flash, DeepSeek + 3 Ollama models). Test the smart router and LLM-as-Judge.

Day 5–6 — RAG Pipeline. Populate `knowledge_base/{lab_info,faq,visitor_profiles,...}` with **real lab data**. Index. Test retrieval with `omnillm ask --rag "What is the WiFi password?"`.

Day 7 — LangGraph + End-to-End Integration. Build the agent graph (transcribe → classify → route → RAG/LLM → action plan → respond). End-to-end smoke test: speak to Pepper, get a RAG-augmented, gesture-synchronised response.

54.2 Week 2 — Refinement and Experiment Prep (Days 8–14)

Day 8–9 — Refine Robot Behaviours. Map response types to gestures (greeting → wave, information → open palms, pointing → arm raise toward direction). Set up tablet displays. Tune Whisper for your acoustic environment.

Day 10–11 — Experiment Protocol Design. Write the experimenter script. Create task prompts for T1–T4. Design the questionnaire (Google Forms or paper). Create consent forms. Set up automatic logging.

Day 12–13 — Pilot Testing. Run 3–4 pilot sessions with lab colleagues. Identify and fix latency, STT errors, gesture timing. Refine task prompts.

Day 14 — Final Preparations. Schedule participants. Prepare the physical lab space. Create a backup plan: if cloud APIs fail, fall back to Ollama. Finalise data-collection scripts.

54.3 Week 3 — Run Experiments (Days 15–21)

Day 15–21 — Participant Sessions. 3–4 participants per day. Each session ≈ 20–30 minutes (brief intro, 3 conditions × 4 tasks, questionnaire, debrief). Target: **15–20 participants minimum**. Experimenter takes observer notes. **Back up data daily.**

54.4 Week 4 — Analysis and Writing (Days 22–30)

Day 22–24 — Data Analysis. Aggregate the automatic metrics. Run repeated-measures ANOVA or Friedman test on the Likert means. Compute ELO rankings from pairwise preferences. Build the Embodied LLM Leaderboard. Cross-correlate with Chatbot Arena rankings.

Day 25–28 — Write. Introduction + Related Work (use the literature in Chapter 48). System Architecture (use the diagrams in Parts II and Appendix G). Experimental Design (use Chapters 51–53). Results.

Discussion. Conclusions + Future Work (Chapter 60).

Day 29–30 — Polish and Submit. Create figures. Proofread. Prepare the code repository for open-source release. Record a demo video.

Chapter 55 — Recommended Tools and the OpenClaw Question

55.1 The Recommended Stack

Layer	Tool	Why
Robot	Pepper (NAOqi 2.5, Python 2.7 client)	The platform under study
Orchestration	OmniLLM	Your central platform
LLM Gateway	LiteLLM	Unified API across 100+ providers
Agent Graph	LangGraph	Stateful conversation with conditional routing
RAG	LangChain + ChromaDB	Document loaders + vector store
Local LLMs	Ollama	Llama 3:8b, Qwen 2.5:7b, Mistral:7b, DeepSeek-R1 14b
Cloud APIs	OpenAI, Anthropic, Google, DeepSeek	Cloud baselines
Speech-to-Text	OpenAI Whisper (local or API)	Cross-lingual, robust
Bridge	Flask	Python 2.7 ↔ 3.x HTTP
Surveys	Google Forms or paper	Standard HRI questionnaire delivery

55.2 The OpenClaw Question

If you have access to **OpenClaw** (an agent-orchestration framework sometimes mentioned alongside LangGraph):

- Use it as an alternative to LangGraph for the multi-step pipeline: perceive → classify → retrieve → generate → act.
- It can manage tool-calling: the agent calls RAG retrieval as a tool, calls the OmniLLM router as a tool, calls the robot action API as a tool.
- If OpenClaw supports multi-agent patterns, you can split the work across agents: a "Retrieval Agent" (handles RAG), a "Router Agent" (selects LLM), a "Robot Agent" (generates action plans), a "Judge Agent" (evaluates quality).
- If OpenClaw has its own LLM routing capabilities, you can compare OpenClaw routing vs. OmniLLM routing as **an additional experimental condition**.

OmniLLM does not depend on OpenClaw; LangGraph already covers the same ground. Use whichever you are more comfortable with.

55.3 Why the Robot Is Essential (Direct Answers)

The user always asks: *"Could you do this without the robot?"*

You cannot because:

1. **The core research question is about embodied evaluation** — does physical presence change LLM quality perception? Removing the robot removes the question.
2. **Gesture-speech coordination timing** is a robot-specific variable.
3. **Navigation and pointing tasks** (T2) require a physical body.
4. **The tablet displays visual information** (maps, schedules) synchronised with speech.
5. **Physical-presence effects** (documented extensively in HRI literature — Bartneck, Wainer, Li, Bainbridge) are the confound that makes this study novel compared to any chatbot study.

You can, however, build and debug **most of OmniLLM** without a robot. Chapter 58 covers the Choregraphe-only path.

55.4 Risks and Mitigations

Risk	Mitigation
Pepper's speech recognition is unreliable	Use external microphone + Whisper API
Python 2.7 ↔ 3.x bridge latency	Optimise HTTP calls; co-locate the Flask server on a fast LAN
1-month timeline is tight	Use a minimal viable task set; recruit from lab visitors
Cloud API costs	Use mini / haiku / flash tier models + Ollama local for most testing
Small sample size	Use within-subjects design to increase statistical power
Hardware failure (battery, motors)	Have a backup Pepper or rescheduling plan; keep charger plugged
Participant no-shows	Over-recruit by 25 %; have a confirmation email 24 h before

55.5 Expected Contributions

If the experiment succeeds:

1. **First systematic multi-LLM comparison in embodied social robot HRI.**
2. **First evidence on whether embodied LLM rankings differ from text-only rankings.**

3. **First application of smart routing in social robotics.**
4. **First RAG-augmented social robot with multi-model backend comparison.**
5. **Open-source architecture** (OmniLLM + Pepper bridge) for reproducibility.
6. **Novel "Embodied LLM Arena" evaluation methodology** — a direct extension of LMSYS Chatbot Arena to physical space.

Suitable venues: **HRI Conference**, **ICRA**, **IEEE RO-MAN**, **Frontiers in Robotics and AI**, **MDPI Robotics**.

PART XI — EXTENDED BEGINNER MATERIAL

This part is for the **beginner / intermediate** reader. It assumes you have read Parts I-III but want **deeper** explanations. Four chapters expand on areas that beginners typically struggle with:

- **Chapter 56** walks one user question through every Python line.
- **Chapter 57** is a much larger glossary of the AI stack with the "why we chose this" reasoning.
- **Chapter 58** is a Choregraphe-only walkthrough — you can do this without a real Pepper.
- **Chapter 59** is a complete pretend experimental session — what commands to type, what the JSON looks like, what to enter into the questionnaire.
- **Chapter 60** is the future-work / improvement plan.

Chapter 56 — "Reading the Code": A Single Question's Journey, Line by Line

AT A GLANCE. Take one user input — "Where is Room 305?" — and trace it through every Python function that touches it. If you can read this chapter top to bottom and recognise every step, you understand OmniLLM.

56.1 The Scenario

You are running:

```
# Terminal 1
python -m omnillm.server.app

# Terminal 2 (drives the AI server with curl)
curl -X POST http://localhost:5000/interact \
-H "Content-Type: application/json" \
-d '{"text": "Where is Room 305?",
    "participant_id": "DEMO",
    "session_id": "demo-1",
    "condition": "C",
    "rag_enabled": true}'
```

What happens between the curl and the JSON response?

56.2 Stage 1 — Flask Receives the Request

File: `omnillm/server/app.py`, function `interact()`

Flask matches `POST /interact` to a Python function. The body is parsed from JSON into a dict:

```
@app.route("/interact", methods=["POST"])
def interact():
    data = request.get_json(force=True) or {}
    utterance = data.get("text", "") # "Where is Room 305?"
    audio_b64 = data.get("audio", "") # "" (text-only)
    participant_id = data.get("participant_id", "unknown")
    session_id = data.get("session_id", "")
    condition = data.get("condition", "A")
    rag_enabled = bool(data.get("rag_enabled", True) and rag is not None)
```

What you see in Python you may not recognise:

- `@app.route(...)` is a **decorator**. It wraps the function below it with Flask's URL routing logic. The function is still callable directly, but Flask also knows "if a `POST /interact` comes in, call `interact`".
- `data.get("text", "")` returns `data["text"]` if it exists, otherwise `""`. This is **safer than** `data["text"]` because the latter raises `KeyError` if the field is missing.

56.3 Stage 2 — Hand Off to the LangGraph Pipeline

File: `omnillm/server/app.py`, function `interact()` (continued)

```
audio_bytes = base64.b64decode(audio_b64) if audio_b64 else b""
```

```
graph = _get_graph()           # lazily compile LangGraph the first time
if graph is not None:
    state = {
        "utterance": utterance,
        "audio_bytes": audio_bytes,
        "participant_id": participant_id,
        "session_id": session_id,
        "condition": condition,
        "rag_enabled": rag_enabled,
        "model_id": _default_model,
    }
    result = _run_async(graph.ainvoke(state))
    return jsonify(result.get("robot_action",
                              {"speech": result.get("response_text", "")}))
```

What's happening:

- `_get_graph()` returns a compiled LangGraph DAG. The first call takes about 10 seconds (imports + compile); subsequent calls are instant.
- `graph.ainvoke(state)` runs the **entire graph asynchronously** starting from the `state` dict. Each node reads keys it cares about and writes back updates.
- `_run_async(coro)` is a tiny helper that creates a fresh event loop per request — needed because Flask is synchronous but LangGraph is asynchronous.

56.4 Stage 3 — Inside the Graph (Node by Node)

File: `omnilm/hri/agent_graph.py`, function `build_hri_graph()`

Compilation builds this graph:

```
START
v
transcribe_audio
v
detect_language
v
classify_task
v (conditional)
+> rag          (T1, RAG on)
+> nav_rag      (T2, RAG on)
+> direct_llm   (T3 or condition E)
+> multilingual_llm (T4)
v
smart_router
v
generate_action_plan
v
log_interaction
v
END
```

56.4.1 Node 1 — `transcribe_audio`

```
async def transcribe_audio(state: dict) -> dict:
    audio = state.get("audio_bytes", b"")
    if not audio:
        # text-only mode — utterance already set
        return {"_start_time": time.monotonic()}
    # ...else call Whisper STT (skipped here)
```

Because we passed `text="Where is Room 305?"` (no audio), `audio` is empty, so this node just records the start time and passes through.

56.4.2 Node 2 — detect_language

```
async def detect_language(state: dict) -> dict:
    utterance = state.get("utterance", "")
    detector = LanguageDetector()
    result = detector.detect(utterance)
    return {"detected_language": result.language_code}
```

The detector inspects the script and word frequencies and returns "en". Almost any English text is identified in under 10 ms by Tier 1 (Unicode script check) — no LLM call required.

56.4.3 Node 3 — classify_task

```
async def classify_task(state: dict) -> dict:
    utterance = state.get("utterance", "")
    classifier = HRTaskClassifier()
    result = classifier.classify(utterance, detected_language="en")
    return {"task_type": result.task_type.value,
            "task_confidence": result.confidence}
```

The classifier hits the regex `\broom\s+\d+\b` ("Room 305") and the keyword `where`. Score for **navigation** wins → returns `{"task_type": "navigation", "task_confidence": 0.92}`.

56.4.4 Conditional Edge — Pick the Branch

```
def _route_by_task_type(state: dict) -> str:
    task_type = state.get("task_type", "info_retrieval")
    condition = state.get("condition", "A")
    rag_enabled = state.get("rag_enabled", True)
    if condition == "E":
        return "direct_llm"
    routing = {
        "info_retrieval": "rag" if rag_enabled else "direct_llm",
        "navigation": "nav_rag" if rag_enabled else "direct_llm",
        "social_conversation": "direct_llm",
        "multilingual": "multilingual_llm",
    }
    return routing.get(task_type, "direct_llm")
```

`task_type="navigation"` and `rag_enabled=True` → return `"nav_rag"`. LangGraph dispatches to the `nav_rag` node.

56.4.5 Node 4 — nav_rag (RAG + Gesture Planner)

```
async def run_nav_rag(state: dict) -> dict:
    utterance = state.get("utterance", "")
    rag_resp = await rag.query(utterance, model_id=state.get("model_id"))
    planner = GesturePlanner()
    gesture, led = planner.plan("navigation", rag_resp.answer)
    return {
        "rag_context": "\n\n".join(f"[{c.source}]: {c.text}"
                                    for c in rag_resp.retrieved_chunks),
        "rag_chunks": rag_resp.retrieved_chunks,
        "rag_faithfulness": rag_resp.faithfulness_score,
        "response_text": rag_resp.answer,
        "model_id": rag_resp.model_id,
        "latency_ms": rag_resp.latency_ms,
        "gesture": gesture,
        "led_color": led,
    }
```

What `rag.query()` does internally:

1. Convert "Where is Room 305?" into an embedding vector.
2. Search ChromaDB for the 4 most similar chunks.
3. Build an augmented prompt: `system + "Use only this context: ..." + user`.
4. Call `gateway.query(model_id, messages)` — this returns the answer.
5. (Optional) Score the answer's faithfulness with a judge LLM.

The result is something like:

```
"Room 305 is on the third floor of Building C. Take the elevator
on your left, then turn left at the corridor."
```

Then `planner.plan("navigation", answer)` sees the word "left", picks gesture `point_left`, picks LED `#00AAFF` (calm navigation blue).

56.4.6 Node 5 — `smart_router`

For `condition="C"`, the router fires:

```
from omnillm.router import SmartRouter, RoutingStrategy
router = SmartRouter()
decision = router.route_for_hri_task(
    task_type="navigation",
    strategy=RoutingStrategy.TASK_TYPE,
)
target = decision.model_id      # e.g. "gemini-flash" per hri_task_routing
resp = await gateway.query(target, messages)
return {"response_text": resp.content, "model_id": resp.model_id, ...}
```

The router reads `config/models.yaml`'s `hri_task_routing.navigation` and finds `gemini-flash` is the optimal model for navigation. It re-asks the question with that model. The answer overrides what `nav_rag` produced.

(For condition A or B, the `smart_router` node is a no-op — the previous answer wins.)

56.4.7 Node 6 — `generate_action_plan`

```
async def generate_action_plan(state: dict) -> dict:
    return {
        "robot_action": {
            "speech": state["response_text"],
            "gesture": state["gesture"],
            "emotion_led": state["led_color"],
            "metadata": {
                "task_type": state["task_type"],
                "model_id": state["model_id"],
                "rag_enabled": state["rag_enabled"]},
        }
    }
```

This is just packaging — the heavy work is done.

56.4.8 Node 7 — log_interaction

```
async def log_interaction(state: dict) -> dict:
    total_latency = (time.monotonic() - state["_start_time"]) * 1000
    logger.log_interaction(
        session_id=state.get("session_id", ""),
        participant_id=state.get("participant_id", ""),
        condition=state.get("condition", "A"),
        task_type=state.get("task_type"),
        utterance=state.get("utterance", ""),
        response=state.get("response_text", ""),
        model_id=state.get("model_id", ""),
        latency_ms=total_latency,
        ...
    )
    return {"latency_ms": total_latency}
```

The logger appends an `InteractionRecord` to the in-memory list (written to disk at end of session via `logger.save()`).

56.5 Stage 4 — Flask Returns the JSON

Back in `app.py`:

```
return jsonify(result.get("robot_action", ...))
```

The dict is serialised to JSON and returned with HTTP 200:

```
{
  "speech": "Room 305 is on the third floor of Building C. Take the elevator on your left, then turn left at the co",
  "gesture": "point_left",
  "emotion_led": "#00AAFF",
  "metadata": {
    "task_type": "navigation",
    "model_id": "gemini-flash",
    "rag_enabled": true
  }
}
```

56.6 Stage 5 — What Pepper Does With It

If a real Pepper was at the other end, `naoqi_client.py` would:

```
def _execute_action(self, action):
    speech = action.get("speech", "")
    gesture = action.get("gesture")
    led_color = action.get("emotion_led")
    if led_color: self._set_leds(led_color)
    if gesture and gesture in GESTURE_TO_BEHAVIOR:
        self._run_behavior_async(GESTURE_TO_BEHAVIOR[gesture])
    if speech: self._speak(speech)
```

- Eyes fade to `#00AAFF` over 300 ms.
- `ALBehaviorManager.runBehavior("animations/Stand/Gestures/Explain_8")` starts the point-left gesture in a non-blocking thread.
- `ALAnimatedSpeech.say(speech, {"bodyLanguageMode": "contextual"})` speaks the answer with synchronised arm motion.

56.7 The Whole Picture

That is one question, end-to-end. Sub-second latency (≈ 1.4 s typical), all of it logged, every model swap configurable, every step testable in isolation.

If you can find the code for each step in the repo without re-reading this chapter, you understand OmniLLM.

Chapter 57 — AI-Stack and Frameworks — Expanded Glossary with Rationale

AT A GLANCE. Every important tool OmniLLM depends on, with a "what it is", "why we picked it", "what we'd use otherwise", and "where to read more".

57.1 LiteLLM

What it is. A Python library that exposes 100+ LLM providers behind one identical function call. You pass `model="openai/gpt-4o"` or `model="anthropic/claude-haiku"` or `model="ollama/llama3:8b"` — the call is the same.

Why we picked it. Without it, OmniLLM would need separate adapter code for each provider, each with its own pagination, error-handling, and token-counting quirks. LiteLLM normalises this in one library that's actively maintained.

Alternatives. Direct provider SDKs (`openai`, `anthropic`, `google-generativeai`) — more control, more code. `aisuite`, `openrouter` — similar idea, smaller communities.

Where used. Imported by `omnillm/gateway.py`.

57.2 LangGraph

What it is. A library on top of LangChain for building **stateful agent graphs** — directed graphs whose nodes are async functions sharing a single state dict, with conditional edges between them.

Why we picked it. OmniLLM's HRI pipeline is genuinely a graph (transcribe → detect → classify → branch → answer → plan → log) with conditional routing. Writing this as one 200-line async function with `if/elif` branches would work but be untestable, unreadable, and impossible to visualise. LangGraph gives us named nodes, typed state, and `add_conditional_edges`.

Alternatives. Plain async functions (more verbose, no diagram). LangChain's deprecated `AgentExecutor` (limited routing). Custom DAG libraries — most don't handle async well.

Where used. `omnillm/hri/agent_graph.py`.

57.3 ChromaDB

What it is. A pure-Python embedded vector database. Stores embeddings + their source text + metadata. Supports cosine similarity search.

Why we picked it. Embedded (no separate process), persistent across restarts, zero configuration, pip-installable. For a single researcher with a < 10 k-document knowledge base, this is plenty.

Alternatives. **Qdrant**, **Weaviate**, **Pinecone** — heavier, require running a server. **FAISS** — fast but lower-level, no metadata. For OmniLLM's scale, ChromaDB is the sweet spot.

Where used. `omnillm/rag/pipeline.py`.

57.4 LangChain

What it is. A broader framework for LLM apps. OmniLLM uses *only* its document loaders (PyPDFLoader, CSVLoader, TextLoader) — not the larger Chain abstractions, which are opinionated and easy to outgrow.

Why we picked it. The document loaders save dozens of lines per file format. Tested by a large community.

Alternatives. Write your own loaders. For PDFs, pypdf directly. For CSVs, pandas. We could do this — but LangChain's loaders normalise the output (each chunk has .page_content, .metadata) which makes the rest of the pipeline cleaner.

Where used. omnillm/rag/pipeline.py.

57.5 Sentence-Transformers

What it is. A Python library that turns text into embedding vectors using small (~100 MB), fast, locally-runnable models.

Why we picked it. ChromaDB's default embedding is all-MiniLM-L6-v2 from sentence-transformers — fast (~5 ms per short text on CPU), small, and free. No need to pay OpenAI for embeddings during indexing.

Alternatives. OpenAI's text-embedding-3-small API — higher-quality embeddings, \$0.02 per million tokens, requires internet. For a small lab KB, the quality difference doesn't justify the cost.

Where used. Indirectly via ChromaDB.

57.6 OpenAI Whisper

What it is. A speech-to-text model trained by OpenAI on 680 k hours of multilingual audio. Available as a downloadable model (pip install openai-whisper) or via the OpenAI API.

Why we picked it. Best-in-class accuracy at the **multilingual** level, robust to noise, free if you run it locally. Pepper's onboard ALSpeechRecognition is widely regarded as unusable for open-domain speech.

Alternatives. Vosk (lighter, offline), Google Speech (cloud), Faster-Whisper (a faster C++ port — recommended for production but adds setup complexity). For a research prototype, plain Whisper is fine.

Where used. omnillm/robotics/whisper_stt.py.

57.7 Asyncio

What it is. Python's standard library for **non-blocking I/O**. async def declares a coroutine; await pauses it until an I/O operation completes; asyncio.gather(...) runs many coroutines concurrently.

Why we picked it. LLM calls are I/O-bound — most of the time is spent waiting for the network. Asyncio lets us query 5 models in the wall-clock time of the slowest one, not the sum.

Alternatives. Threading (more complex, GIL contention), multiprocessing (overkill, slow startup). For pure-I/O work, asyncio wins.

Where used. Throughout — gateway.py, consensus.py, evaluator.py, agent_graph.py, server/app.py.

57.8 Flask

What it is. A minimalist Python web framework. Serves HTTP endpoints with decorators (`@app.route`).

Why we picked it. Tiny, well-documented, easy to test, single worker is the right concurrency model for one robot.

Alternatives. **FastAPI** (async-native, faster, more modern) — genuinely better for high-throughput services. We didn't pick it because Flask's simplicity matched the single-robot single-worker use case, and the synchronous-to-async bridge (`_run_async`) is easy to write. If you ever serve many robots simultaneously, switch to FastAPI + uvicorn.

Where used. `omnillm/server/app.py`.

57.9 Click

What it is. A Python library for building command-line interfaces with decorators (`@click.command`, `@click.option`).

Why we picked it. Better than `argparse` (less boilerplate), supports nested command groups (`omnillm ask`, `omnillm evaluate`), auto-generates `--help` text.

Alternatives. **Typer** (built on Click, uses type hints — slightly nicer for greenfield projects), **argparse** (`stdlib`, more verbose).

Where used. `omnillm/cli.py`.

57.10 Rich

What it is. A Python library for beautiful coloured terminal output — tables, panels, progress bars, syntax-highlighted code.

Why we picked it. OmniLLM's CLI is a research dashboard. A coloured table of model latencies is dramatically more readable than `print(dict)`.

Alternatives. **tabulate** (tables only), plain `print()` (grey, unreadable for long output).

Where used. `omnillm/cli.py`.

57.11 Ollama

What it is. A free local LLM runtime. `ollama serve` starts a daemon on `localhost:11434`; `ollama pull llama3:8b` downloads a model. Exposes an OpenAI-compatible HTTP API.

Why we picked it. Free, private, fast on a laptop with a GPU. The "Condition B" baseline of the experiment depends on it.

Alternatives. **vLLM** (faster but heavier), **LM Studio** (GUI), **GGUF + llama.cpp** (more control, more setup). Ollama is the easiest for a research lab.

Where used. Via LiteLLM in `omnillm/gateway.py`. Models registered in `config/models.yaml` with `provider: ollama`.

57.12 ReportLab + xhtml2pdf

What it is. Two pure-Python libraries that together produce PDFs from HTML. xhtml2pdf orchestrates; ReportLab does the actual rendering. We register DejaVu Sans + DejaVu Sans Mono with ReportLab so Unicode glyphs render.

Why we picked it. No native dependencies (libgobject, Pango, Cairo) — installable on Windows with pip alone. WeasyPrint produces better output but requires GTK on Windows which the user couldn't install.

Alternatives. **WeasyPrint** (best quality, needs GTK), **mPDF** (PHP), **wkhtmltopdf** (Webkit, large binary), **LaTeX** (best quality, huge install). xhtml2pdf is the lowest-friction choice.

Where used. `book/build_pdf.py`.

57.13 pytest + pytest-asyncio + pytest-mock

What it is. The Python testing trinity. pytest is the framework; pytest-asyncio runs async test functions; pytest-mock provides the `mock` fixture.

Why we picked it. Standard in the Python ecosystem. The 278 tests in `tests/` run in ~5 s without a single real API call — because pytest-mock fakes out the LLM calls.

Where used. `tests/`.

Chapter 58 — Choregraphe-Only Walkthrough (No Robot Needed)

AT A GLANCE. You can develop and demo 90 % of OmniLLM without a physical robot. Choregraphe's virtual robot is enough. This chapter is the complete recipe.

58.1 Why This Path Exists

A real Pepper costs \$20–30 k second-hand (and rising, since production stopped). For development you only need to verify:

- The AI server returns sensible JSON.
- The JSON, when replayed against NAOqi, produces the expected speech
- gesture + LED.

Choregraphe's virtual robot is sufficient for both.

58.2 Install Choregraphe 2.5

Choregraphe is shipped by SoftBank / Aldebaran. The download is on the SoftBank developer portal (free account required); the version that matches Pepper is **2.5.5.5** or **2.5.10/11**.

- Install at the default location: `C:\Program Files (x86)\Aldebaran\Choregraphe Suite 2.5\`.
- Launch — you should see the four-panel layout (Box Library, Flow Diagram, 3D Robot, Log Viewer).

58.3 Spin Up the Virtual Robot

1. Menu → **Connection** → **Connect to virtual robot**.
2. The 3D view fills with a virtual Pepper.
3. Right-click the robot in the 3D view → **Wake Up**. Stiffness turns on; Pepper holds upright.
4. The status bar shows `Connected to localhost:9559 (virtual)`.

58.4 Start OmniLLM

In a separate terminal:

```
cd C:\Users\akshi\OneDrive\Desktop\OmniLLM
.venv\Scripts\activate
python -m omnillm.server.app --model llama3-8b-local
# (or --model openai-gpt4o-mini if you have an API key)
```

In another terminal, query the server:

```
curl -X POST http://localhost:5000/interact ^
-H "Content-Type: application/json" ^
-d "{\"text\":\"Where is Room 305?\",\"participant_id\":\"DEMO\",\"condition\":\"B\"}"
```

Response:

```
{
  "speech": "Room 305 is on your left on the third floor.",
  "gesture": "point_left",
  "emotion_led": "#00AAFF",
  "metadata": {"task_type": "navigation", "model_id": "llama3-8b-local"}
}
```

58.5 Replay the Action on the Virtual Robot

In Choregraphe, press **Alt+5** to open the Script Editor. Paste:

```
# (replace with the values from your curl response)
speech = "Room 305 is on your left on the third floor."
gesture = "animations/Stand/Gestures/Explain_8" # point_left
led_hex = (0.0, 0.67, 1.0) # #00AAFF

tts = ALProxy("ALAnimatedSpeech", "localhost", 9559)
bm = ALProxy("ALBehaviorManager", "localhost", 9559)
leds = ALProxy("ALLeds", "localhost", 9559)

leds.fadeRGB("FaceLeds", led_hex[0], led_hex[1], led_hex[2], 0.3)
bm.post.runBehavior(gesture) # non-blocking
tts.say(speech, {"bodyLanguageMode": "contextual"})
```

You should see:

- The virtual Pepper's eyes turn blue (#00AAFF).
- Its left arm raises (the Explain_8 gesture).
- Your PC speakers play the speech.

This proves OmniLLM ↔ NAOqi end-to-end **without a real robot**.

58.6 Fully Automating Replay

A helper script that closes the loop — Choregraphe-side Python 2.7 client that polls the AI server like the real `naoqi_client.py`:

```
# choregraphe_replay.py -- run in Choregraphe Script Editor
import json, urllib2, time
from naoqi import ALProxy

SERVER = "http://localhost:5000/interact"
GESTURE_MAP = {
  "point_left": "animations/Stand/Gestures/Explain_8",
  "point_right": "animations/Stand/Gestures/Explain_7",
  "wave": "animations/Stand/Gestures/Hey_1",
  "nod": "animations/Stand/Emotions/Positive/Enthusiastic_1",
  "think": "animations/Stand/Emotions/Neutral/Thinking_1",
}

tts = ALProxy("ALAnimatedSpeech", "localhost", 9559)
bm = ALProxy("ALBehaviorManager", "localhost", 9559)
leds = ALProxy("ALLeds", "localhost", 9559)

def ask(text):
    payload = json.dumps({"text": text, "participant_id": "DEMO",
                          "session_id": "demo", "condition": "C",
                          "rag_enabled": True})
    req = urllib2.Request(SERVER, payload,
                          {"Content-Type": "application/json"})
```

```

    return json.loads(urllib2.urlopen(req).read())

def execute(action):
    led = action.get("emotion_led", "#44AAFF").lstrip("#")
    r, g, b = (int(led[i:i+2], 16) / 255.0 for i in (0,2,4))
    leds.faderRGB("FaceLeds", r, g, b, 0.3)
    gname = action.get("gesture")
    if gname and gname in GESTURE_MAP:
        bm.post.runBehavior(GESTURE_MAP[gname])
    if action.get("speech"):
        tts.say(action["speech"], {"bodyLanguageMode": "contextual"})

# Demo loop -- four task types
for q in [
    "What time does the lab open?",
    "Where is Room 305?",
    "How are you today, Pepper?",
    "Bonjour Pepper, comment allez-vous?",
]:
    print(q)
    execute(ask(q))
    time.sleep(2)

```

Save this in Choregraphe's Script Editor and run it. The virtual Pepper will speak, gesture and change LED colour for each of the four questions, end-to-end through OmniLLM. **This is the zero-hardware demo loop.**

58.7 What You Cannot Test Without a Real Robot

Feature	Virtual works?	Why / Why not
ALAnimatedSpeech (speech + body language)	yes (PC speakers)	rendered
ALMotion (joints)	yes	rendered in 3D
ALLeds	yes	LED colour shows in 3D
ALBehaviorManager (pre-installed gestures)	yes	runs the animation
ALAudioDevice (microphone)	no	no virtual mic
ALAudioRecorder	no	no audio stream
ALFaceDetection	no	no camera
ALTabletService	partial	tablet view exists but limited

The bottom line: anything *output* works on the virtual robot; anything *input from the world* (audio, vision, touch) does not. So you cannot run a full STT-driven session on the virtual robot. Use the `--text` curl path instead.

Chapter 59 — A Complete First Experimental Session — Worked Example

AT A GLANCE. A pretend participant P001 walks into your lab. We run them through all three conditions of one slot of the Latin square. Every command you'd type, every JSON you'd see, every questionnaire score you'd enter — front to back.

59.1 Before the Participant Arrives

```
# 0. Pre-flight check
cd C:\Users\akshi\OneDrive\Desktop\OmniLLM
.venv\Scripts\activate

omnillm models                # all 19 models load
pytest tests/ -v 2>&1 | tail -3 # 278 tests pass

# 1. Make sure Pepper is awake (or virtual Choregraphe is connected)
C:\Python27\python.exe -c "from naoqi import ALProxy; ^
ALProxy('ALMotion', '192.168.1.100', 9559).wakeUp()"

# 2. Start the AI server, this terminal
python -m omnillm.server.app --host 0.0.0.0 --port 5000

# expected log lines:
# INFO:omnillm.rag.pipeline:Indexed 12 chunks from lab_info.txt
# INFO:omnillm.rag.pipeline:Indexed 8 chunks from faq.txt
# INFO:omnillm.rag.pipeline:Knowledge base loaded -- 42 total chunks
# * Running on all addresses (0.0.0.0)
# * Running on http://127.0.0.1:5000
```

You also have a paper copy of:

- Consent form (two copies)
- The 5-item questionnaire (printed)
- A Latin-square sheet that says "P001: blocks A, C, E in that order"

59.2 Participant Arrives — Briefing (5 minutes)

"Thank you for joining. Today you'll talk with Pepper in three different modes. After each mode you'll fill in a short questionnaire. The whole session is about 30 minutes. You can stop at any time without giving a reason. There are no right or wrong answers — we're studying Pepper, not you. Please sign here..."

Hand over the consent form. Note that P001 speaks English and French.

59.3 Block 1 — Condition A (GPT-4o-mini, RAG ON)

In a second terminal, launch the NAOqi client (or, in our virtual demo, the `choregraphe_replay.py` from §58.6 — but pass the participant ID and condition):

```
# (real Pepper)
C:\Python27\python.exe omnillm\server\naoqi_client.py ^
--robot-ip 192.168.1.100 ^
```

```
--server-ip 192.168.1.50 ^
--participant P001 ^
--condition A
```

Pepper says: "Hello! I am Pepper, powered by OmniLLM. How can I help you today?"

You prompt the participant for each task in turn:

Task	Your verbal cue to participant
T1	"Ask Pepper a factual question about the lab — like the WiFi password, or what time it opens."
T2	"Ask Pepper to direct you somewhere — like Room 305 or the cafeteria."
T3	"Have a casual chat with Pepper — say hello, ask how it is."
T4	"You said you speak French — please ask any of those questions again in French."

P001 says: "What's the lab WiFi password?"

Server logs (truncated):

```
INFO:omnillm.server.app:[interact] participant=P001 condition=A task=info_retrieval
INFO:omnillm.rag.pipeline:Retrieved 4 chunks, top score 0.78
INFO:litellm:openai-gpt4o-mini in=320 out=42 latency=987ms cost=$0.0001
INFO:omnillm.utils.experiment_logger:Logged interaction len(records)=1
```

Pepper responds: "The visitor WiFi is 'UniGuest' with password 'Welcome2026!'."

After all four tasks for block A, hand over the **questionnaire**:

```
Condition (you write): A
Participant ID: P001
Q1 The robot's answers were accurate [ ] 1 [ ] 2 [ ] 3 [ ] 4 [ ] 5 [ ] 6 [ ] 7
Q2 The robot was natural to talk to [ ] 1 [ ] 2 [ ] 3 [ ] 4 [ ] 5 [ ] 6 [ ] 7
Q3 I trust the information the robot gave me [ ] 1 [ ] 2 [ ] 3 [ ] 4 [ ] 5 [ ] 6 [ ] 7
Q4 The robot's gestures were appropriate [ ] 1 [ ] 2 [ ] 3 [ ] 4 [ ] 5 [ ] 6 [ ] 7
Q5 The robot responded quickly enough [ ] 1 [ ] 2 [ ] 3 [ ] 4 [ ] 5 [ ] 6 [ ] 7
```

P001 ticks 6 / 5 / 6 / 5 / 7.

Save those into OmniLLM via the `/evaluate` endpoint:

```
curl -X POST http://localhost:5000/evaluate ^
-H "Content-Type: application/json" ^
-d "{\"session_id\":\"P001-A\",\"participant_id\":\"P001\",\"condition\":\"A\",\"scores\":{\"accuracy\":6,\"natur
```

Or do it programmatically at the end (Section 59.7).

59.4 Block 2 — Condition C (Smart-Routed)

Kill the NAOqi client (Ctrl+C). Restart with `--condition C`:

```
C:\Python27\python.exe omnillm\server\naoqi_client.py ^
--robot-ip 192.168.1.100 ^
--server-ip 192.168.1.50 ^
--participant P001 ^
--condition C
```

Same four tasks. This time the smart router picks:

Task	Picked model	Why (from <code>hri_task_routing</code> in <code>models.yaml</code>)
T1 Info Retrieval	openai-gpt4o-mini	accurate, cheap, RAG-friendly
T2 Navigation	gemini-flash	fast, strong spatial language
T3 Social Chat	claude-haiku	natural conversational tone
T4 Multilingual	gemini-flash	strong multilingual coverage

P001 ticks the questionnaire: 6 / 6 / 6 / 6 / 6.

59.5 Block 3 — Condition E (RAG-OFF Control)

```
C:\Python27\python.exe omnillm\server\naoqi_client.py ^
--robot-ip 192.168.1.100 ^
--server-ip 192.168.1.50 ^
--participant P001 ^
--condition E
```

Same four tasks again. **The interesting moment** — P001 asks "What's the WiFi password?" Without RAG, the LLM has never seen your lab's WiFi password. It says something like *"I'm sorry, I don't have that information."* That's the data point.

P001 ticks: 4 / 5 / 3 / 5 / 7. Notice the drop in **accuracy** and **trust** vs. block A — exactly H3's prediction.

59.6 Pairwise Preference

"Which version of Pepper did you prefer overall? Why?"

P001: "The second one was the most natural."

That maps to **Condition C wins over A and E**. Record it:

```
curl -X POST http://localhost:5000/evaluate ^
-H "Content-Type: application/json" ^
-d "{\"type\": \"pairwise\", \"session_id\": \"P001\", \"participant_id\": \"P001\", \"condition_a\": \"A\", \"condition_
```

(And similarly C > E.) These feed the ELO scorer's `embodied_hri` category.

59.7 End of Session — Save Everything

```
# Dump all the logged interactions for the day
curl http://localhost:5000/export > results/P001_full.json

# (Inside the AI server REPL, or as a separate script)
python -c "
from omnillm.scorer import EloScorer
elo = EloScorer()
elo.record_match('condition-C', 'condition-A', 'model_a', category='embodied_hri')
elo.record_match('condition-C', 'condition-E', 'model_a', category='embodied_hri')
elo.save('results/elo.json')
print(elo.get_category_leaderboard('embodied_hri'))
"
```

59.8 What the JSON Logs Look Like

A typical InteractionRecord row in results/P001_full.json:

```
{
  "timestamp": "2026-05-14T10:32:17.512Z",
  "session_id": "P001-A",
  "participant_id": "P001",
  "condition": "A",
  "task_type": "info_retrieval",
  "utterance": "What's the lab WiFi password?",
  "response": "The visitor WiFi is 'UniGuest' with password 'Welcome2026!'.",
  "model_id": "openai-gpt4o-mini",
  "latency_ms": 987.3,
  "input_tokens": 320,
  "output_tokens": 42,
  "cost_usd": 0.0001,
  "rag_enabled": true,
  "rag_faithfulness": 0.94,
  "rag_chunk_count": 4,
  "judge_score": 0.91,
  "task_success": true,
  "language": "en",
  "gesture_used": "nod"
}
```

After 20 participants × 3 conditions × 4 tasks, you have **240 such rows** plus 100 questionnaire rows plus ~60 pairwise preferences. That is the dataset.

59.9 Final Backup Step

```
# Copy results/ to a separate location at the end of every day
xcopy /E /Y results\ D:\thesis-backup\results-2026-05-14\
```

Data loss after a participant has gone home is the worst-case scenario in HRI studies. Don't be the person it happens to.

Chapter 60 — Future Scope and Improvement Opportunities

AT A GLANCE. Where OmniLLM goes next, in priority order. Some are weekend tasks; some are PhD theses in themselves.

60.1 Near-Term (Weekend Tasks)

Real audio capture in `naoqi_client.py`. The current `_record_audio()` is a 5-second placeholder. Replace with `ALAudioRecorder` (simpler, +file transfer) or `processRemote` callback (real-time, more code). Critical before actual user studies.

Counterbalancing generator. Add a `LatinSquareGenerator` to `omnilm/hri/experiment.py` so the experimenter doesn't manage a CSV by hand.

A `/admin` dashboard. A small Flask page showing live: model costs, recent interactions, current ELO, participant progress.

Realtime API support. Both OpenAI's Realtime API and Google Gemini Live can do speech-to-speech without an explicit STT step. For Pepper, this would collapse the audio path to a single low-latency loop, dropping ~400 ms of latency.

60.2 Mid-Term (Sprint Tasks)

Streaming TTS responses. Pepper's TTS could start speaking as soon as the first sentence is generated, rather than waiting for the full response. The Buddy bridge already does this; Pepper would need a parallel `stream_response` path. Cuts perceived latency by 30–60 %.

Multi-modal LLMs with vision. Pepper has two cameras + a depth sensor. The agent graph could consume "what does Pepper see" as an additional input. GPT-4o, Gemini and Claude all accept image inputs. This unlocks "describe the person in front of me" tasks.

Adversarial robustness eval. A separate evaluation pass that probes how each condition handles noisy speech, ambiguous questions, off-topic chatter, and minor adversarial inputs. Useful for a deployment readiness report.

Auto-counterbalancing + auto-randomisation. Combine with the admin dashboard for a one-click experimental session.

60.3 Long-Term (Research Programmes)

QiSDK migration. For NAOqi 2.9 (Android) Pepper units, rewrite the NAOqi client in Kotlin/Java using QiSDK. This eliminates the Python 2.7 problem entirely. Substantial rewrite (~3 weeks).

Larger-scale multi-site study. Replicate the methodology in multiple labs (hospitals, universities, retail) and pool the data. The cross-context generalisation is publishable in its own right.

Adaptive routing learning. Replace the static `hri_task_routing` YAML with a learned router: after every interaction, update the router's preference based on observed satisfaction. A multi-armed bandit (Thompson sampling on a per-task basis) could do this in ~100 lines.

Personalisation. A `participant_profile` in `ExperimentManager` that the smart router consults. If P012 historically preferred fast short answers, route their interactions to gemini-flash even on tasks where

another model wins on average.

Multi-robot orchestration. OmniLLM as the brain for a small *team* of robots (Pepper + NAO + Buddy in one room). The agent graph becomes a multi-agent system; each robot has its own `RobotBridge` and `RobotAction`.

Active learning of the knowledge base. When the RAG faithfulness score drops below a threshold, the system logs the question as "not covered" and prompts the human curator (you) to add a document. Knowledge base improves over time.

A general "embodied benchmark" leaderboard. Beyond a single lab's data, a public submission portal (like Chatbot Arena) where any HRI lab can upload their results and the global Embodied LLM Leaderboard updates.

60.4 What Could Make This A Bigger Story

If a single result emerges from your $N = 20$ study, three follow-up papers are obvious:

1. **"Why Smart Routing Wins."** Decompose the C-condition gains: when does smart routing help? On which task types? Under which noise conditions?
2. **"The Embodiment Gap."** Quantify the per-model gap between Chatbot Arena ELO and Embodied ELO. Identify which model capabilities translate to embodiment and which don't.
3. **"Faithfulness Under Pressure."** RAG faithfulness scores per-model under heavy retrieval (top-10 chunks) vs. light retrieval (top-3). Which models stay faithful as context grows?

60.5 When To Stop

You will be tempted to add features forever. The thesis writes itself when you can answer **three** questions with data:

- Does smart routing beat fixed (H2)?
- Does RAG help (H3)?
- Does embodiment change the ranking (H1)?

Everything else is decoration. Add it after submission.

Closing Note (Second Edition)

You have reached the end of the expanded book. By now you should be able to:

- **Explain** in plain English what OmniLLM is, why it exists, and what the robot adds — including to a sceptic asking "why bother with a robot?"
- **Read** any source file and know what each function does without guessing (Part III + Chapter 56).
- **Trace** a single user question through every Python function it touches (Chapter 56).
- **Connect** to a Pepper robot (or simulate one in Choregraphe) and drive it through the AI server (Parts IV, V, Chapter 58).
- **Re-build** the project from scratch in a sensible order (Part VI).
- **Run** a complete experimental session, collect the data, and analyse it (Chapter 59 + Chapter 53).
- **Survey** the related work, place OmniLLM in the literature, and defend the research contribution (Part X).
- **Choose** the right Python / NAOqi / Choregraphe / LangGraph / ChromaDB pattern for any new sub-problem (Part IX + Chapter 57).
- **Plan** the next year of improvements (Chapter 60).
- **Look up** any term, command, or file in seconds (Appendices).

The codebase will continue to evolve; this book will not. When the book and the source disagree, **trust the source**. The book's job was to get you fluent enough that you can read the source for yourself.

Good luck with the thesis. And remember — the most interesting result might be the one you didn't predict. Embodied LLM rankings might agree with text rankings, or they might not. **Both findings advance the field**.

— *End of book, second edition.*